

**Imperial College  
London**

MEng Individual Project

# **Accelerating data centre communication patterns using eBPF**

**Alex Constantin-Gómez**

Under supervision of Marios Kogias

---

## When the kernel becomes the communication bottleneck

- Large-scale distributed applications in data centres rely on 100s or 1000s machines connected over a network.

## When the kernel becomes the communication bottleneck

- Large-scale distributed applications in data centres rely on 100s or 1000s machines connected over a network.
- This leads to a huge volume of network messages, incurring a significant amount of:
  - User-kernel crossings (via system calls)
  - Context switches
  - In-kernel network stack traversals

## When the kernel becomes the communication bottleneck

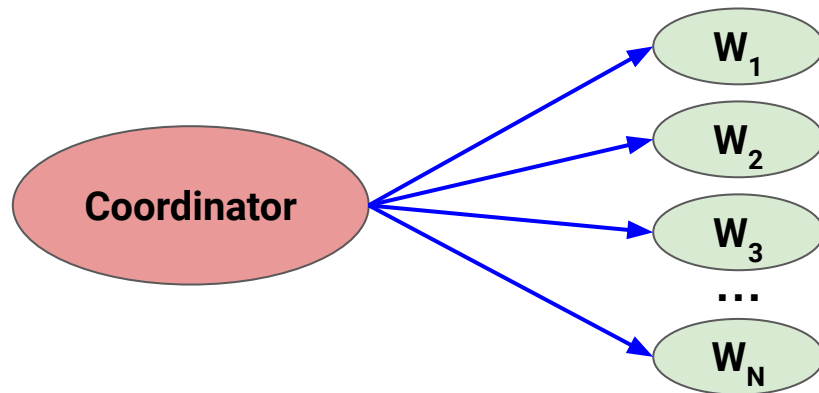
- Large-scale distributed applications in data centres rely on 100s or 1000s machines connected over a network.
- This leads to a huge volume of network messages, incurring a significant amount of:
  - User-kernel crossings (via system calls)
  - Context switches
  - In-kernel network stack traversals
- The default Linux kernel network stack is designed for a small number of long-lived connections, rather than a large number of small messages.
  - High per-packet overhead

## When the kernel becomes the communication bottleneck

- This problem is further exaggerated when high-speed NICs are used
  - This moves the communication bottleneck to the kernel
- This is especially true for applications following *scatter-gather* workloads which generate a large number of messages.
- Scatter-gather communication is a very common pattern in data centre applications (also known as fan-out/fan-in pattern in “*The tail at scale*” article [1])

## Scatter-gather communication

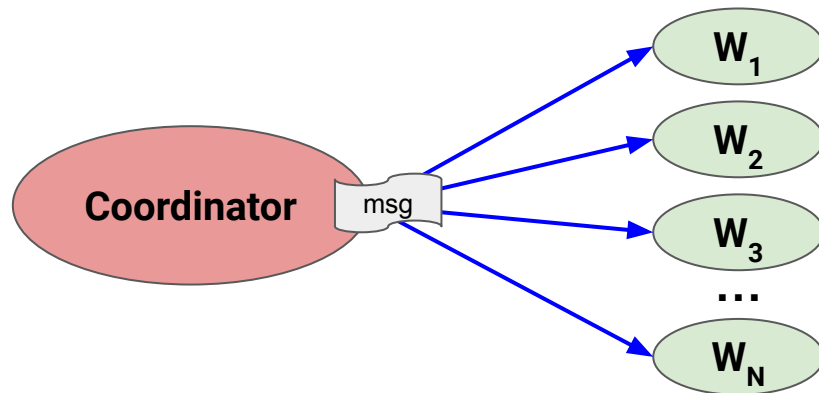
- The *scatter-gather* pattern refers to a single coordinator node scattering a message to multiple worker nodes and then gathering their responses into a single result.
- Typically, the coordinator performs some aggregation logic on the responses to produce a final reduced result (aka *scatter-reduce*).



**Scatter**  
(Fan-out)

## Scatter-gather communication

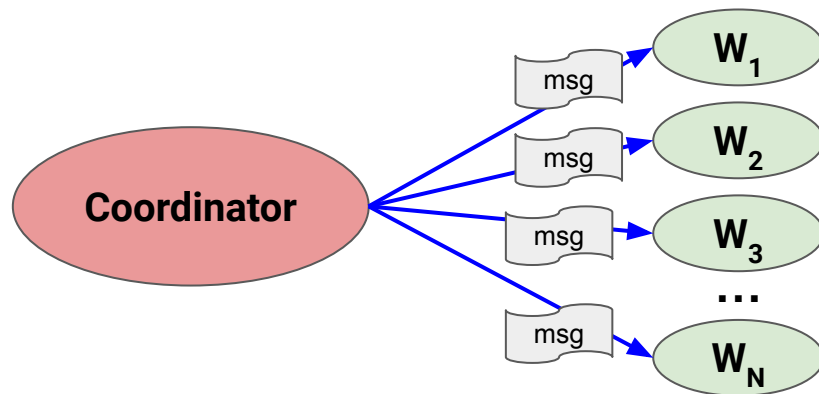
- The *scatter-gather* pattern refers to a single coordinator node scattering a message to multiple worker nodes and then gathering their responses into a single result.
- Typically, the coordinator performs some aggregation logic on the responses to produce a final reduced result (aka *scatter-reduce*).



**Scatter**  
(Fan-out)

## Scatter-gather communication

- The *scatter-gather* pattern refers to a single coordinator node scattering a message to multiple worker nodes and then gathering their responses into a single result.
- Typically, the coordinator performs some aggregation logic on the responses to produce a final reduced result (aka *scatter-reduce*).




**Scatter**  
(Fan-out)



## Scatter-gather communication

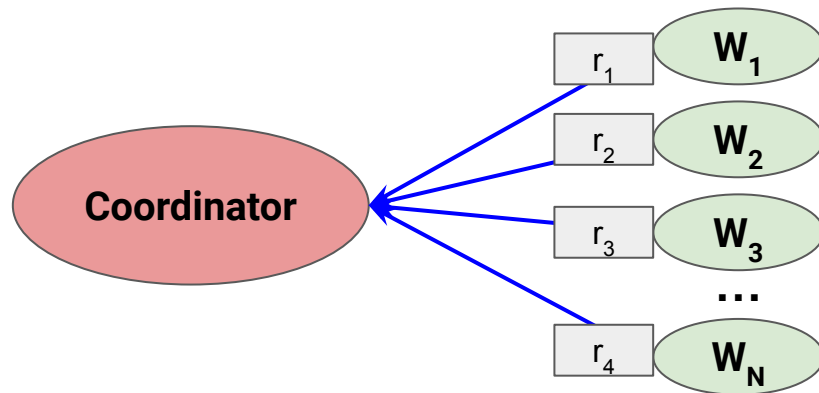
- The *scatter-gather* pattern refers to a single coordinator node scattering a message to multiple worker nodes and then gathering their responses into a single result.
- Typically, the coordinator performs some aggregation logic on the responses to produce a final reduced result (aka *scatter-reduce*).



**Gather**  
(Fan-in)

## Scatter-gather communication

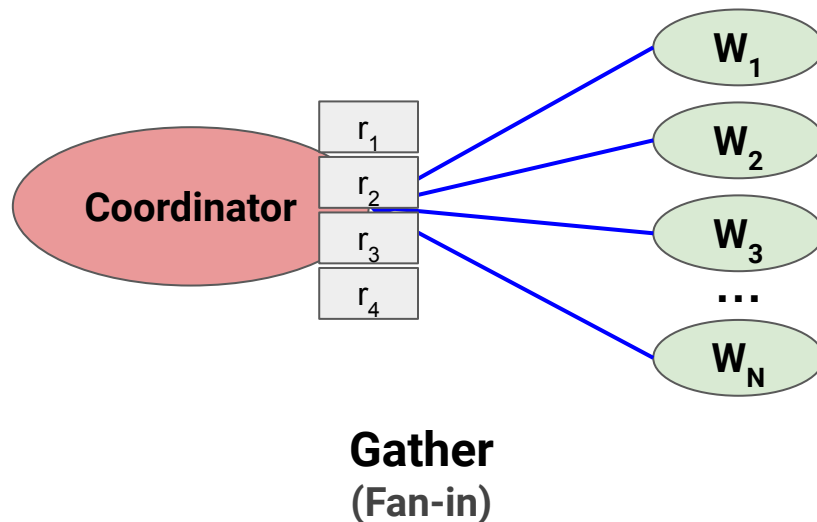
- The *scatter-gather* pattern refers to a single coordinator node scattering a message to multiple worker nodes and then gathering their responses into a single result.
- Typically, the coordinator performs some aggregation logic on the responses to produce a final reduced result (aka *scatter-reduce*).



**Gather**  
(Fan-in)

## Scatter-gather communication

- The *scatter-gather* pattern refers to a single coordinator node scattering a message to multiple worker nodes and then gathering their responses into a single result.
- Typically, the coordinator performs some aggregation logic on the responses to produce a final reduced result (aka *scatter-reduce*).



# Scatter-gather communication

- Many applications follow scatter-gather communication:
  - Distributed machine learning training
  - Distributed graph processing
  - Distributed (network) file systems

## Scatter-gather communication

- Many applications follow scatter-gather communication:
  - Distributed machine learning training
  - Distributed graph processing
  - Distributed (network) file systems
- **Problem:** scatter-gather communication incurs a large volume of messages.
  - ✗ Excessive number of system calls (and context switches)
  - ✗ Excessive number of kernel network stack traversals

## Accelerating scatter-gather communication

- **Goal:** accelerate the performance of scatter-gather workloads using readily-available technologies on commodity Linux machines with minimal deployment overhead.
  - ✗ User-space networking and hardware-offloaded solutions not suitable.

## Accelerating scatter-gather communication

- **Goal:** accelerate the performance of scatter-gather workloads using readily-available technologies on commodity Linux machines with minimal deployment overhead.
  - ✗ User-space networking and hardware-offloaded solutions not suitable.
- Our approach uses **eBPF** as a performance accelerator.

# Accelerating scatter-gather communication

- **Goal:** accelerate the performance of scatter-gather workloads using readily-available technologies on commodity Linux machines with minimal deployment overhead.
  - ✗ User-space networking and hardware-offloaded solutions not suitable.
- Our approach uses **eBPF** as a performance accelerator.
- Presentation outline:
  - a. Introduction to eBPF
  - b. *sgbpf*, an eBPF-accelerated scatter-gather primitive
  - c. Demonstration of the *sgbpf* API
  - d. Performance evaluation



---

## Introduction to eBPF

- eBPF is a technology that allows users to write code that runs inside the Linux kernel.

---

## Introduction to eBPF

- eBPF is a technology that allows users to write code that runs inside the Linux kernel.
- Guarantees both:
  - **Safety**, via the eBPF static verifier
  - **Performance**, via JIT compilation

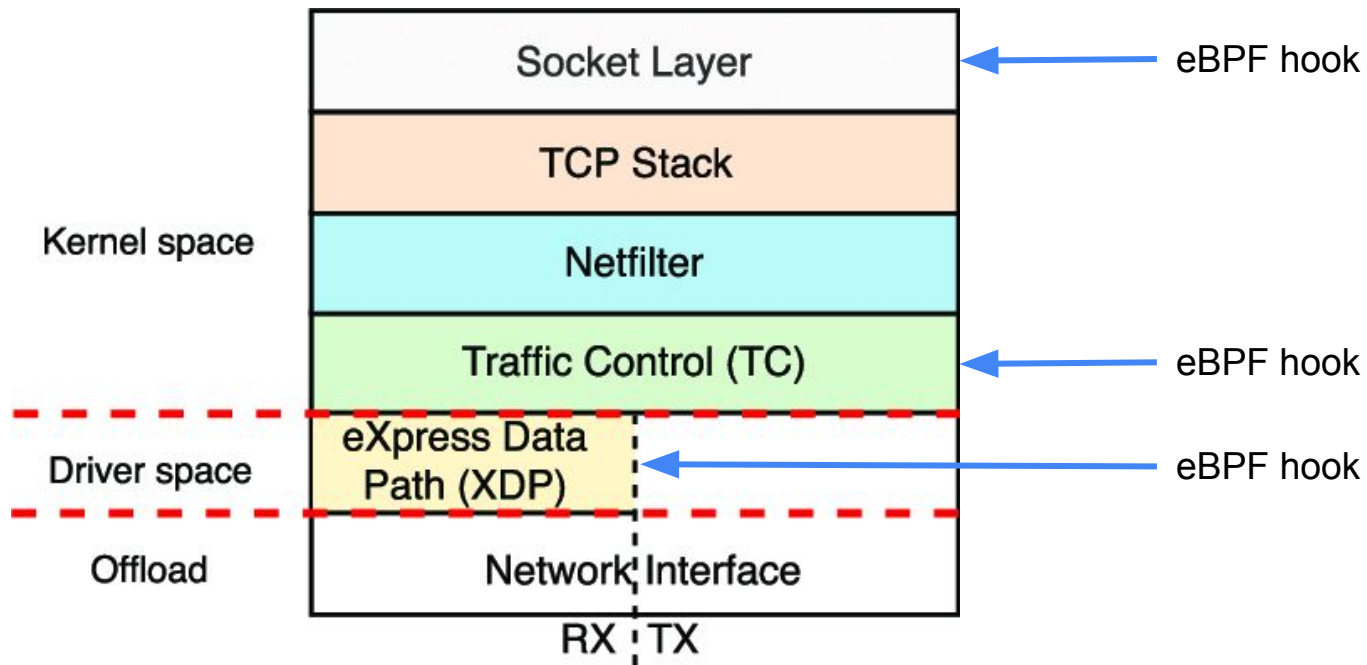
# Introduction to eBPF

- eBPF is a technology that allows users to write code that runs inside the Linux kernel.
- Guarantees both:
  - **Safety**, via the eBPF static verifier
  - **Performance**, via JIT compilation
- Users write their code as eBPF programs, which are event-driven functions that are triggered on the arrival of a kernel event.
  - Programs are attached to a specific hook point inside the kernel
  - Written in a high level language like C and compiled to eBPF bytecode

# Introduction to eBPF

- eBPF is a technology that allows users to write code that runs inside the Linux kernel.
- Guarantees both:
  - **Safety**, via the eBPF static verifier
  - **Performance**, via JIT compilation
- Users write their code as eBPF programs, which are event-driven functions that are triggered on the arrival of a kernel event.
  - Programs are attached to a specific hook point inside the kernel
  - Written in a high level language like C and compiled to eBPF bytecode
- State can be managed using eBPF maps
  - Program  $\leftrightarrow$  program communication
  - User-space  $\leftrightarrow$  program communication

# Network hooks in eBPF



---

## Using eBPF to accelerate network applications

- Applications such as Memcached [2] and Multi-Paxos consensus [3] have seen successful performance boosts with eBPF.

---

## Using eBPF to accelerate network applications

- Applications such as Memcached [2] and Multi-Paxos consensus [3] have seen successful performance boosts with eBPF.
- By performing application logic before the kernel networking stack, high packet processing rates can be achieved by minimising user-kernel crossings.

## Using eBPF to accelerate network applications

- Applications such as Memcached [2] and Multi-Paxos consensus [3] have seen successful performance boosts with eBPF.
- By performing application logic before the kernel networking stack, high packet processing rates can be achieved by minimising user-kernel crossings.
- Memcached example:
  - Fast replying: prepare a response packet in XDP to serve requests quickly by reading from a working copy of a cache in the kernel using eBPF maps.
- Multi-Paxos example:
  - Wait-on-quorums: count ACK messages from followers in the kernel and only wake up the leader application when enough ACKs have been received.



---

## ***sgbpf*: an eBPF-accelerated scatter-gather primitive**

- *sgbpf* is a library that exposes a scatter-gather network primitive for UDP workloads.

---

## ***sgbpf*: an eBPF-accelerated scatter-gather primitive**

- *sgbpf* is a library that exposes a scatter-gather network primitive for UDP workloads.
- Available as a C++ API which can be integrated into network applications.

## ***sgbpf*: an eBPF-accelerated scatter-gather primitive**

- *sgbpf* is a library that exposes a scatter-gather network primitive for UDP workloads.
- Available as a C++ API which can be integrated into network applications.
- It uses eBPF under-the-hood to perform custom packet processing inside the kernel
  - Minimises system calls
  - Minimises network stack traversals

## ***sgbpf*: an eBPF-accelerated scatter-gather primitive**

- *sgbpf* is a library that exposes a scatter-gather network primitive for UDP workloads.
- Available as a C++ API which can be integrated into network applications.
- It uses eBPF under-the-hood to perform custom packet processing inside the kernel
  - Minimises system calls
  - Minimises network stack traversals
- Supports several other features:
  - “All-gather” mode: broadcast final aggregated result to all workers
  - Configurable timeouts and completions policies
  - Multiple data delivery modes available to access aggregated data

---

## Accelerating the scatter phase

- This is essentially a broadcast operation.
- Naive implementation incurs a `sendto()` syscall per worker.

## Accelerating the scatter phase

- This is essentially a broadcast operation.
- Naive implementation incurs a `sendto()` syscall per worker.
- Inefficiencies:

## Accelerating the scatter phase

- This is essentially a broadcast operation.
- Naive implementation incurs a `sendto( )` syscall per worker.
- Inefficiencies:
  - ✗ Performs a syscall per worker (obviously)

## Accelerating the scatter phase

- This is essentially a broadcast operation.
- Naive implementation incurs a `sendto( )` syscall per worker.
- Inefficiencies:
  - ✗ Performs a syscall per worker (obviously)
  - ✗ Performs a user-to-kernel copy of the message per worker



## Accelerating the scatter phase

- This is essentially a broadcast operation.
- Naive implementation incurs a `sendto( )` syscall per worker.
- Inefficiencies:
  - ✗ Performs a syscall per worker (obviously)
  - ✗ Performs a user-to-kernel copy of the message per worker
  - ✗ Performs a network stack traversal per worker

## Accelerating the scatter phase

- This is essentially a broadcast operation.
- Naive implementation incurs a `sendto( )` syscall per worker.
- Inefficiencies:
  - ✗ Performs a syscall per worker (obviously)
  - ✗ Performs a user-to-kernel copy of the message per worker
  - ✗ Performs a network stack traversal per worker
- How can we do better?

## Accelerating the scatter phase

- This is essentially a broadcast operation.
- Naive implementation incurs a `sendto( )` syscall per worker.
- Inefficiencies:
  - ✗ Performs a syscall per worker (obviously)
  - ✗ Performs a user-to-kernel copy of the message per worker
  - ✗ Performs a network stack traversal per worker
- How can we do better?

eBPF

---

## Accelerating the scatter phase


- With APIs such as `io_uring`, we can batch multiple I/O operations into a single syscall

---


## Accelerating the scatter phase

- With APIs such as `io_uring`, we can batch multiple I/O operations into a single syscall
- Our approach with eBPF **tackles all three problems** in one go:

## Accelerating the scatter phase

- With APIs such as `io_uring`, we can batch multiple I/O operations into a single syscall
- Our approach with eBPF **tackles all three problems** in one go:
- Inefficiencies:
  - ✗ Performs a syscall per worker (obviously)  `io_uring` solves this
  - ✗ Performs a user-to-kernel copy of the message per worker
  - ✗ Performs a network stack traversal per worker

## Accelerating the scatter phase

- With APIs such as `io_uring`, we can batch multiple I/O operations into a single syscall
- Our approach with eBPF **tackles all three problems** in one go:
- Inefficiencies:
  - ✗ Performs a syscall per worker (obviously)  `io_uring` solves this
  - ✗ Performs a user-to-kernel copy of the message per worker
  - ✗ Performs a network stack traversal per worker

**eBPF solves all of the above**

## Accelerating the scatter phase

- With APIs such as `io_uring`, we can batch multiple I/O operations into a single syscall
- Our approach with eBPF **tackles all three problems** in one go:
- Inefficiencies:

✗ Performs a syscall per worker (obviously)  `io_uring` solves this

✗ Performs a user-to-kernel copy of the message per worker

✗ Performs a network stack traversal per worker

**eBPF solves all of the above**



---

## Accelerating the scatter phase

- With eBPF, we can perform the broadcast inside the kernel

---

## Accelerating the scatter phase

- With eBPF, we can perform the broadcast inside the kernel
- Specifically, it occurs **after** the main networking stack

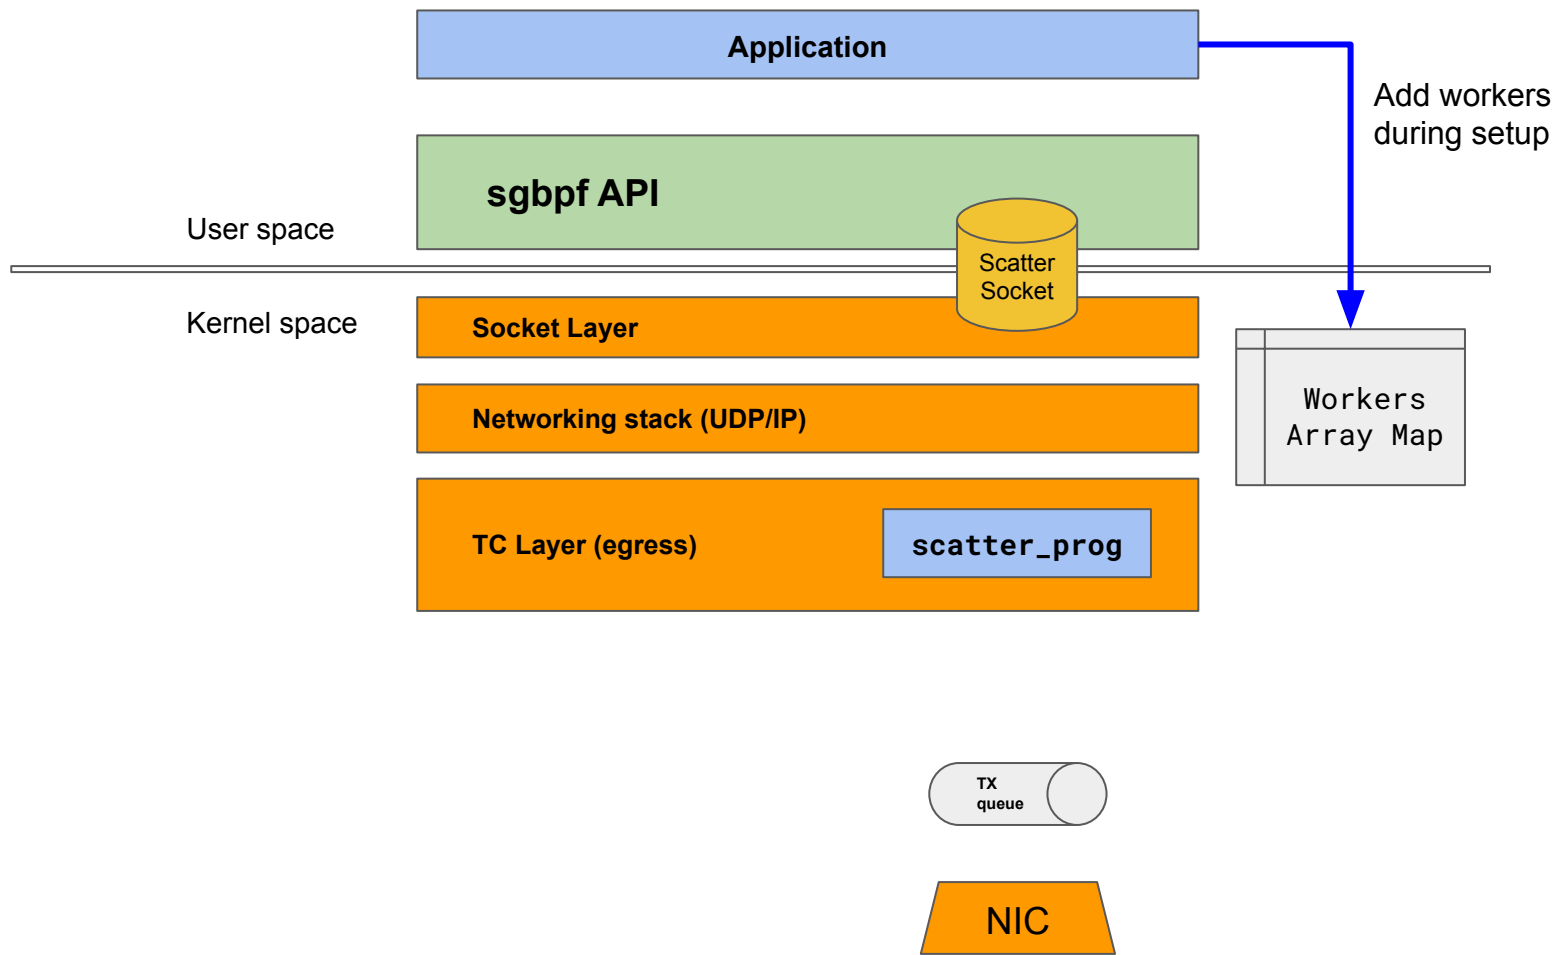
---

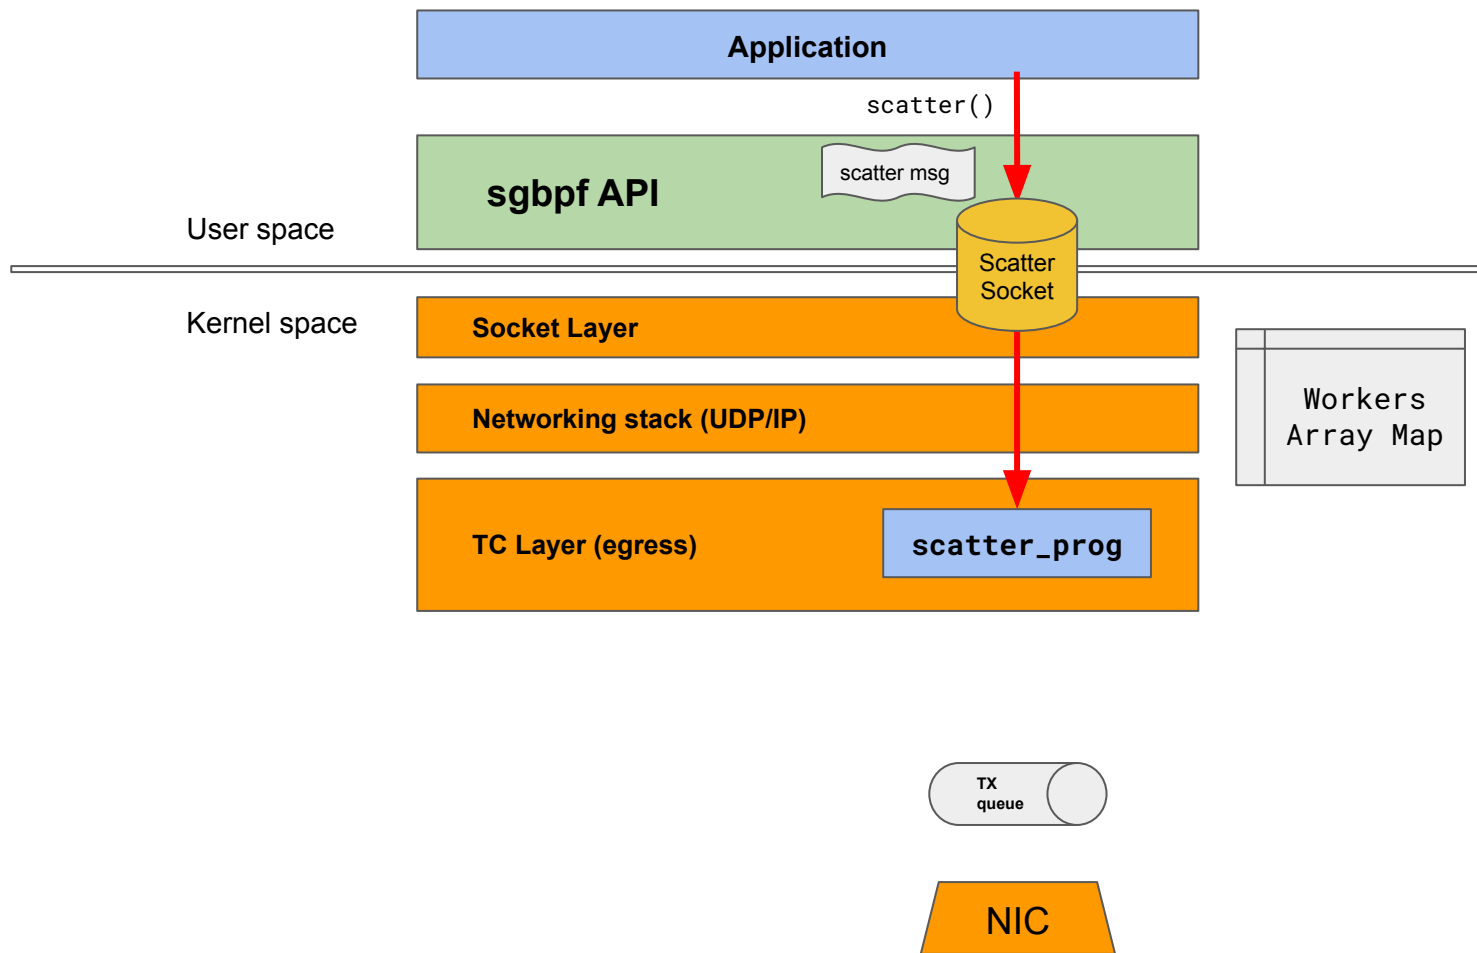
## Accelerating the scatter phase

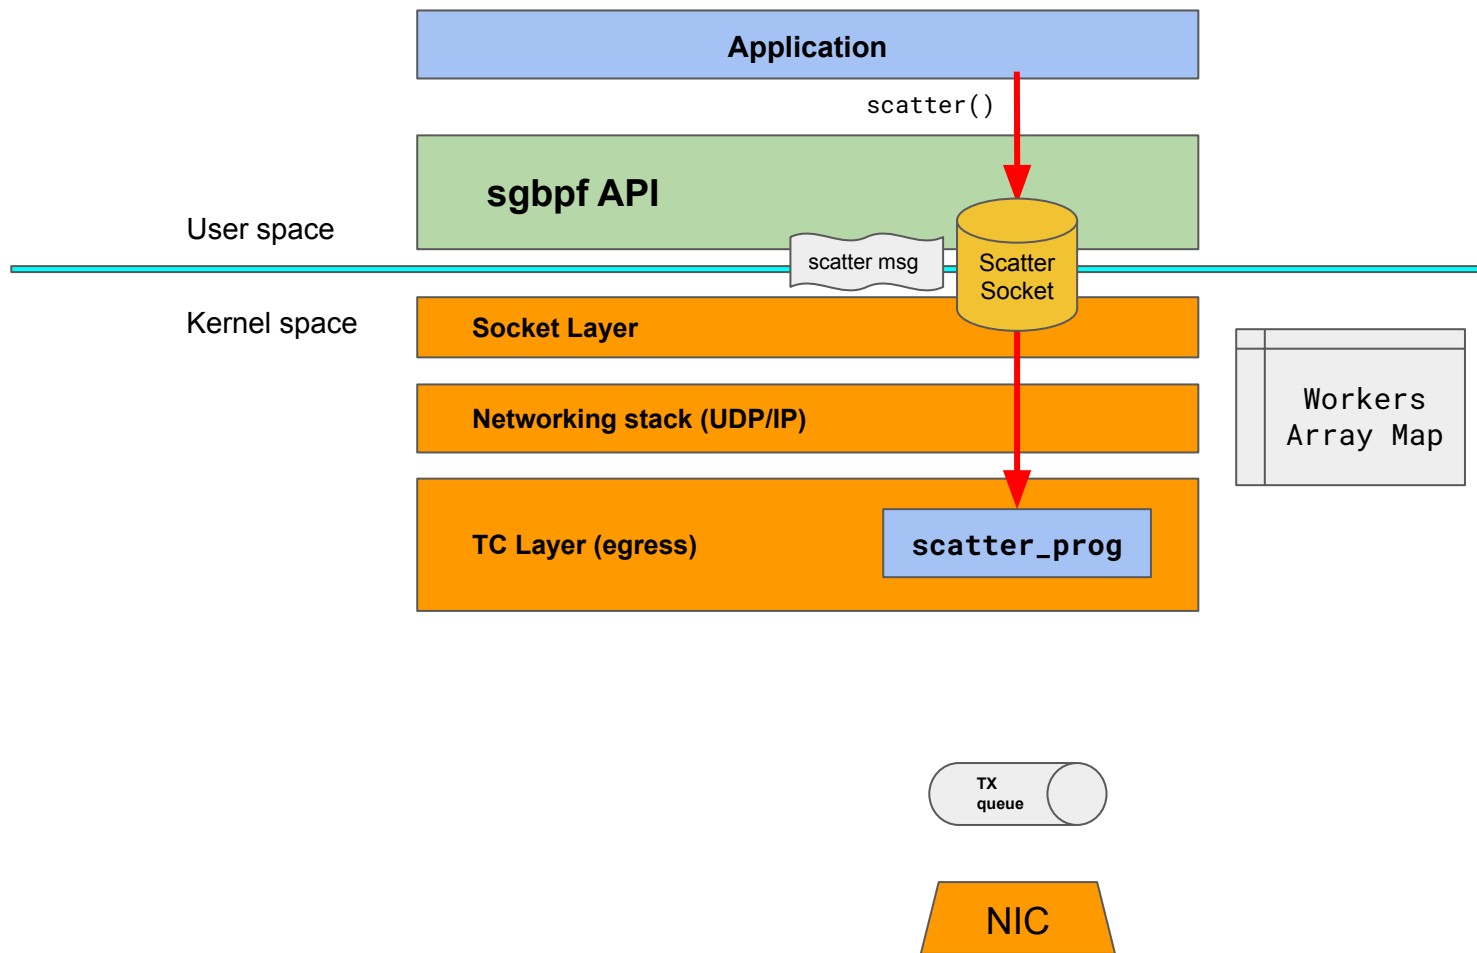
- With eBPF, we can perform the broadcast inside the kernel
- Specifically, it occurs **after** the main networking stack
- The cloned packets do not re-traverse the kernel networking stack

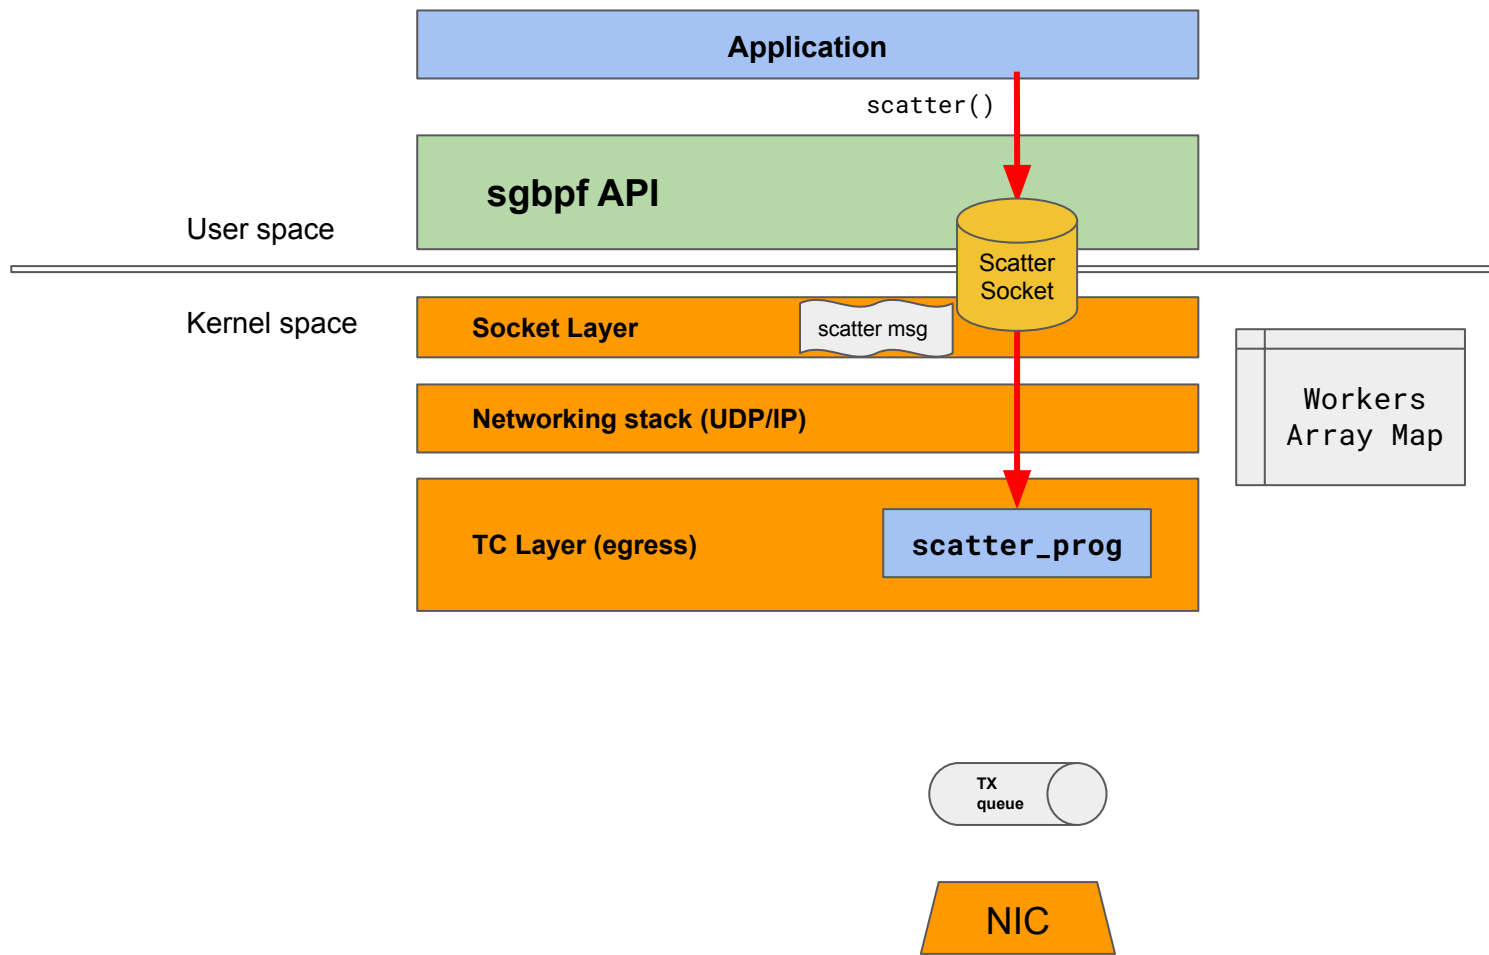
## Accelerating the scatter phase

- With eBPF, we can perform the broadcast inside the kernel
- Specifically, it occurs **after** the main networking stack
- The cloned packets do not re-traverse the kernel networking stack
- Only a single message is written into the socket from the application to trigger the broadcast operation in the eBPF program
  - Hence a single network stack traversal is needed

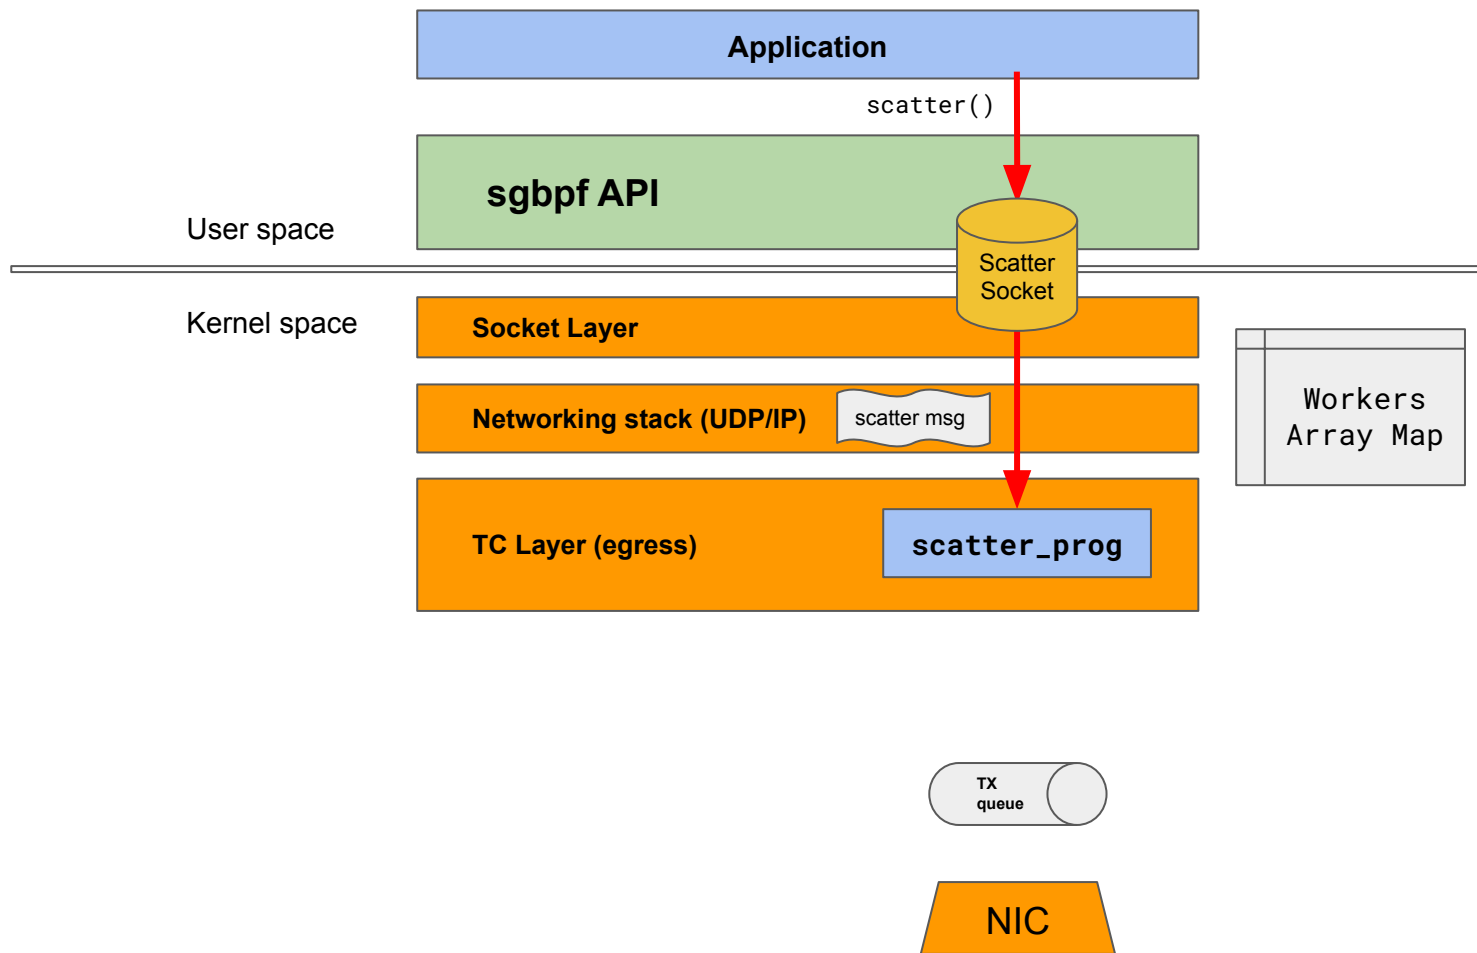


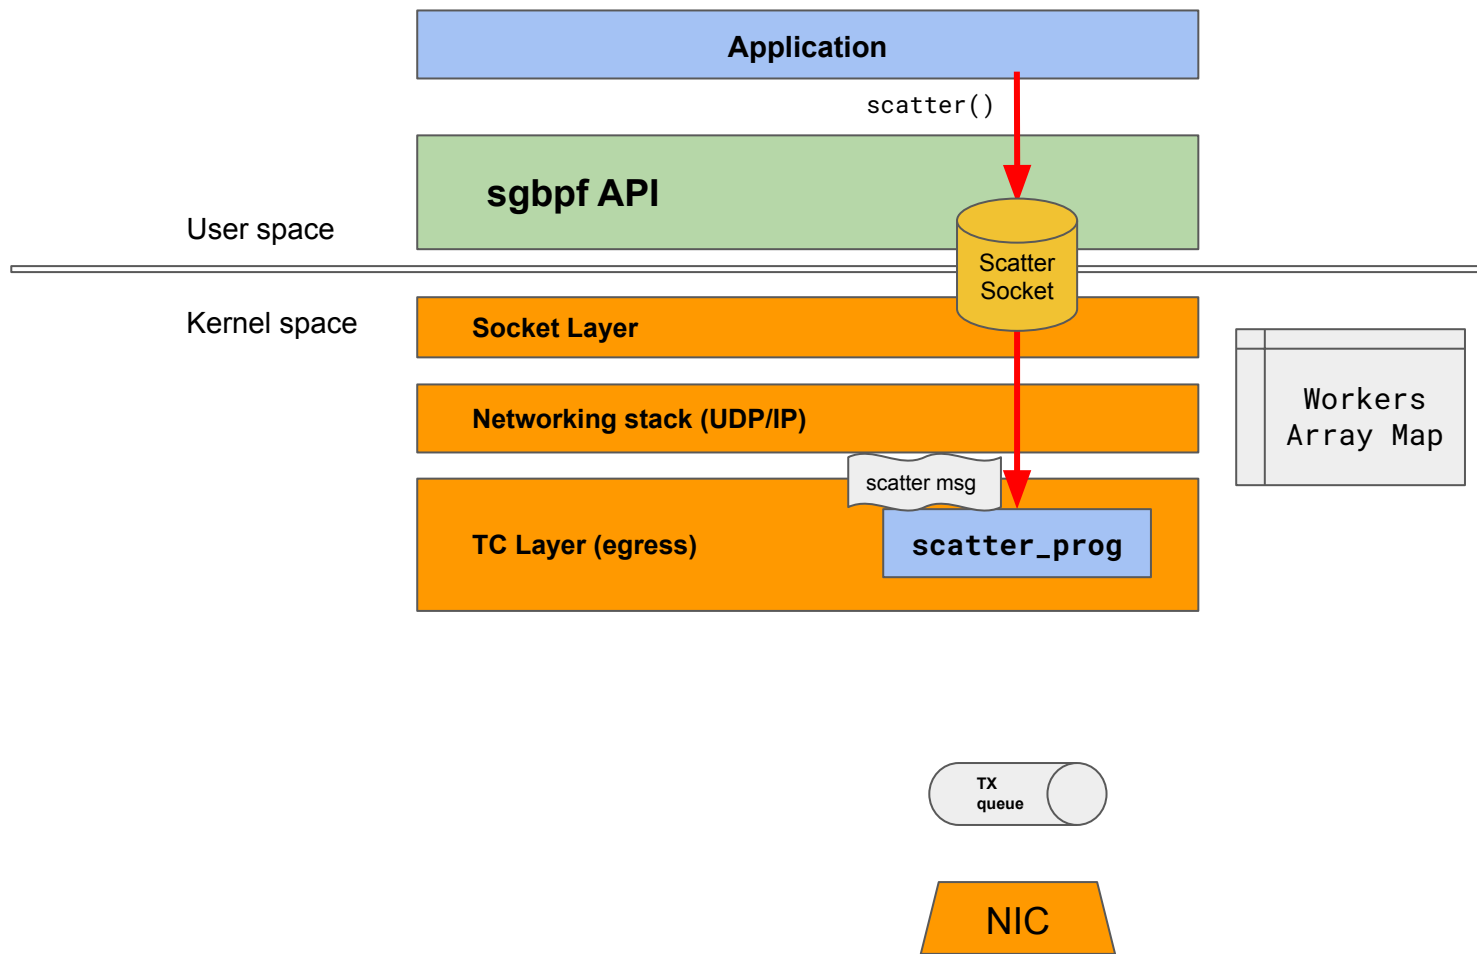


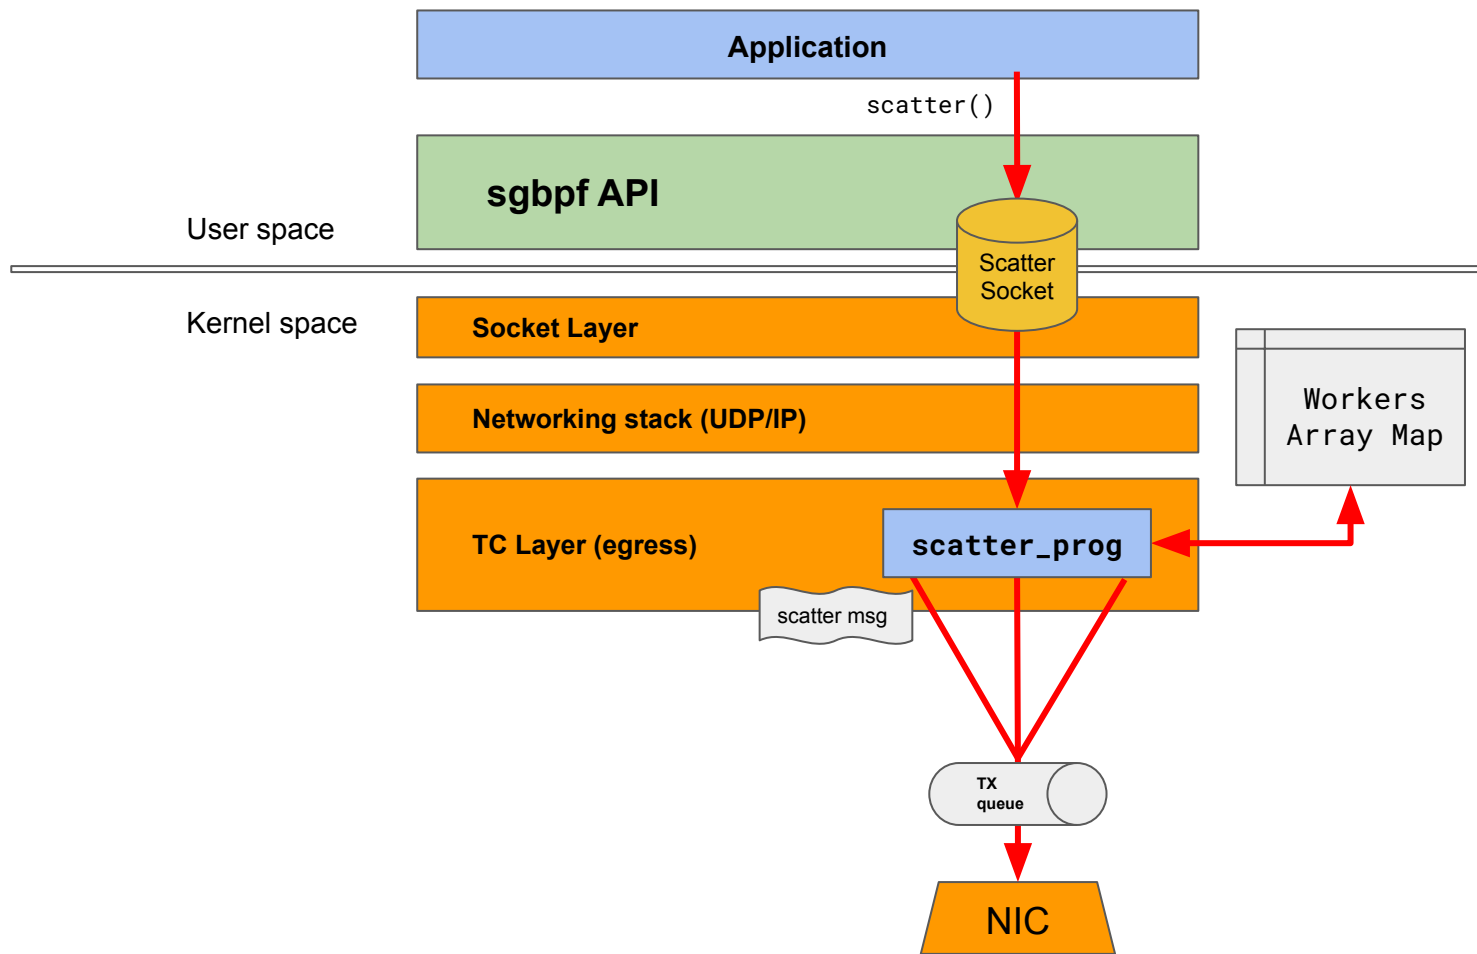


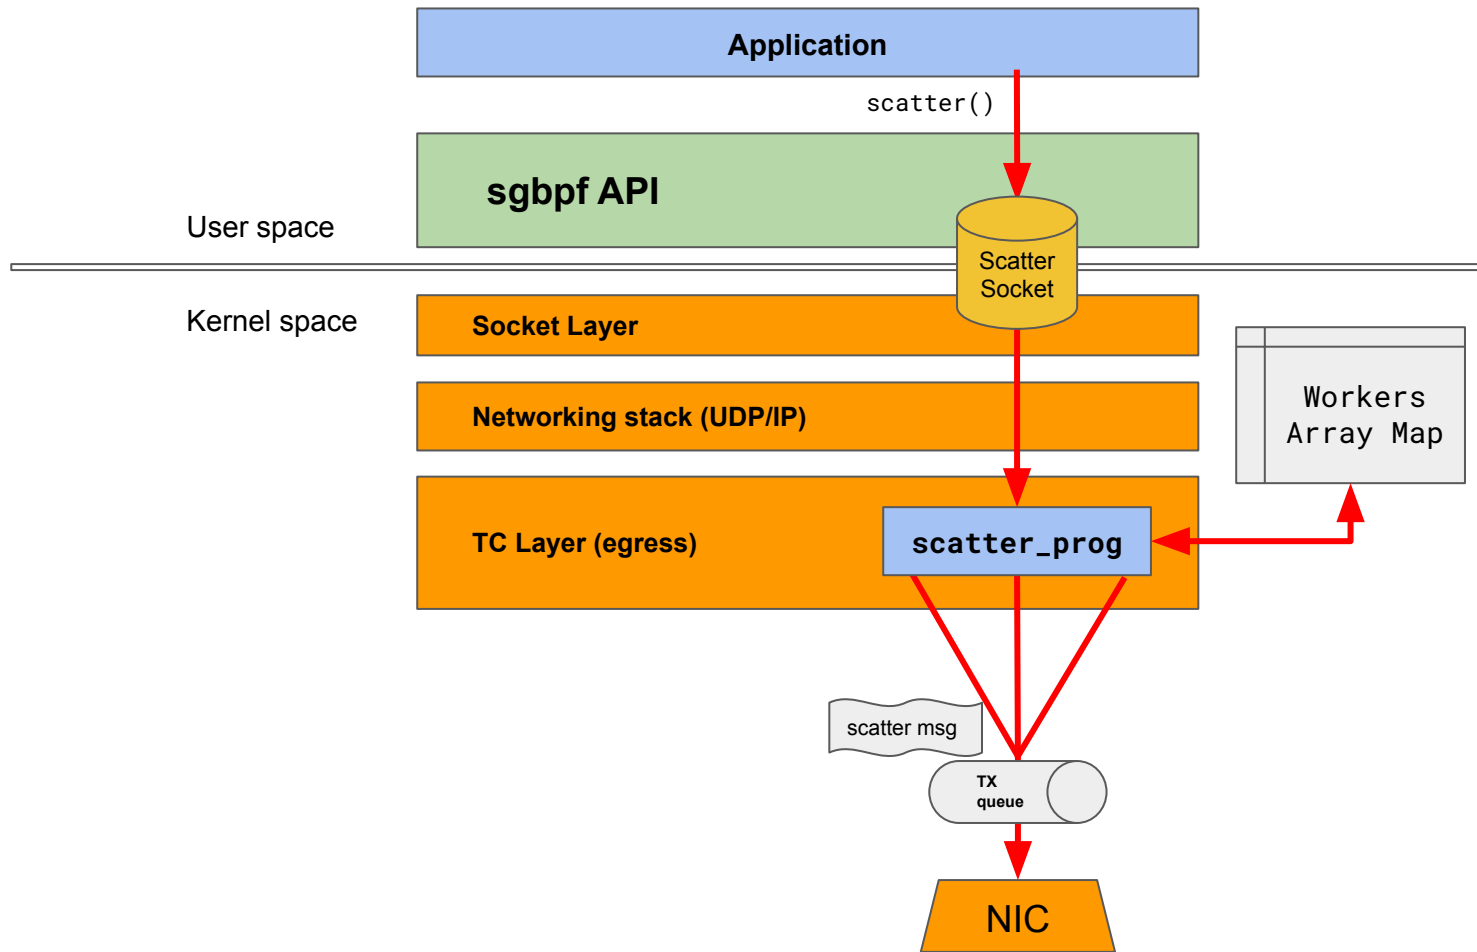


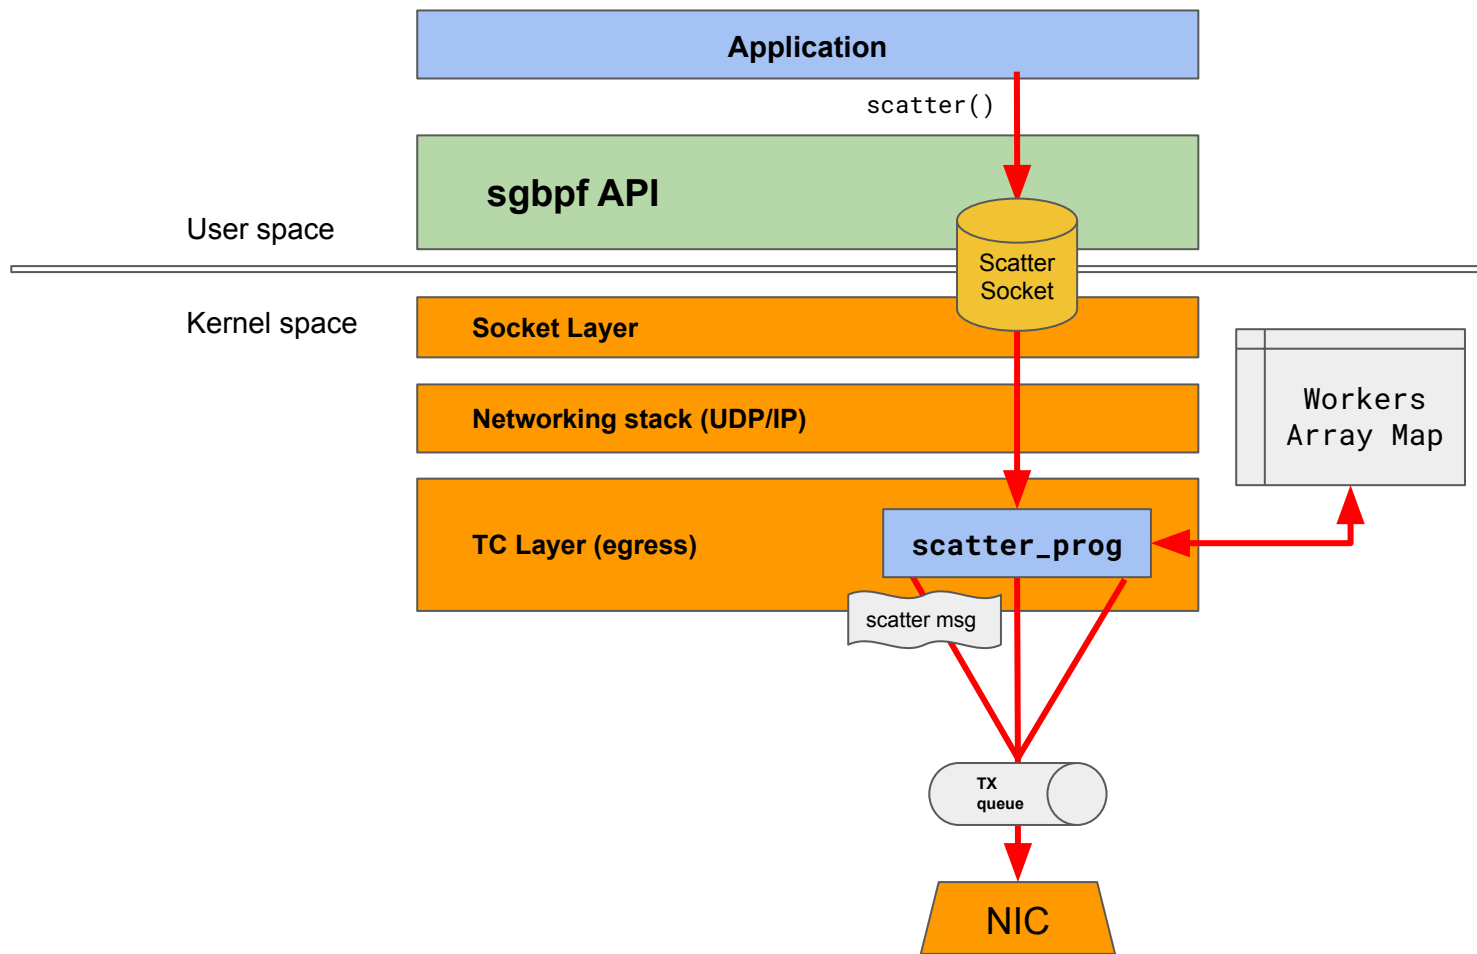


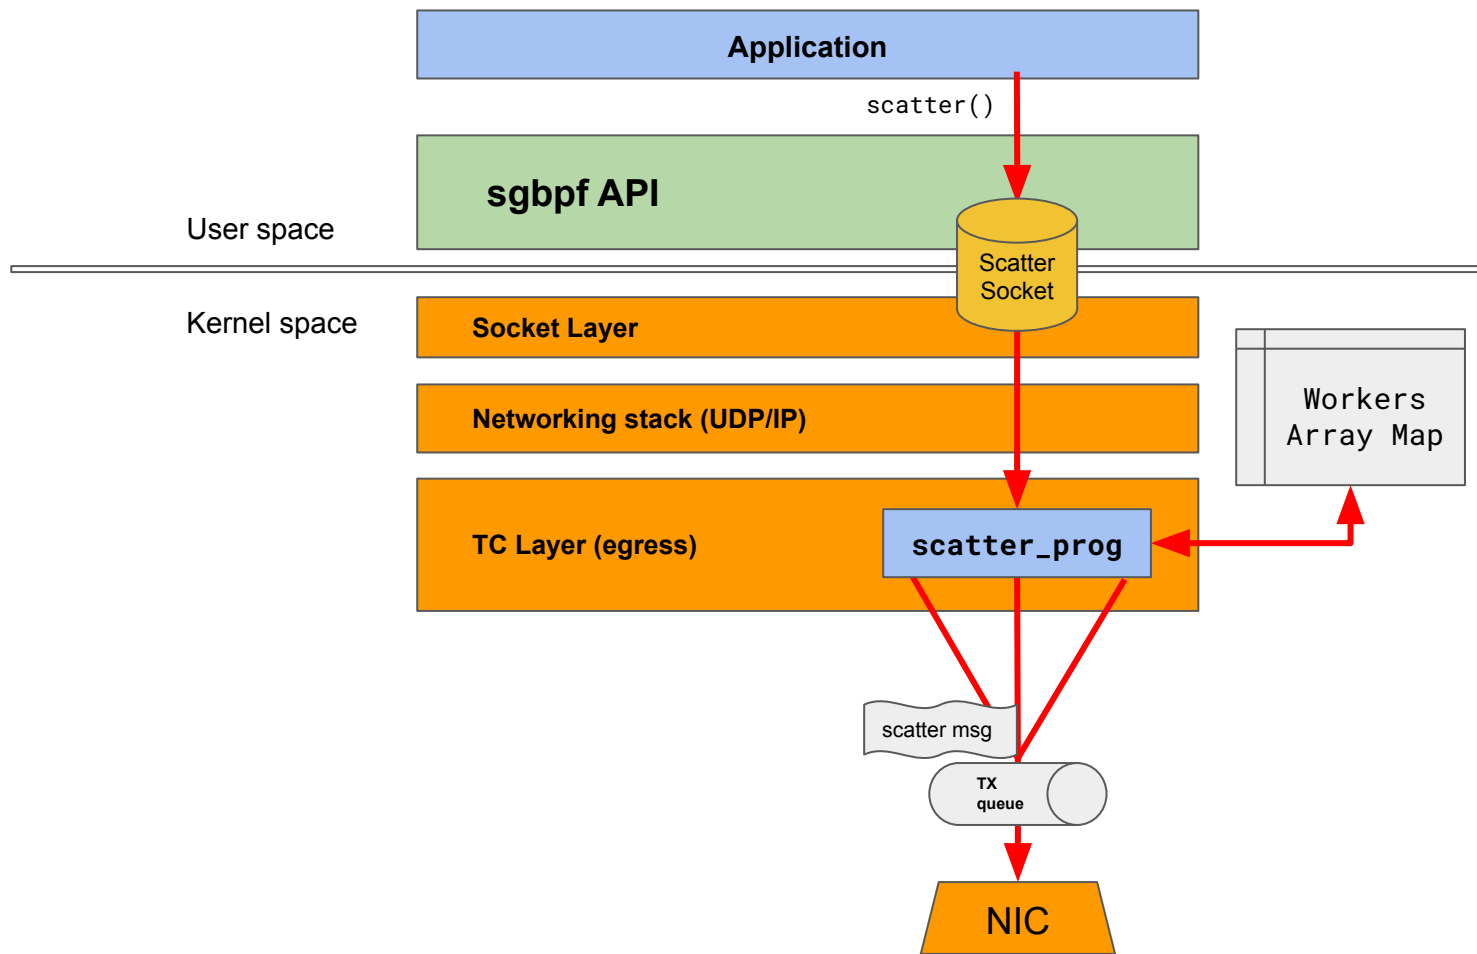


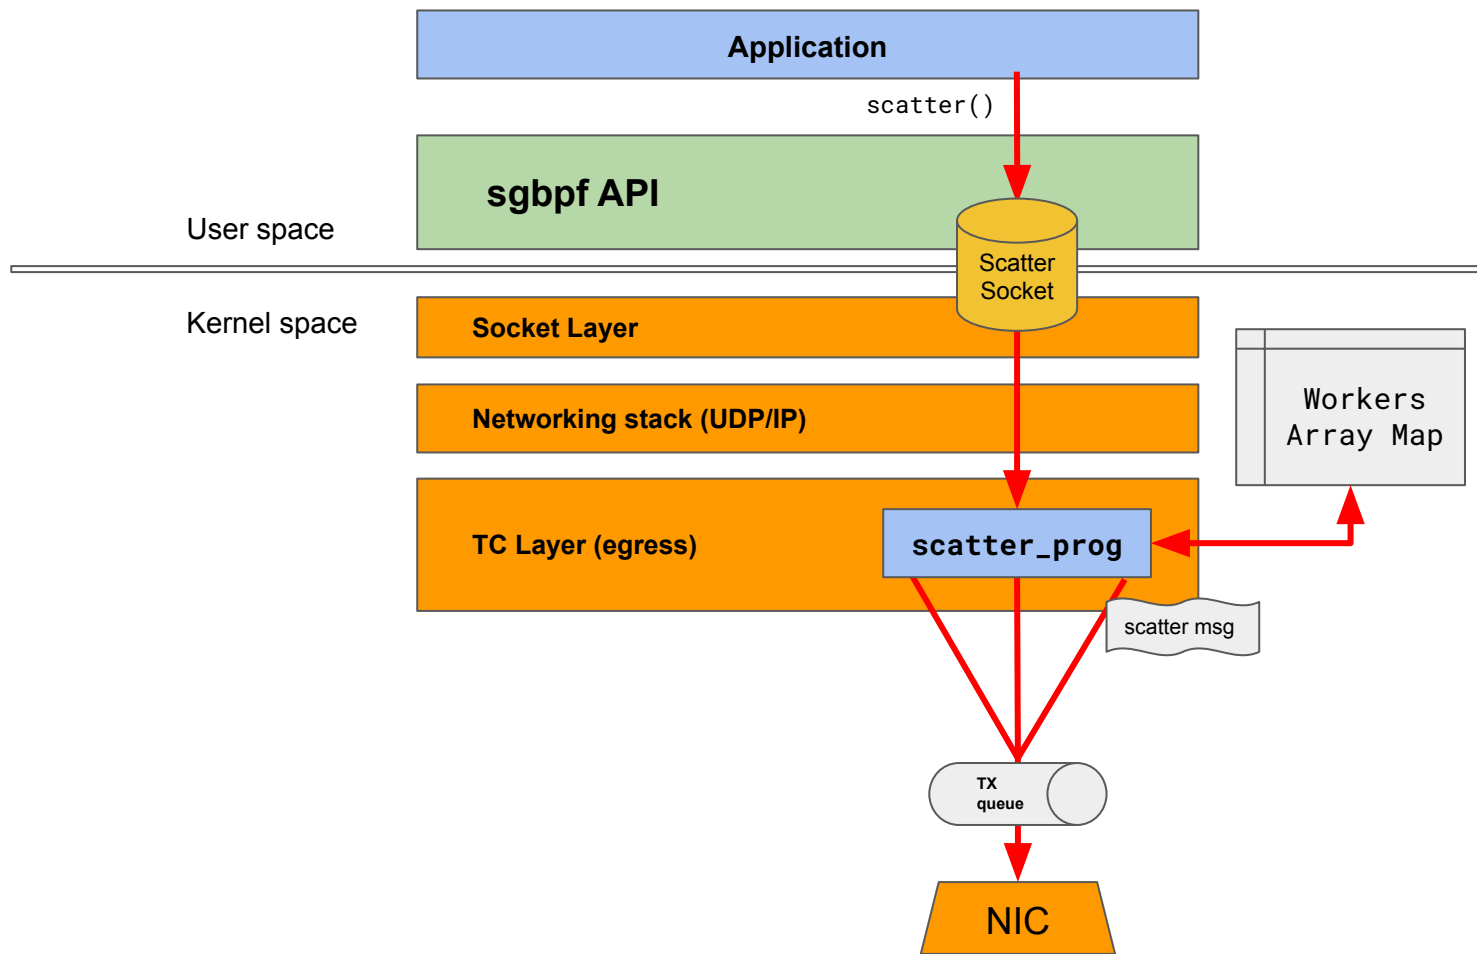


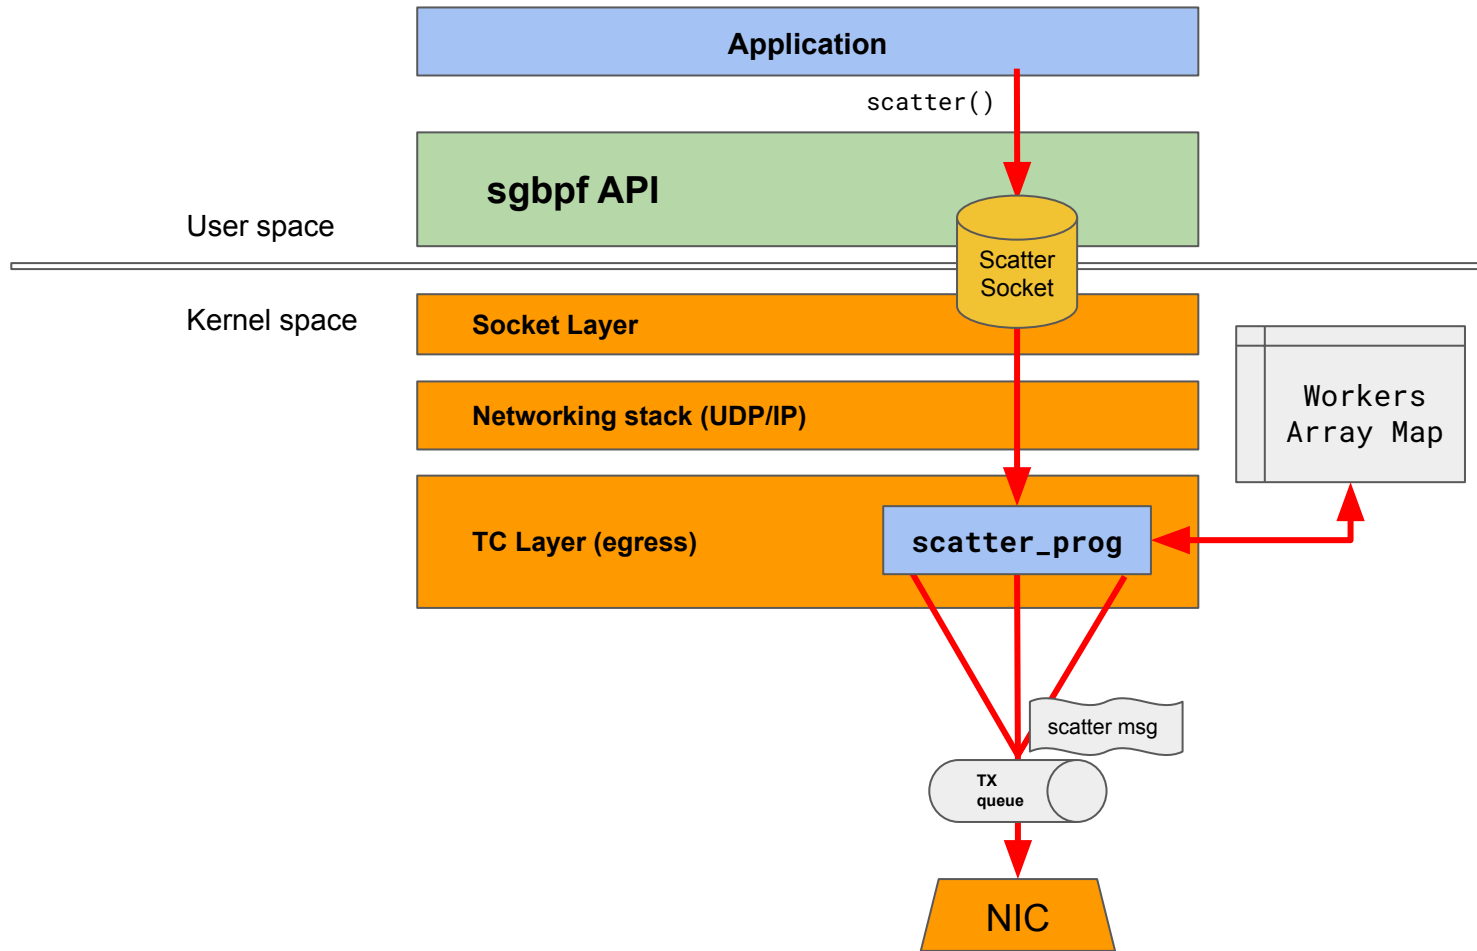














---

## Accelerating the scatter phase

- The eBPF-enabled version now requires only:
  - A single system call
  - A single traversal of the kernel networking stack

## Accelerating the scatter phase

- The eBPF-enabled version now requires only:
  - A single system call
  - A single traversal of the kernel networking stack
- Good start, but the main source of overhead lies in the gather phase.

---

## Accelerating the gather phase

- Can we process incoming responses as early as possible without incurring a kernel-to-user crossing per packet?

## Accelerating the gather phase

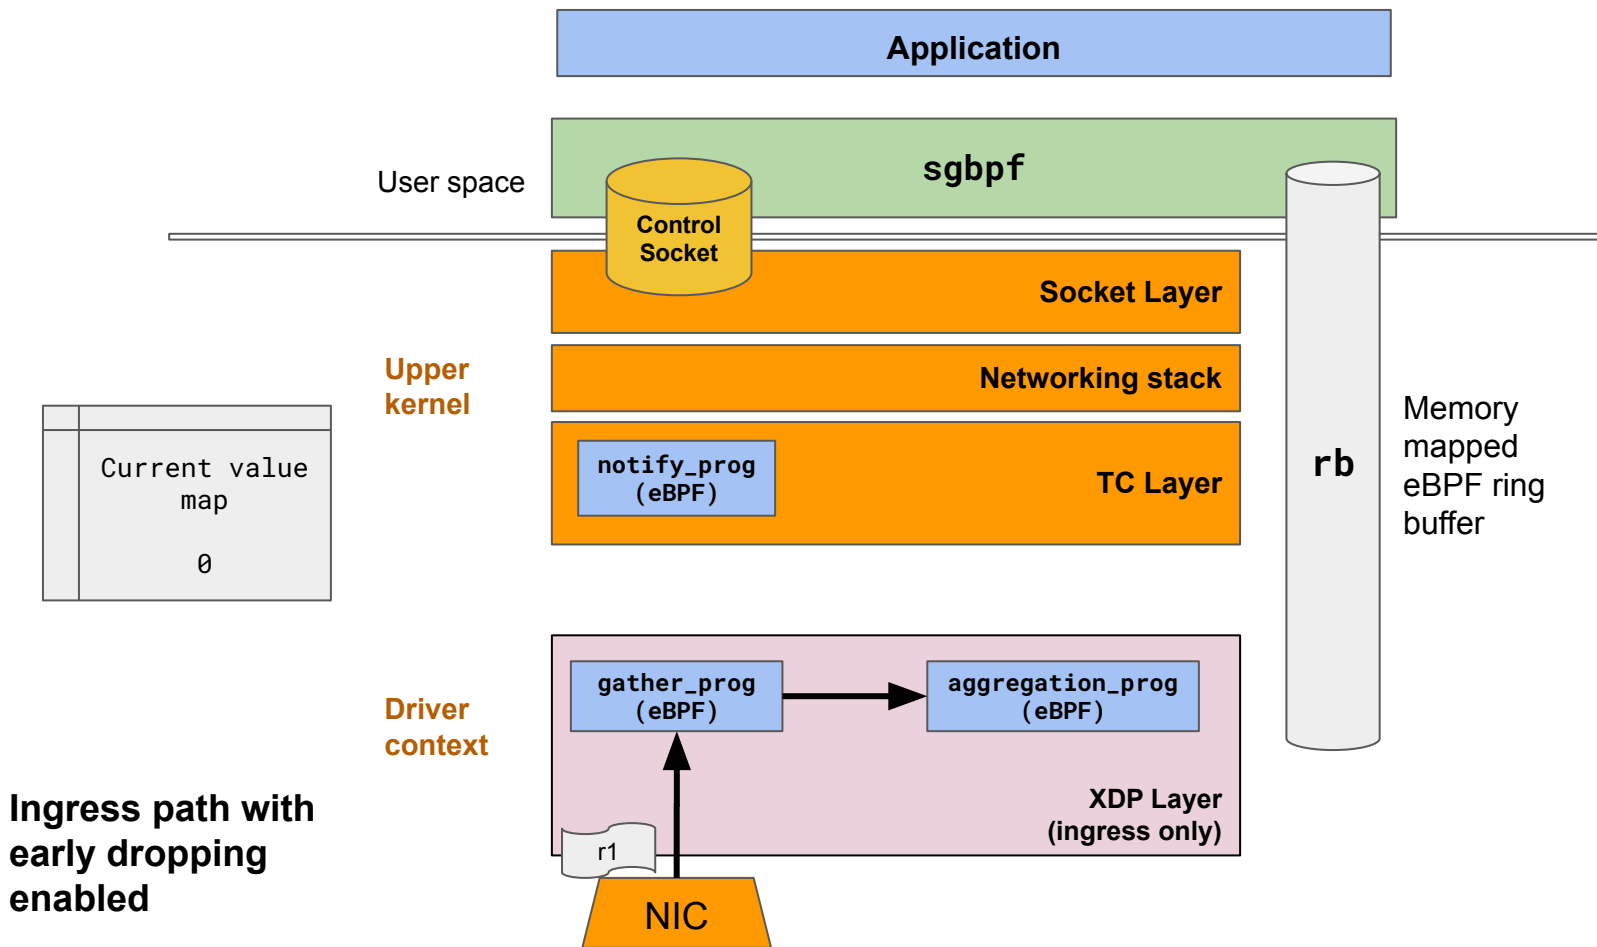
- Can we process incoming responses as early as possible without incurring a kernel-to-user crossing per packet?
- **Idea:** execute the aggregation logic in the kernel using eBPF
  - Perform pre-stack packet processing to bypass the kernel networking stack
  - The application only gets the final aggregated value
  - Minimises the number of syscalls, number of copies and network stack traversals required

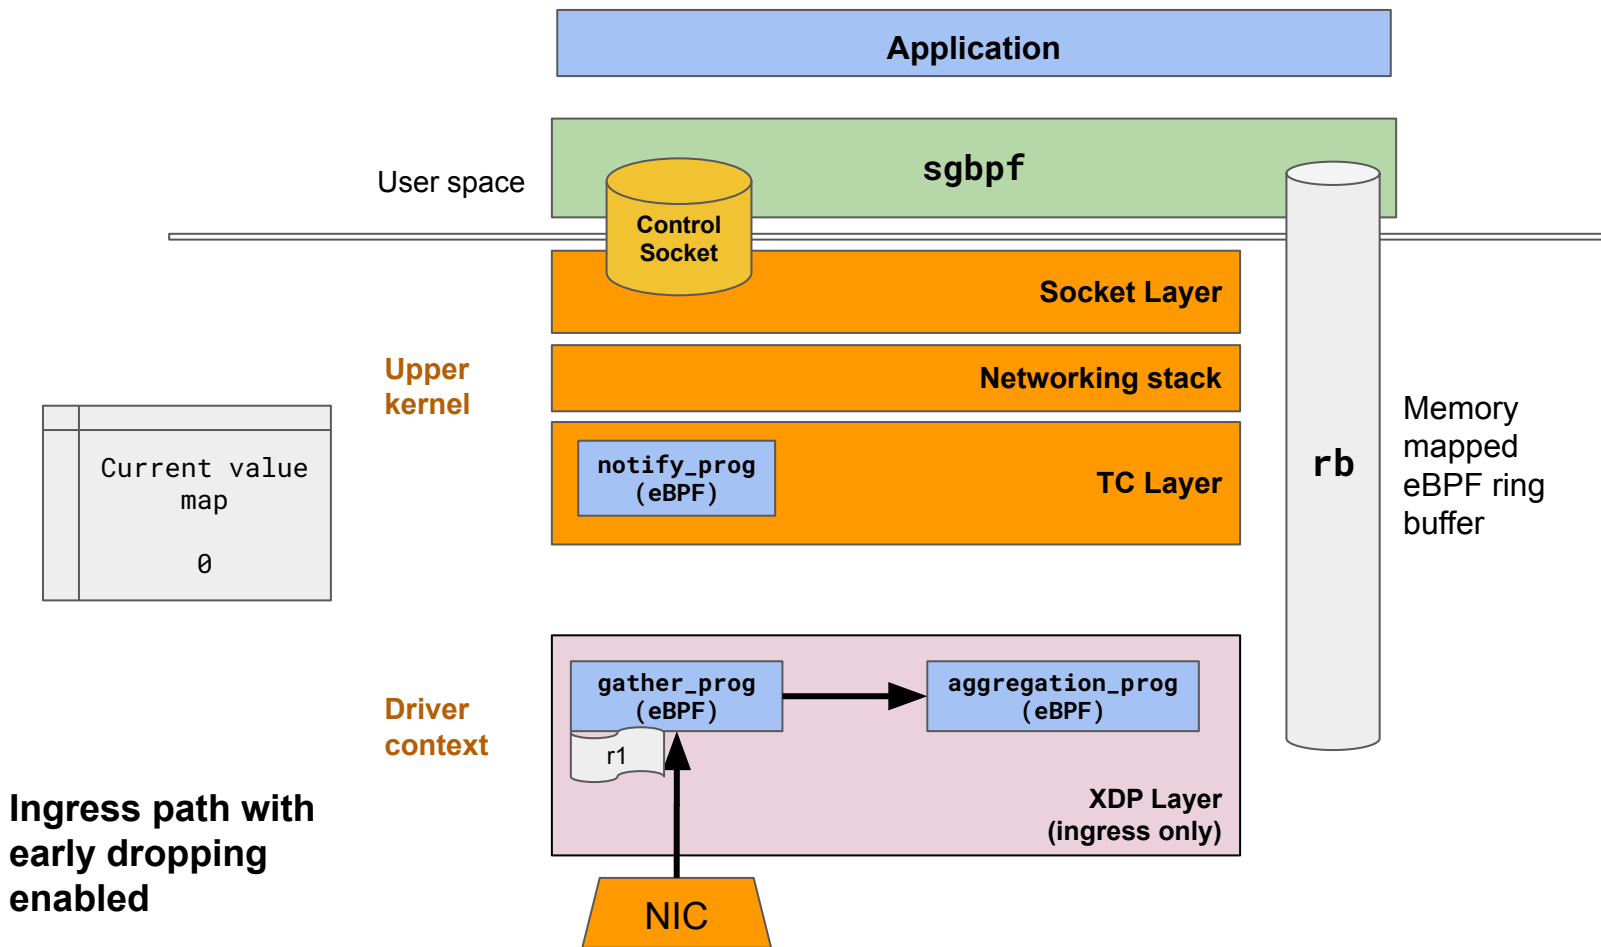
## Accelerating the gather phase

- *sgbpf* relies on pre-stack processing with XDP to aggregate the responses as early as possible
- The aggregation logic is a user-supplied eBPF XDP program
  - Allows flexibility for developer to define the semantics of the aggregation
  - However, must conform to the rules imposed by the eBPF static verifier
- The packet's fate after the aggregation is also specified by the developer:
  - Drop the packet, if no longer needed (referred to as early dropping - recommended)
  - Allow the packet, if needed in the application for further processing

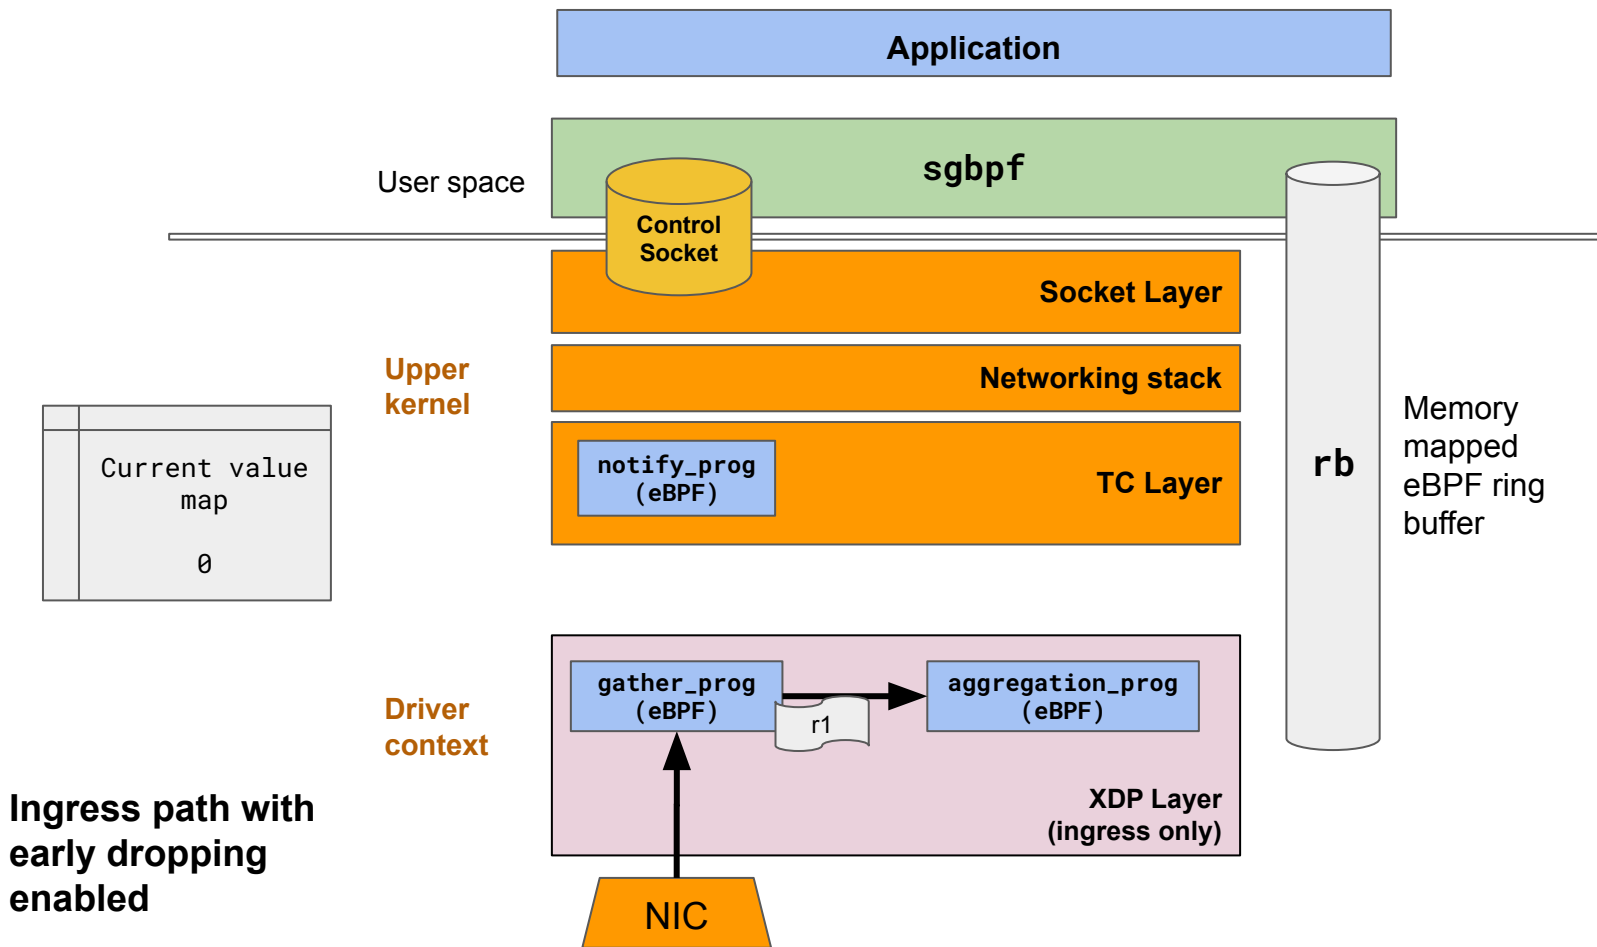
## Accelerating the gather phase

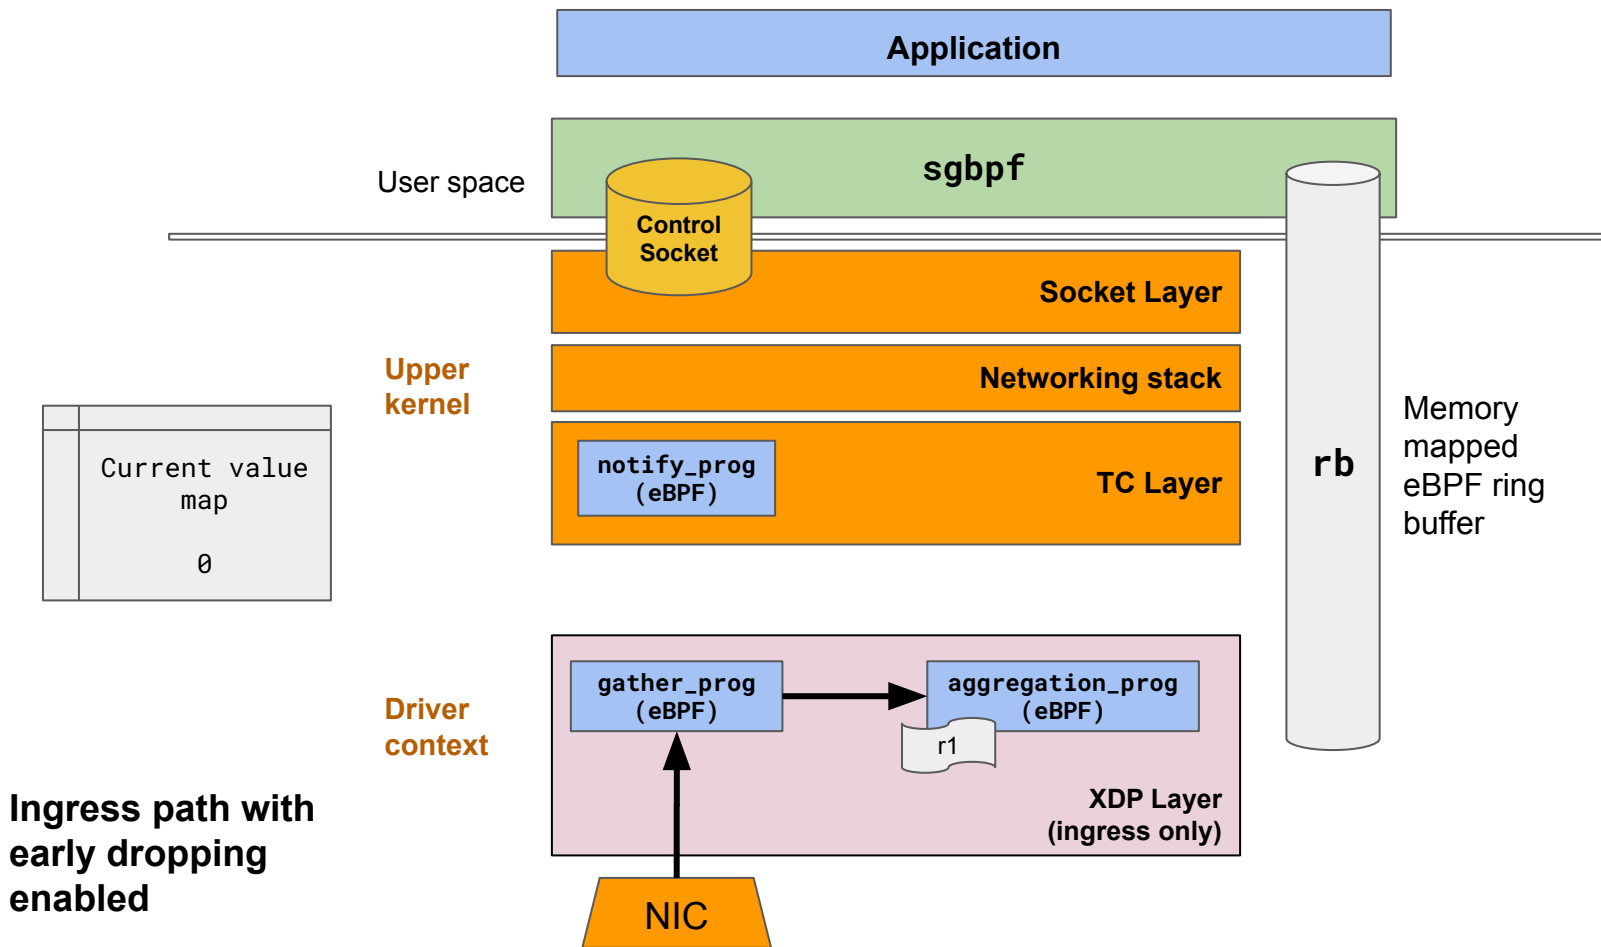
- Once all the worker responses have been aggregated, the final aggregated value is delivered to the user-space application.
- This takes place via one of the following communication channels:
  - **Control socket:** deliver data using a packet redirected into a specific socket monitored by the application
  - **Ring buffer:** use a shared mmaped buffer to deliver data to the application

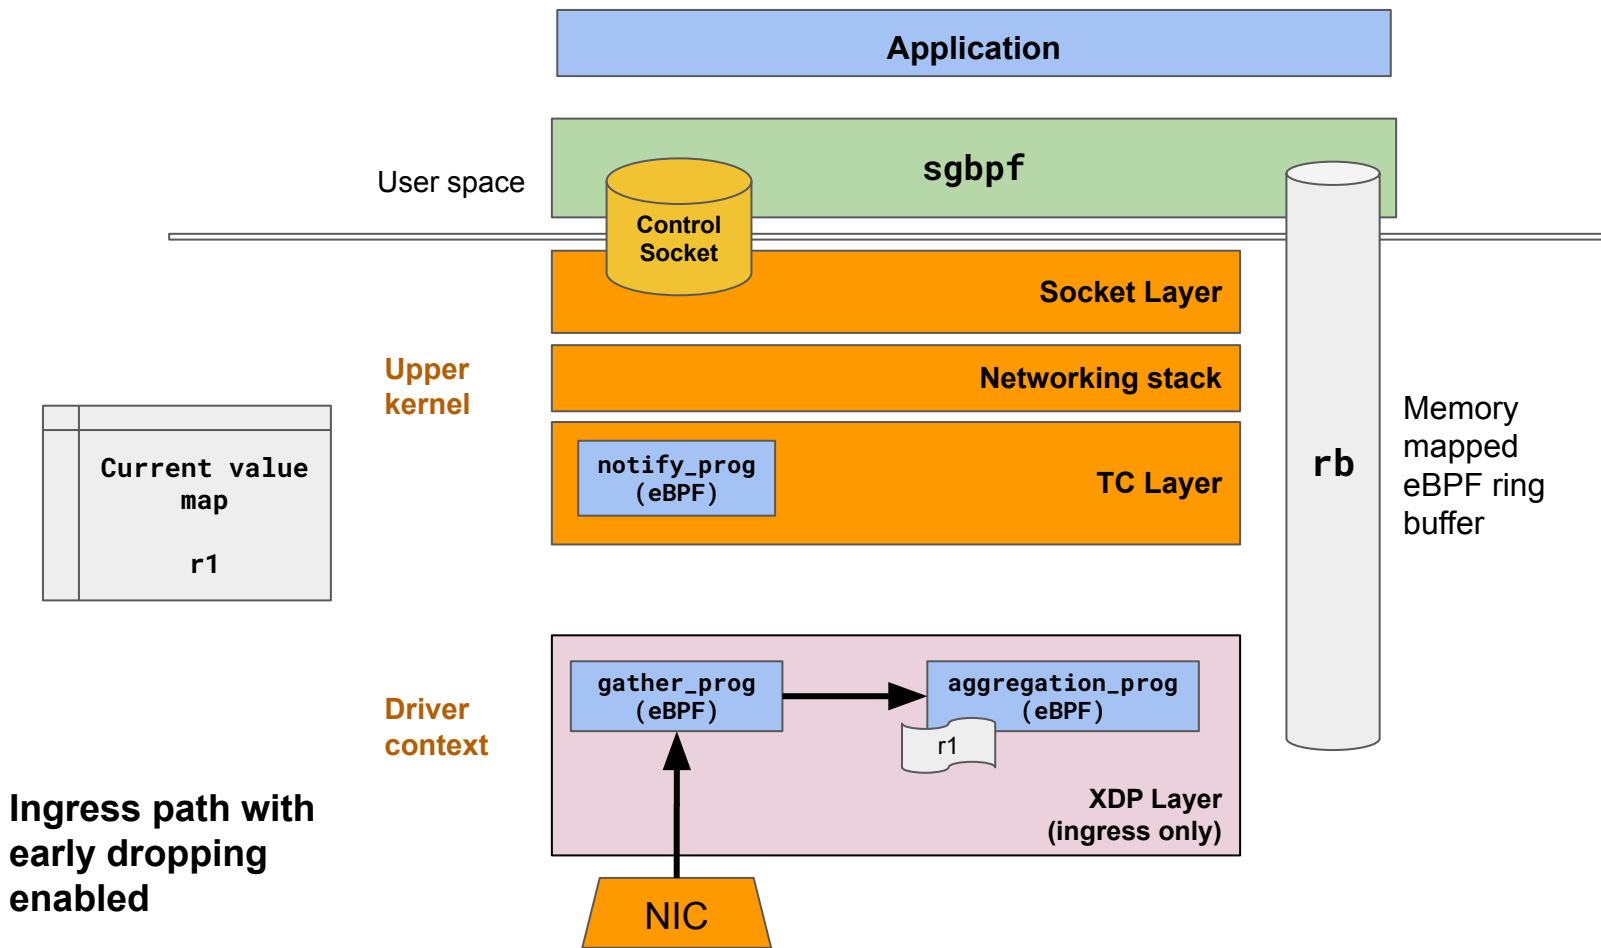


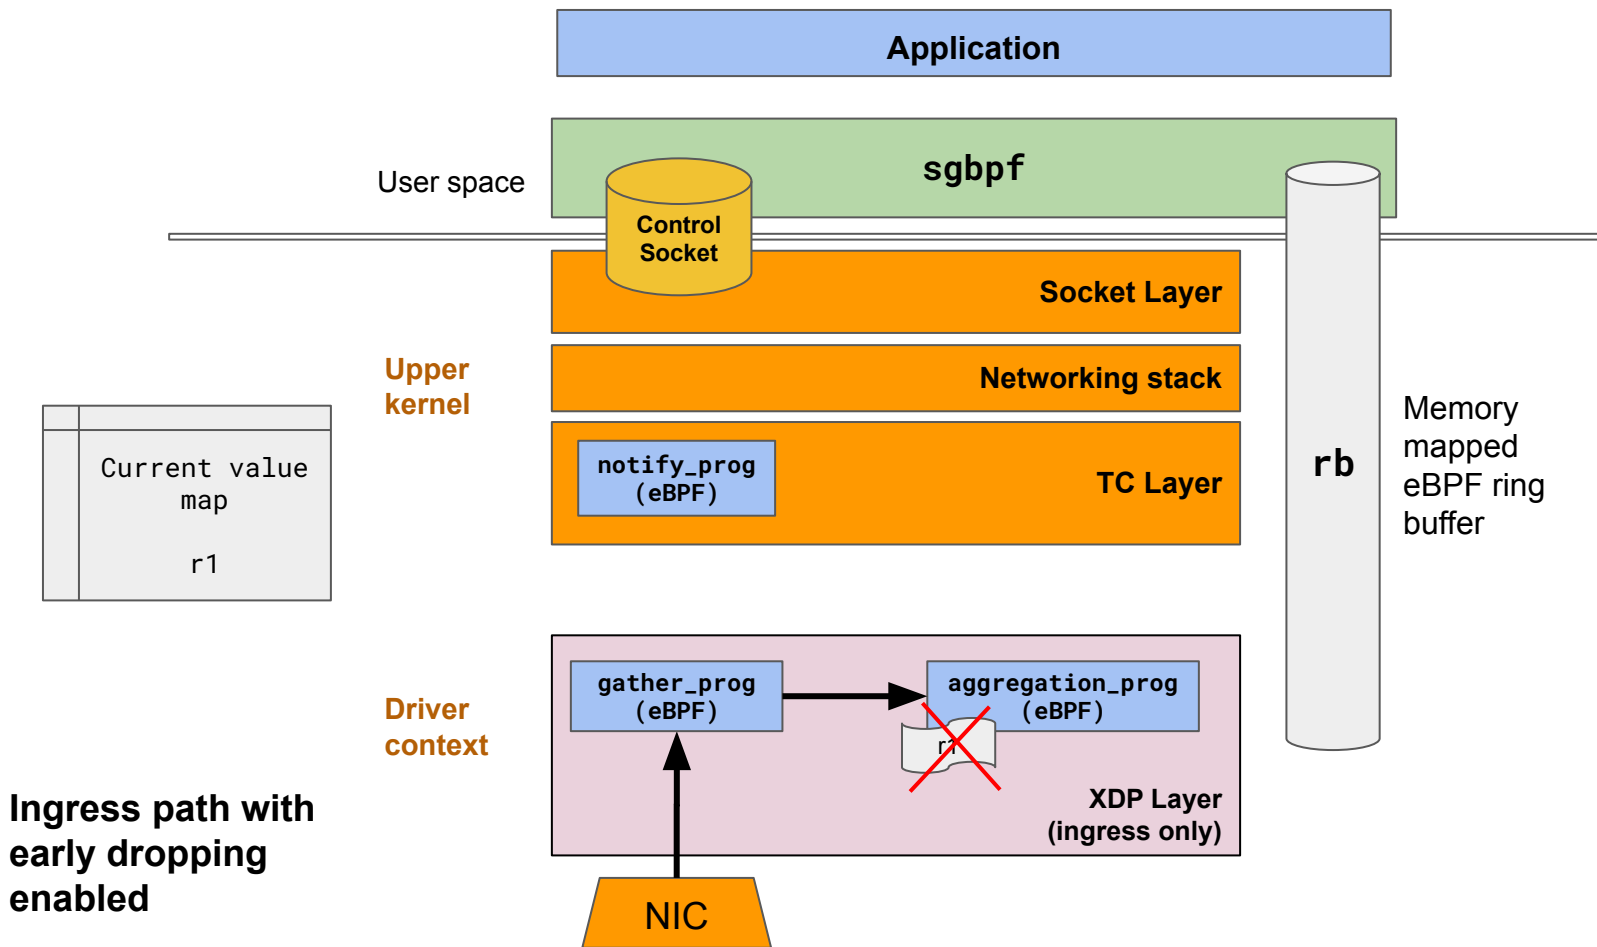


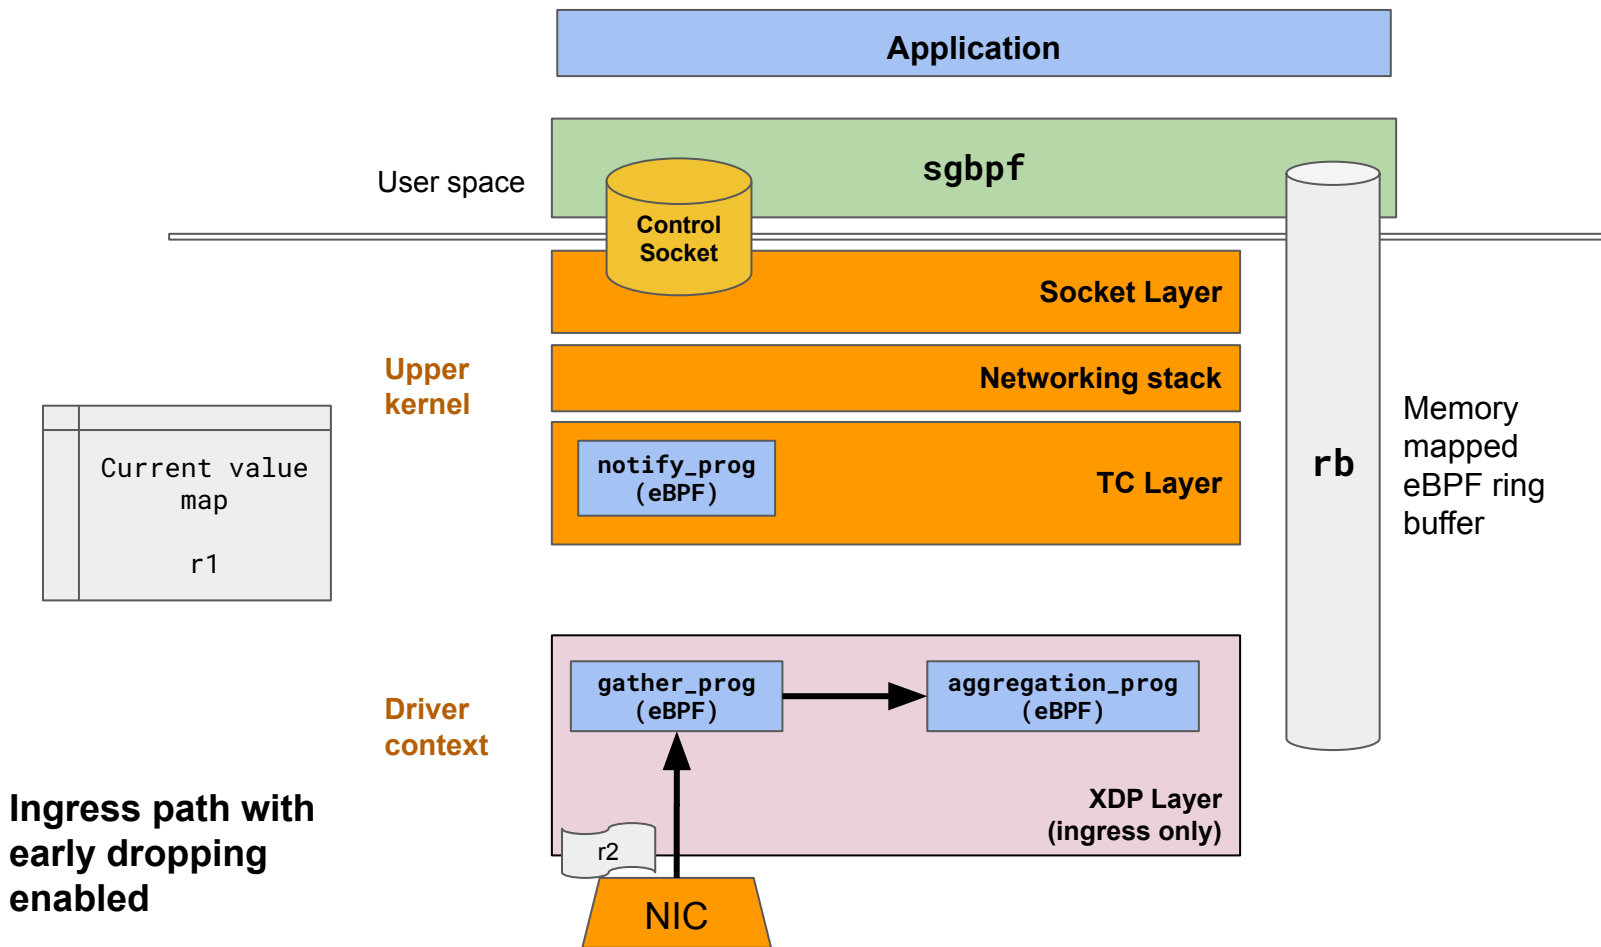


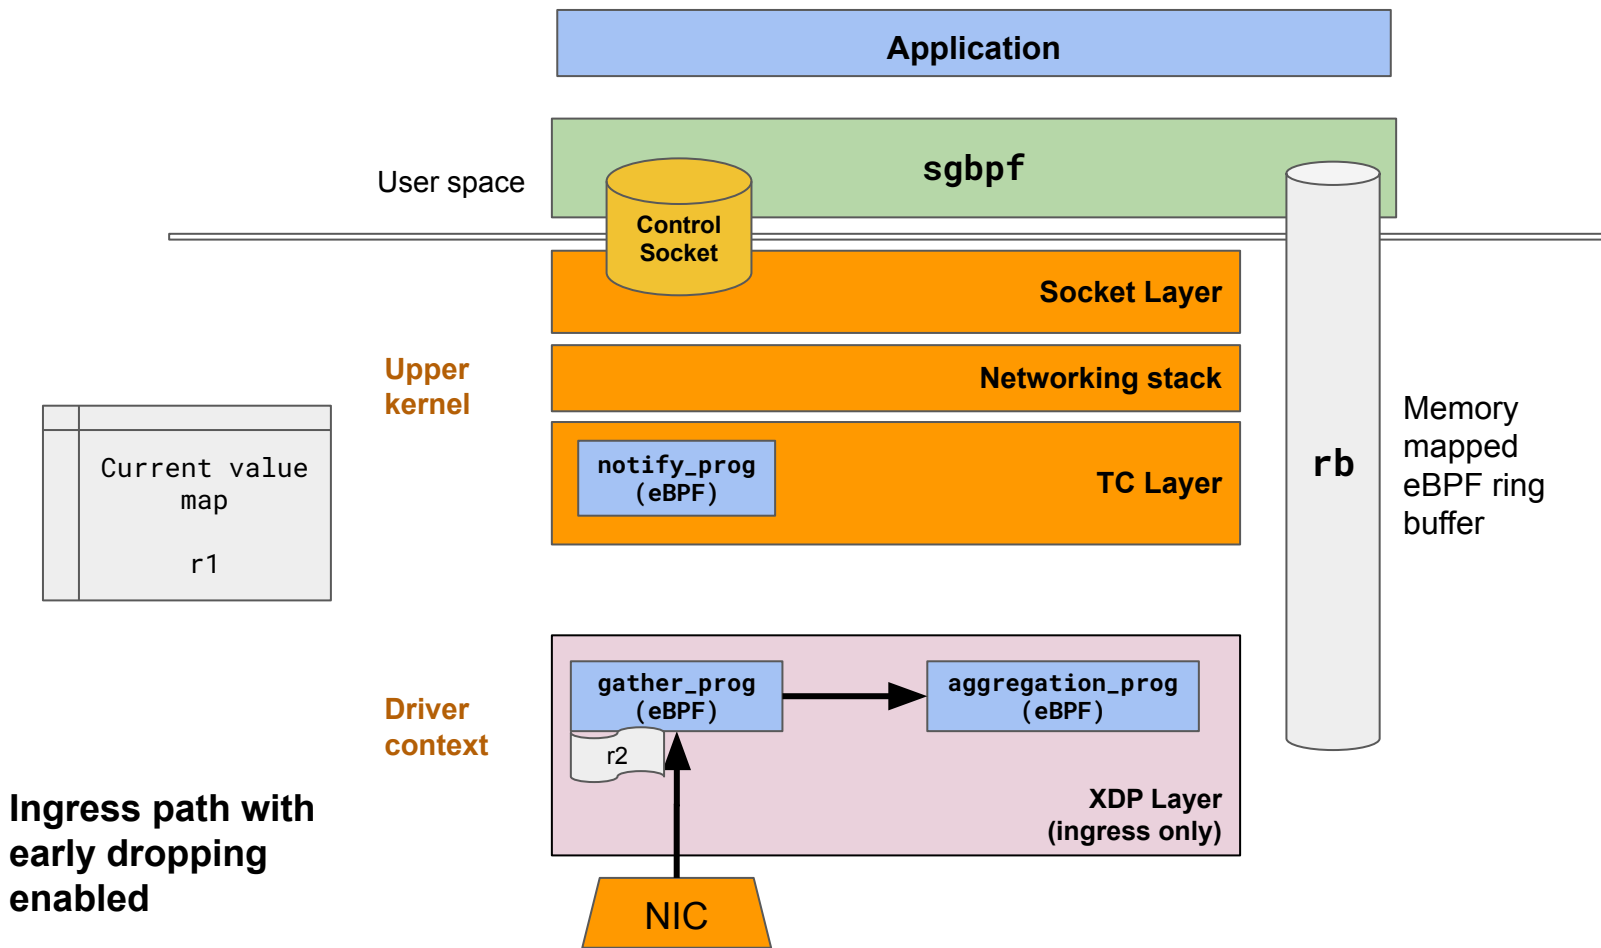


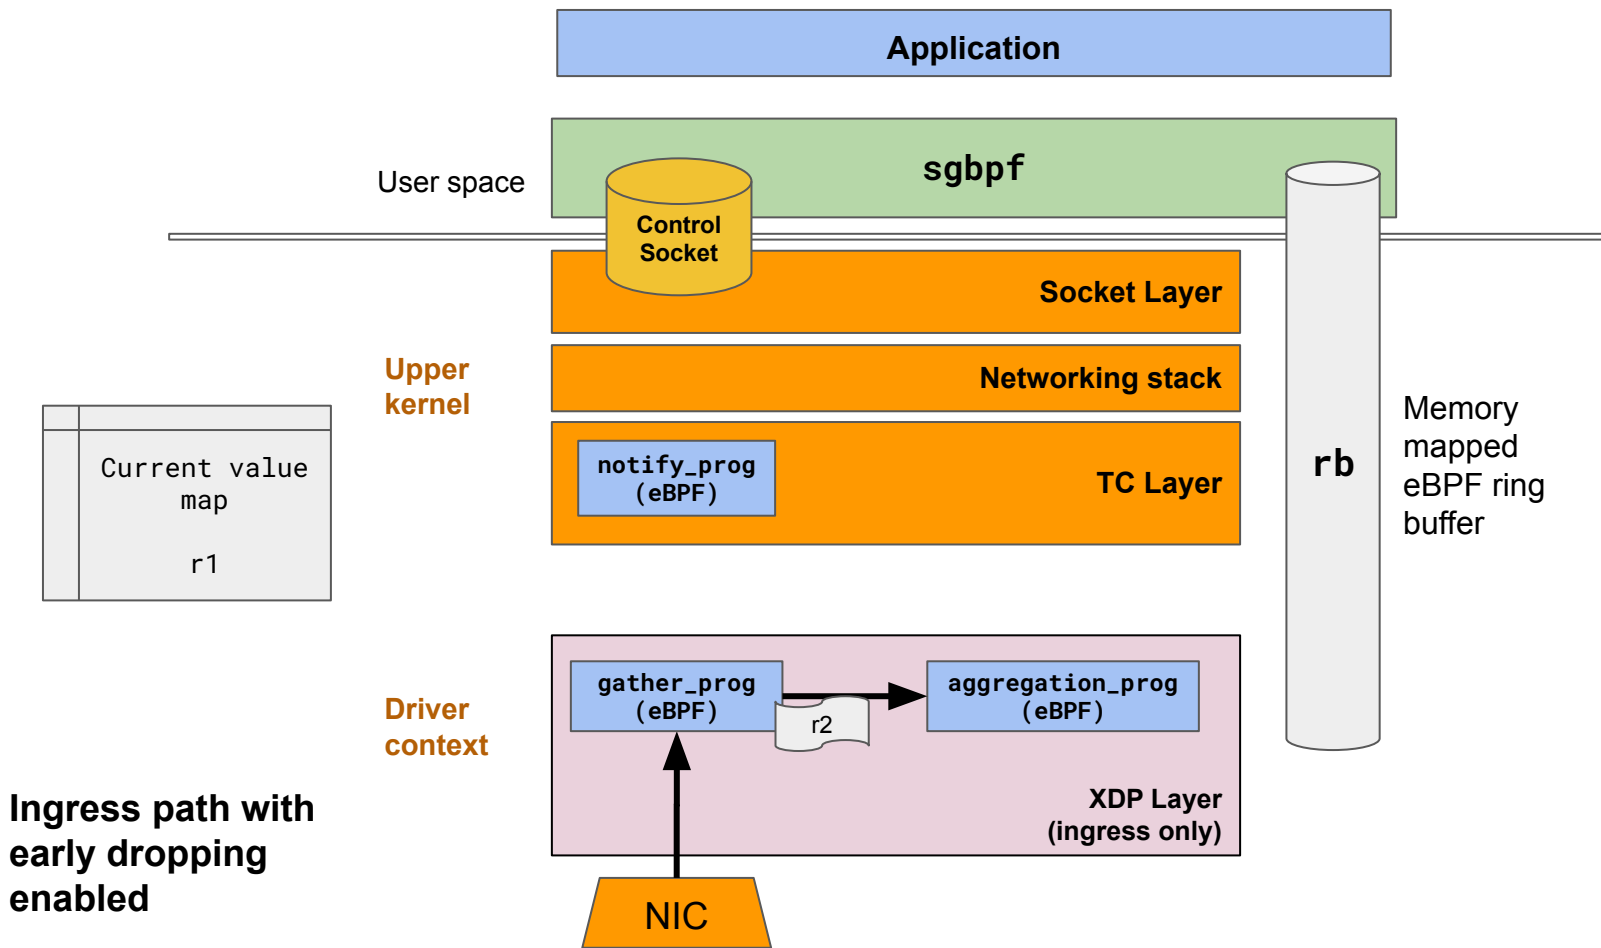


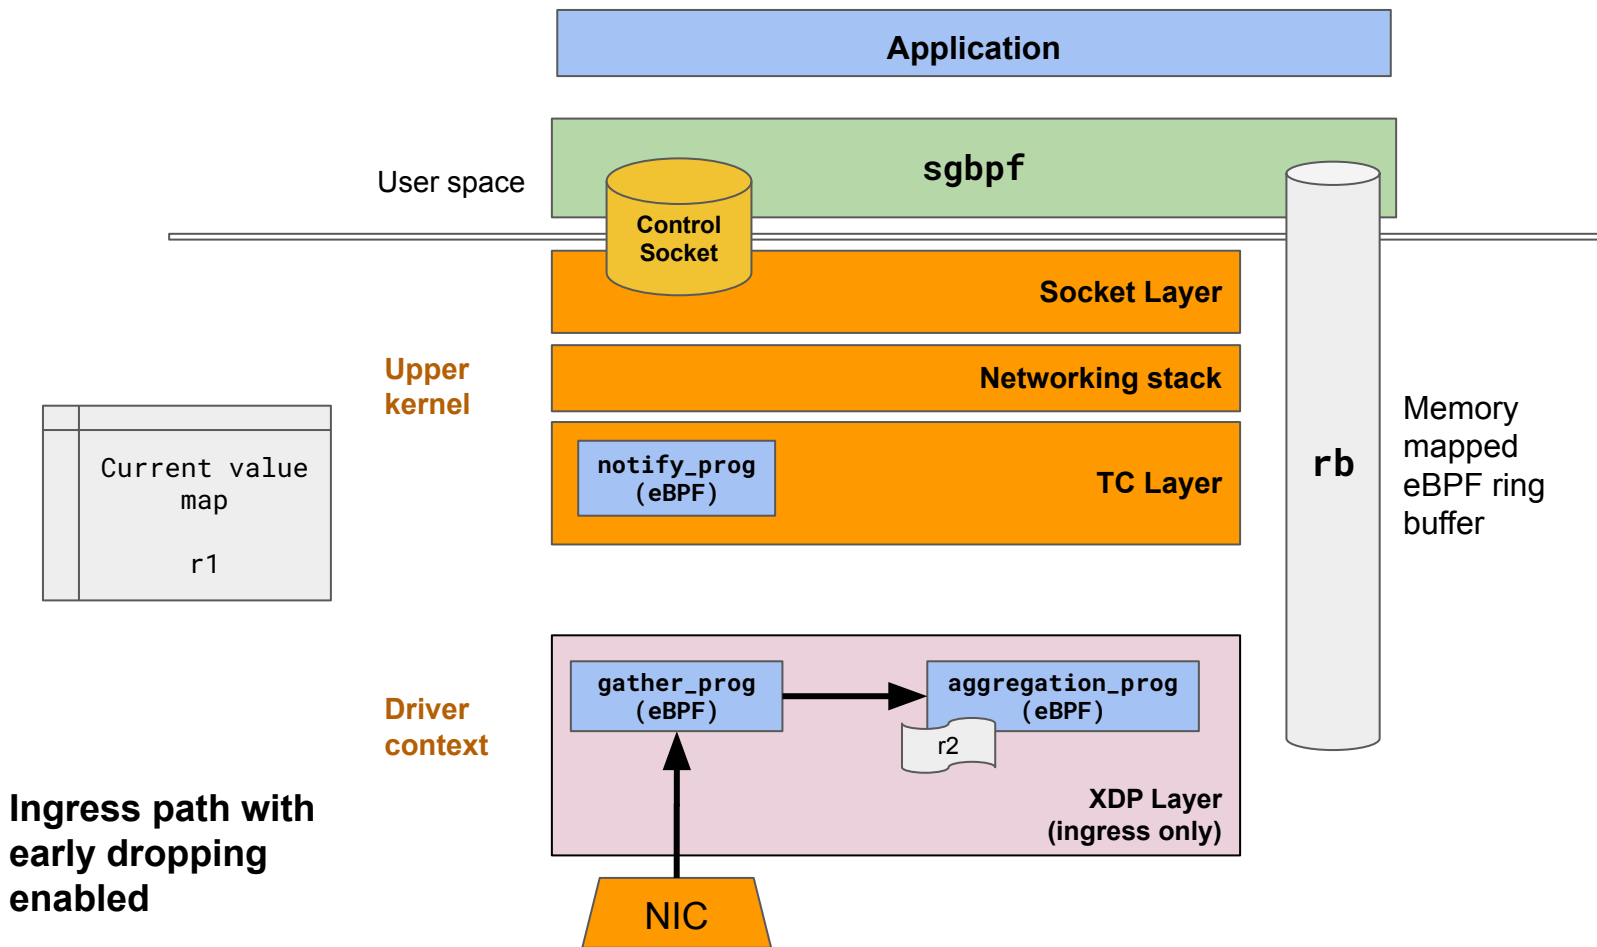




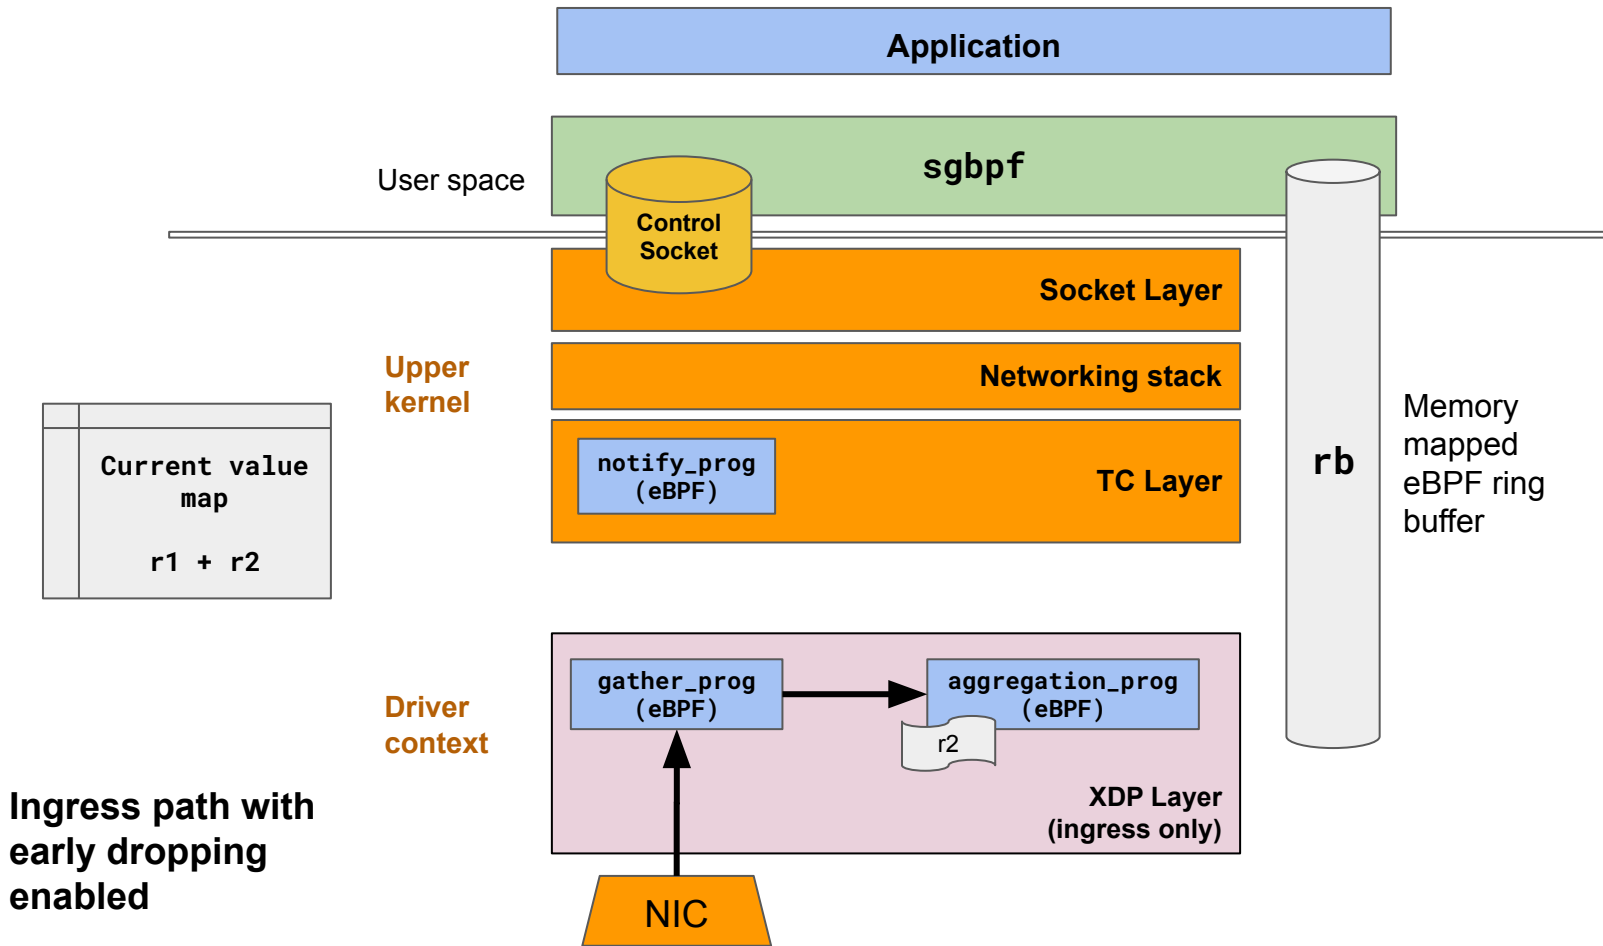


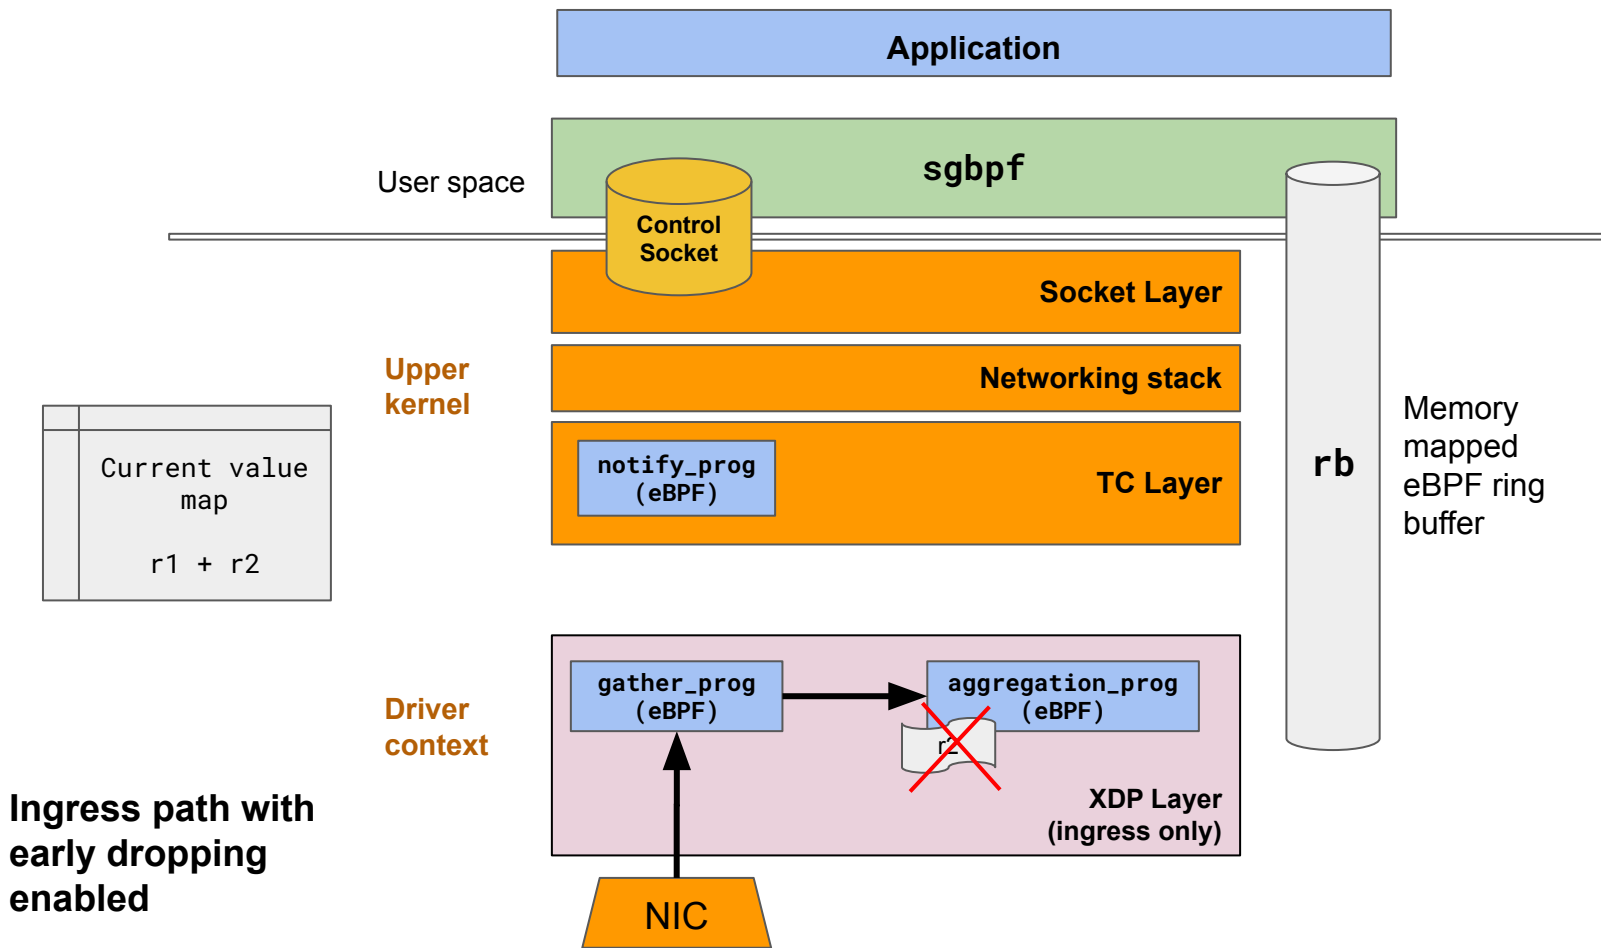


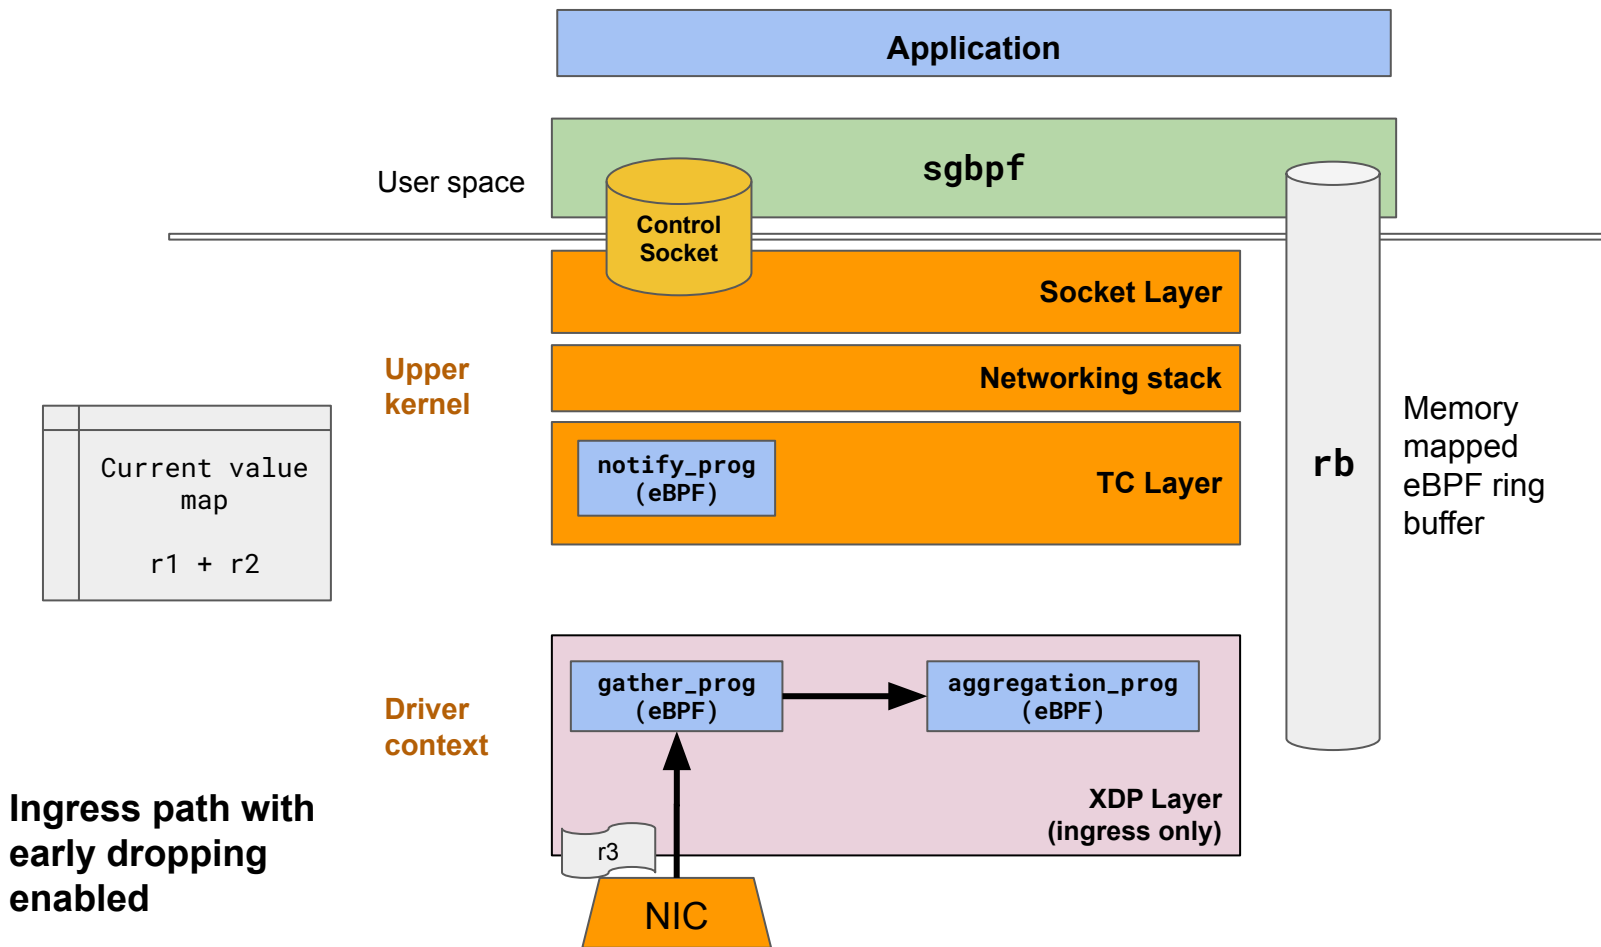


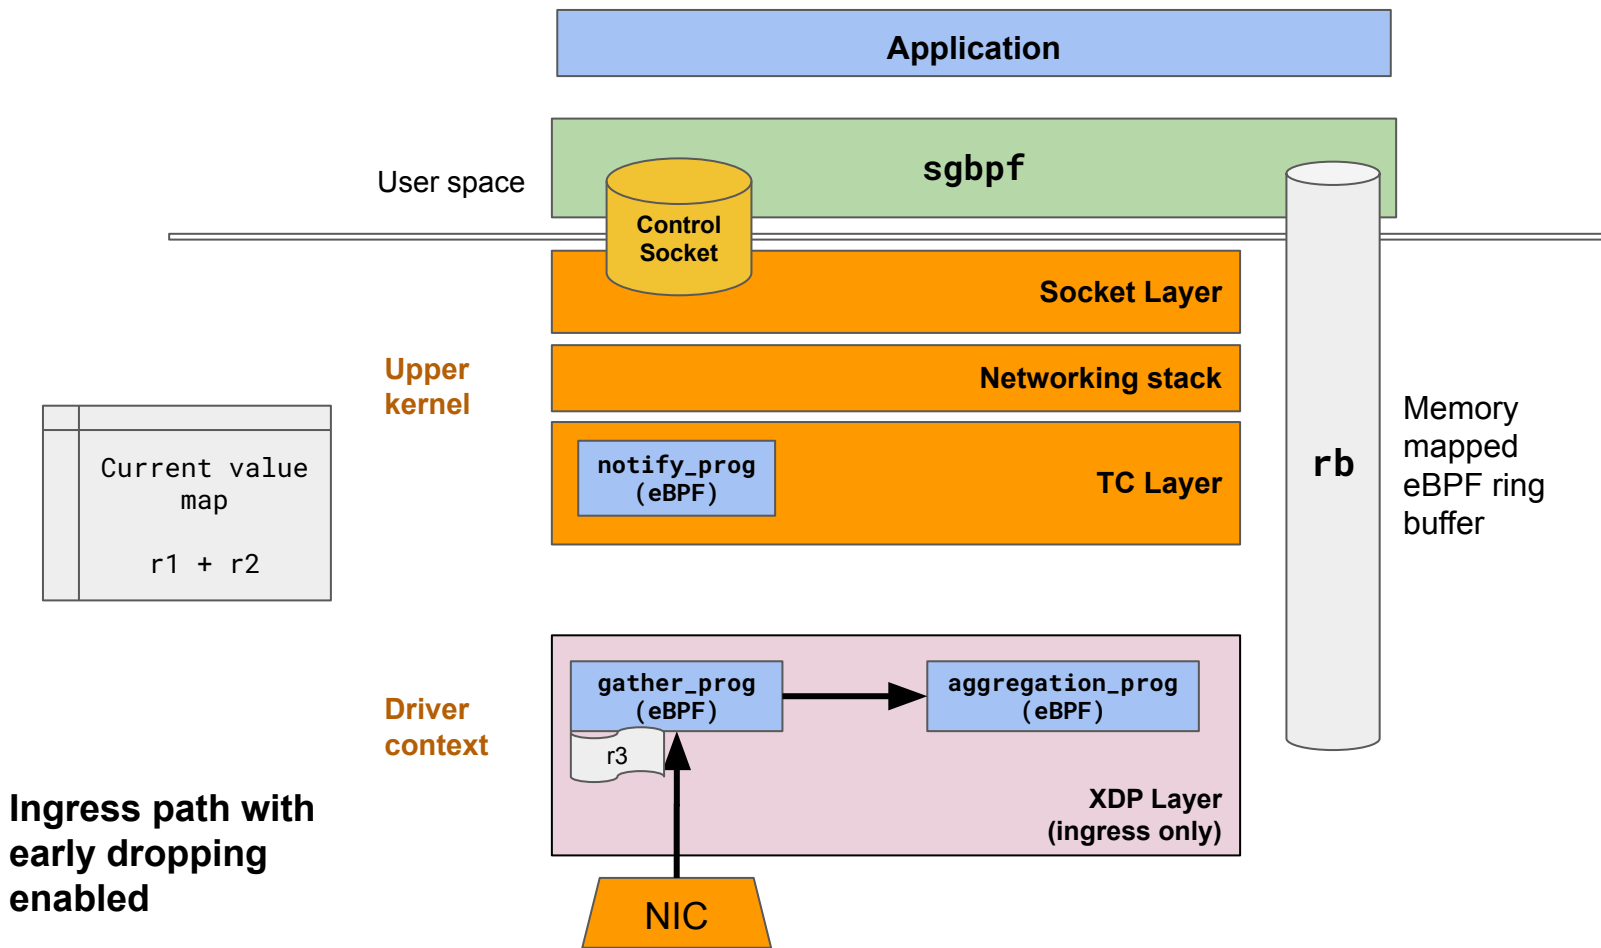


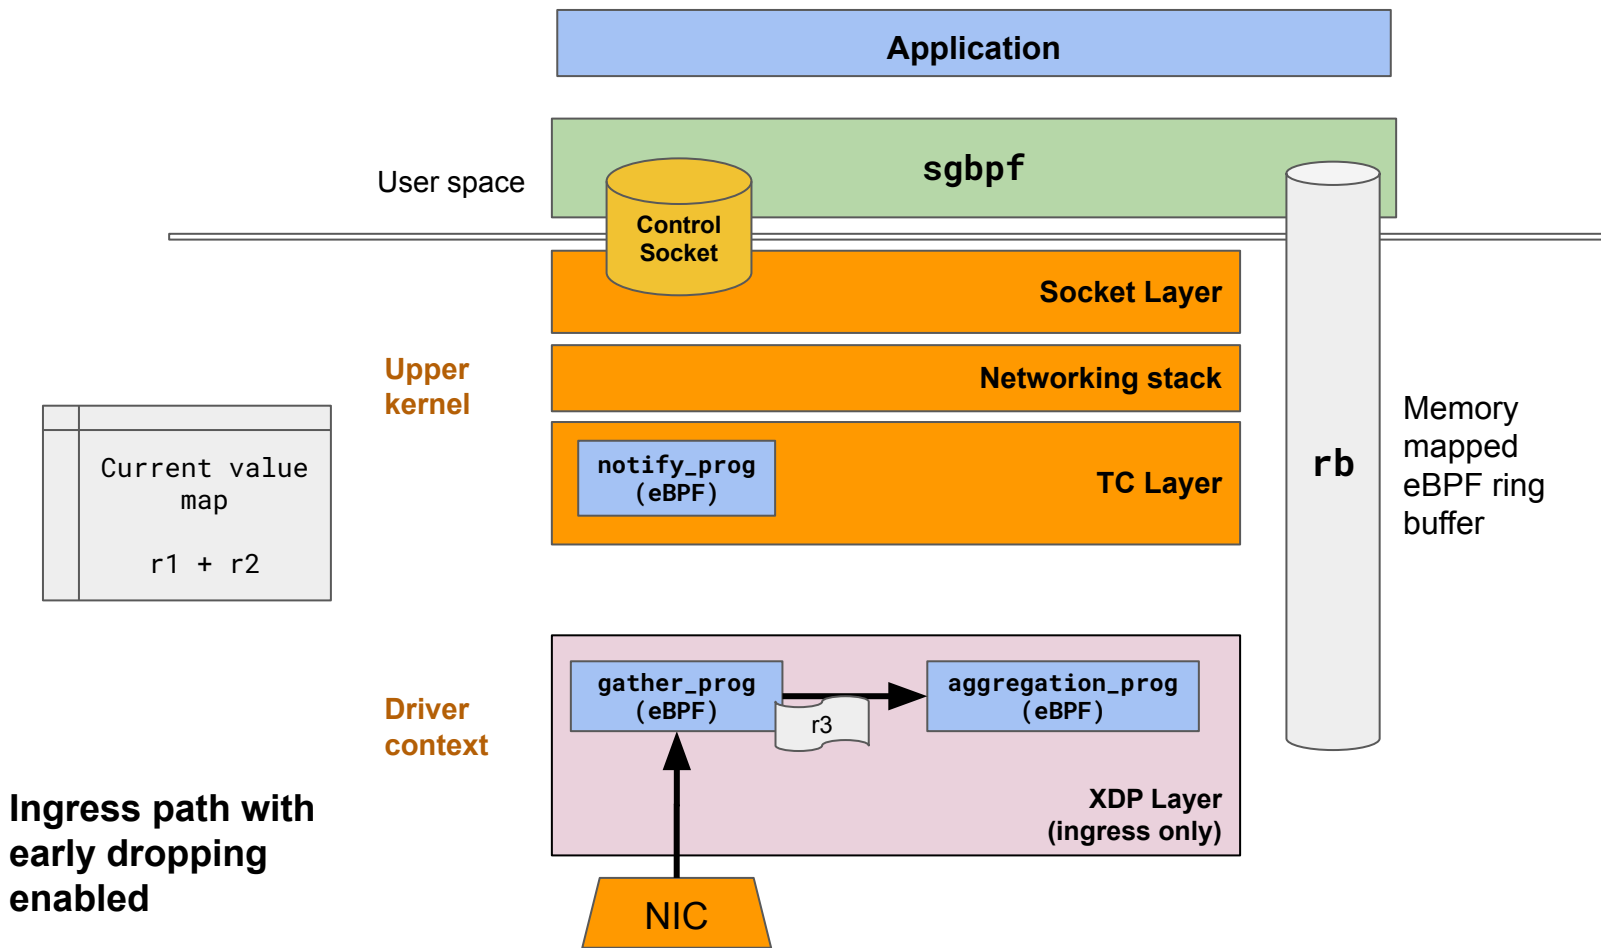


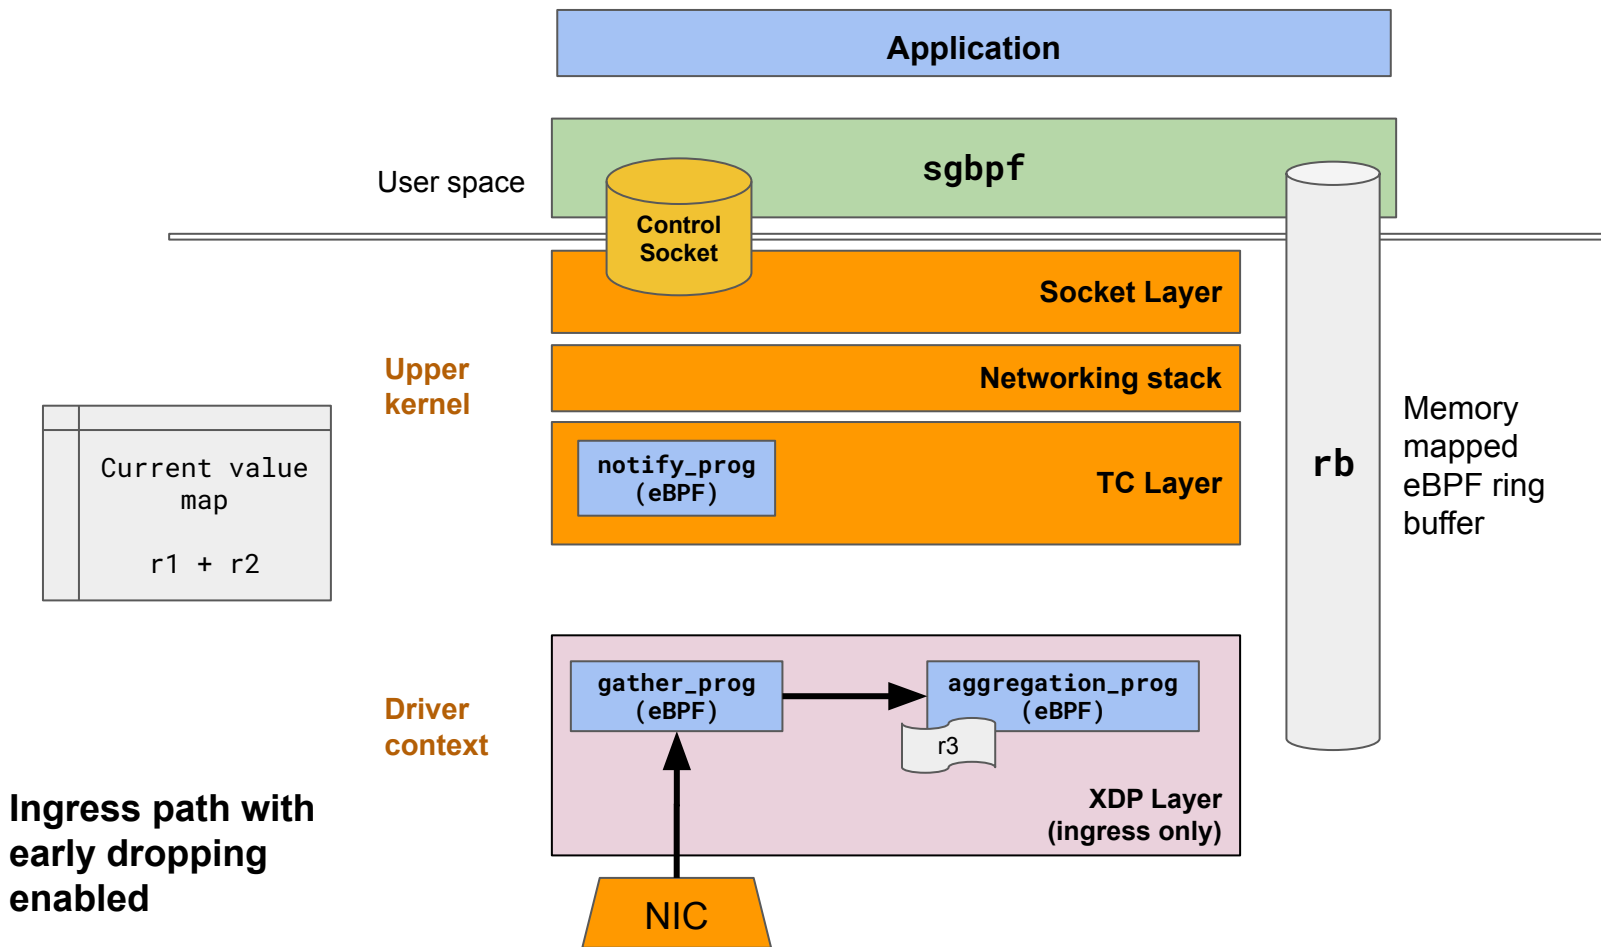


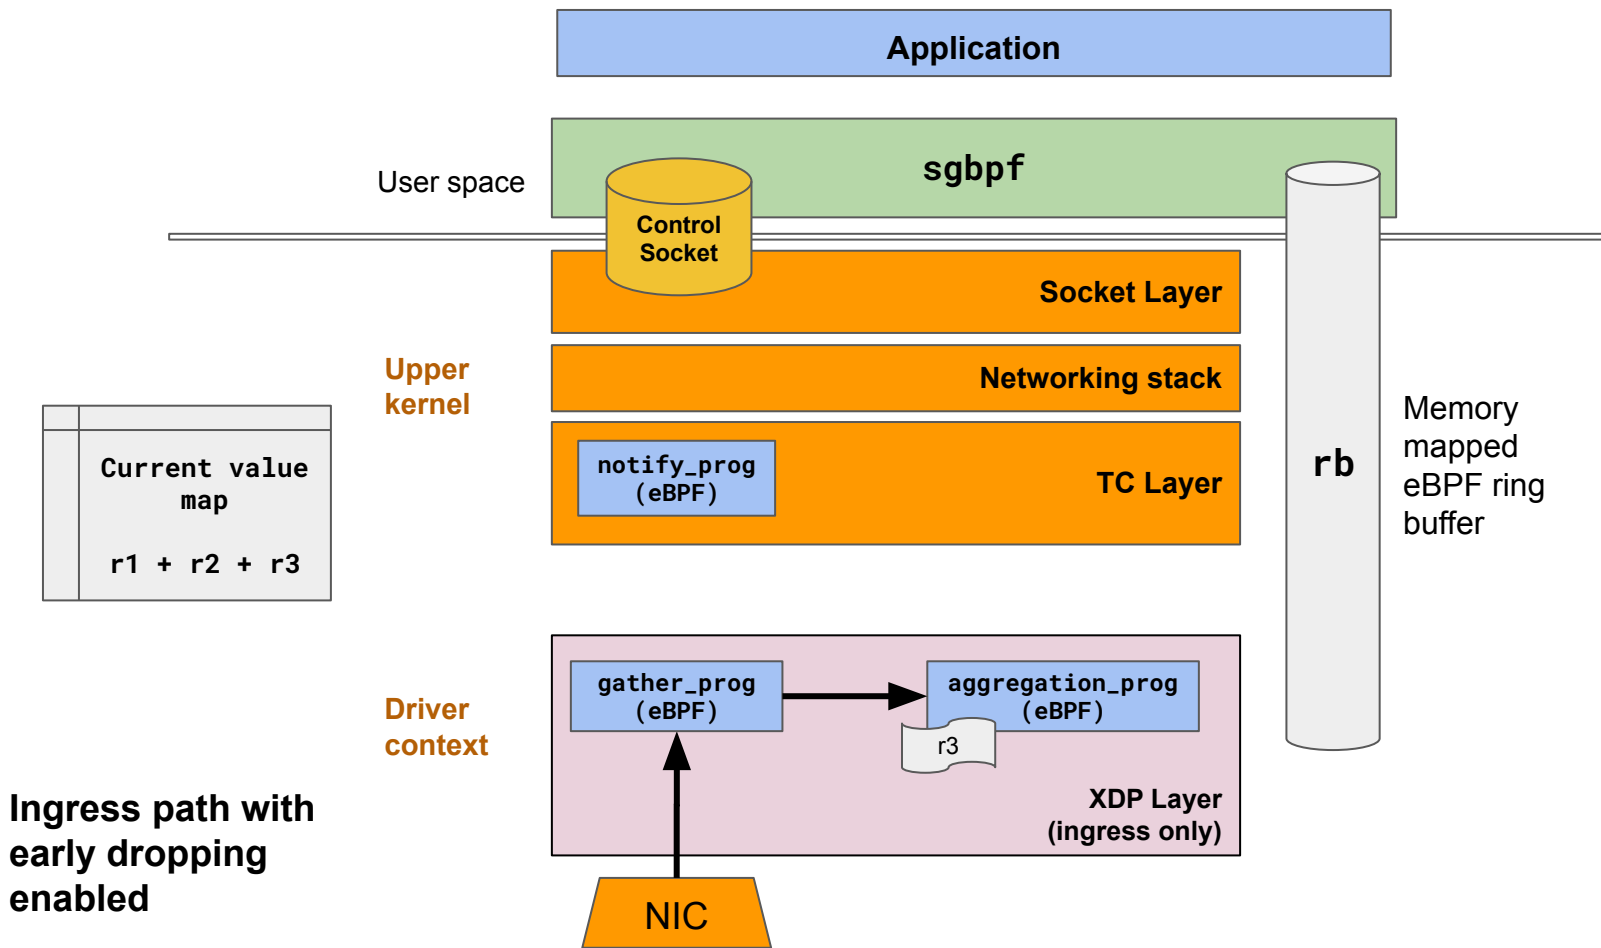


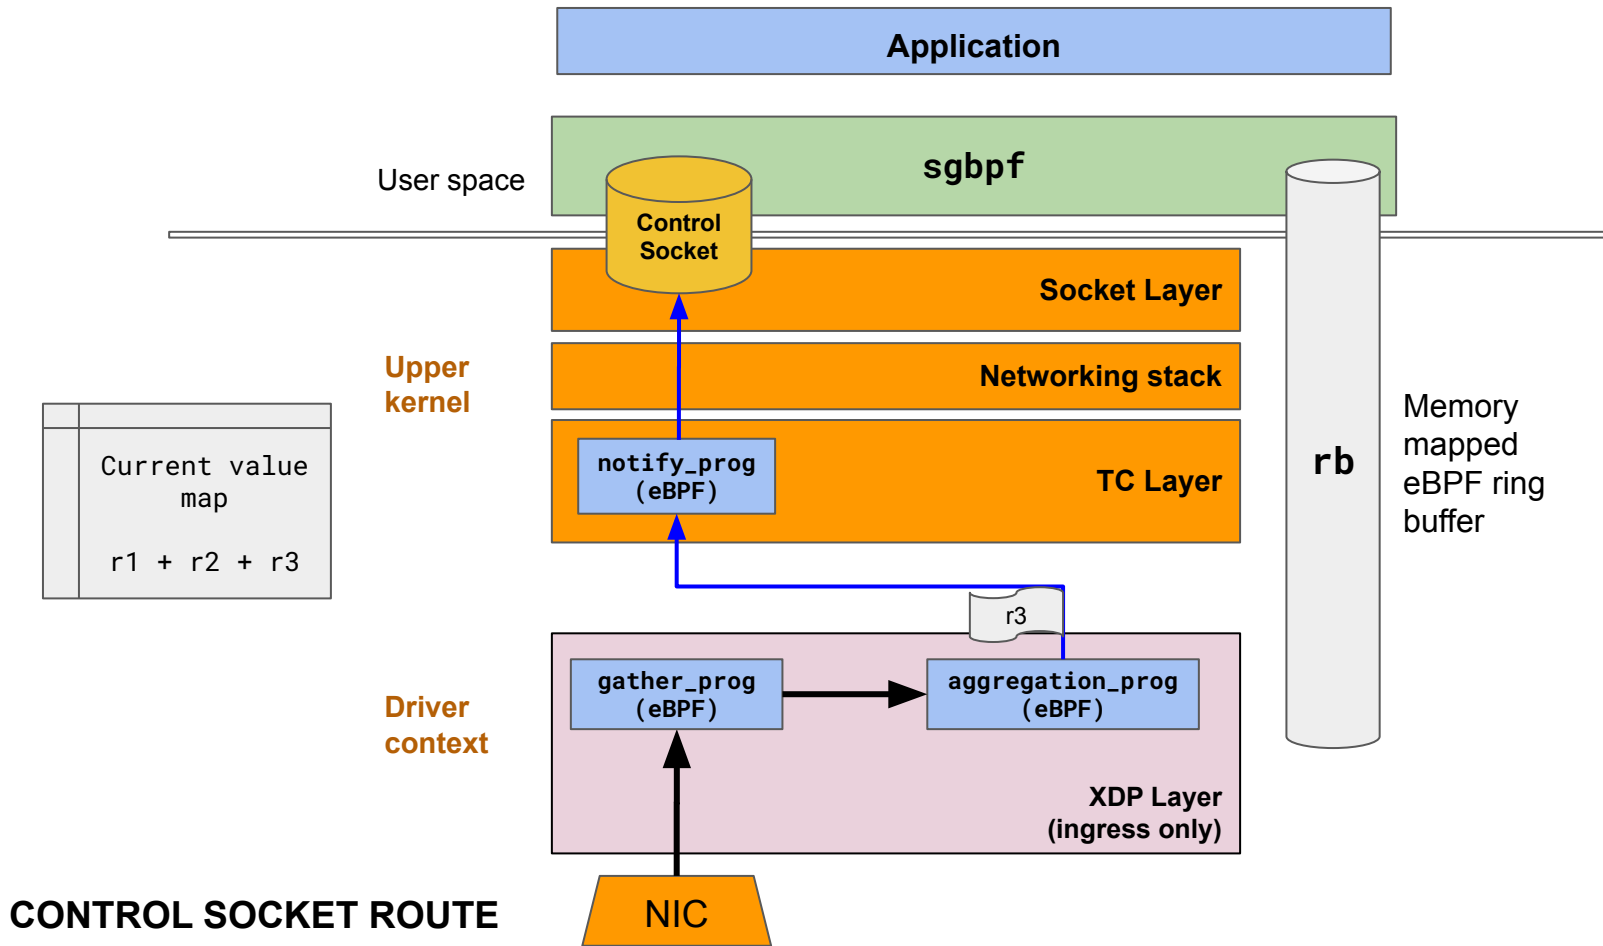




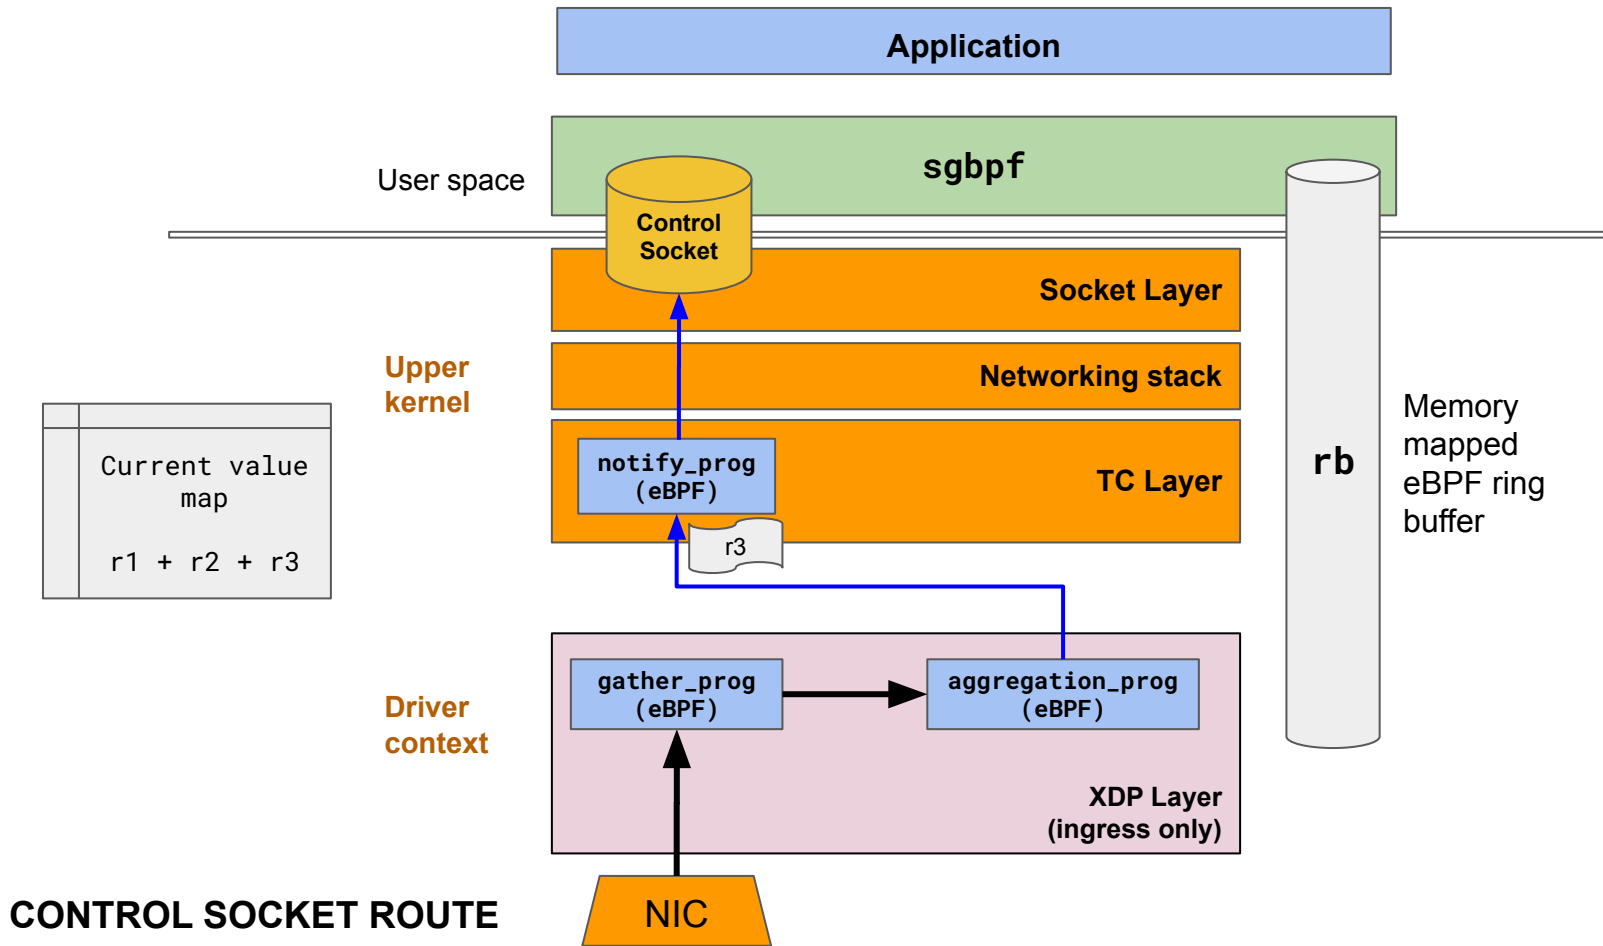


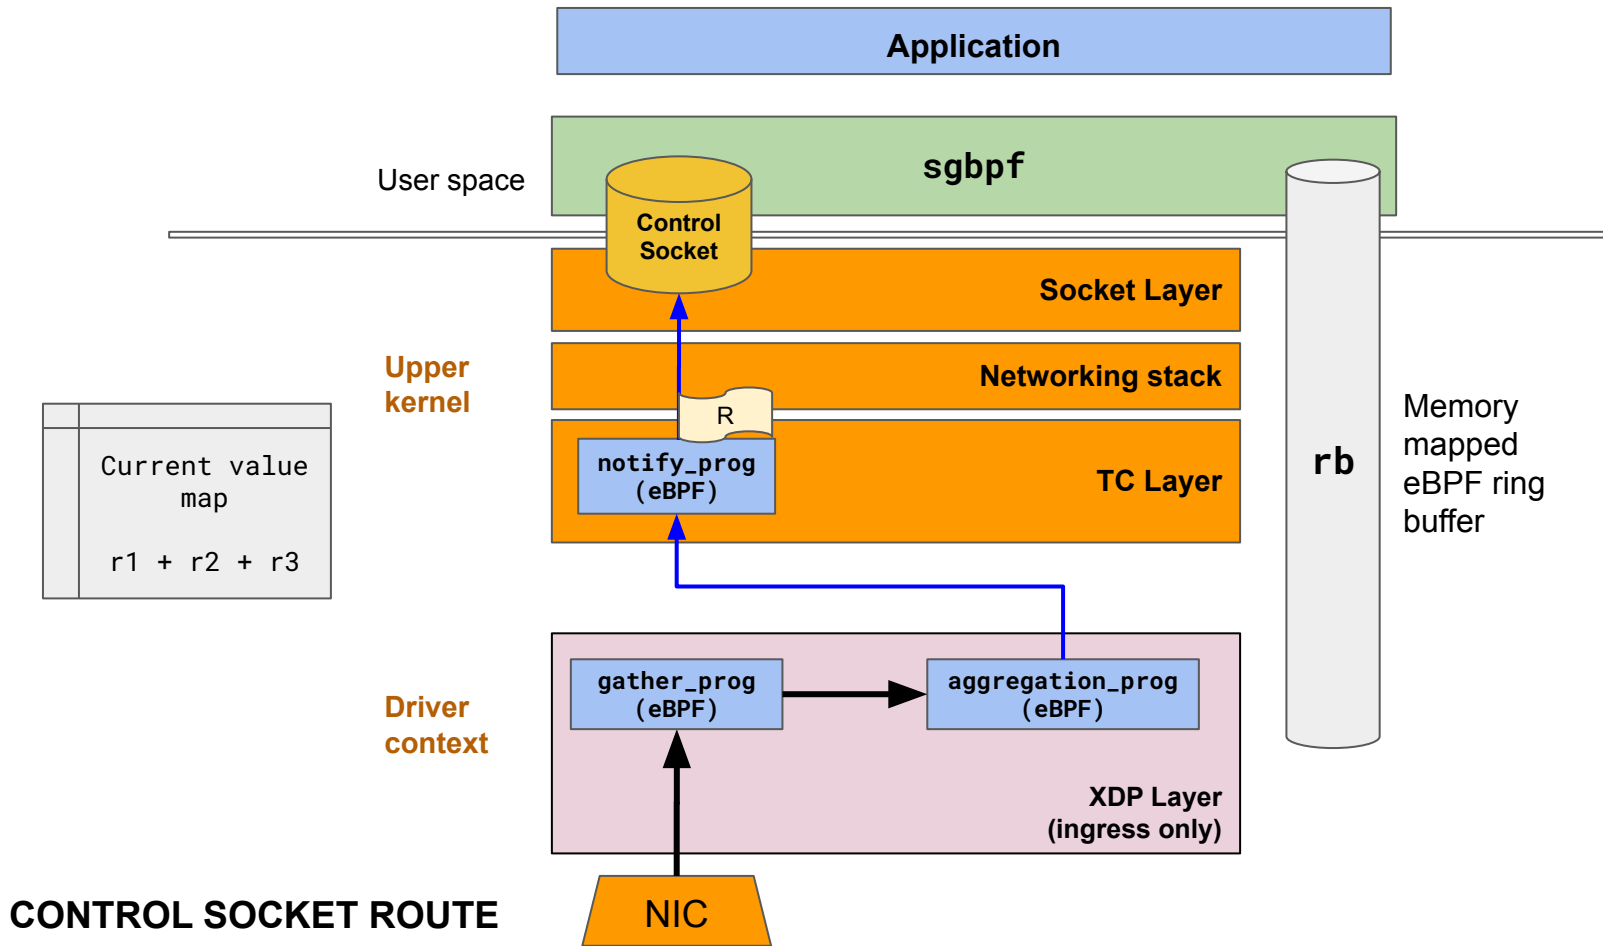


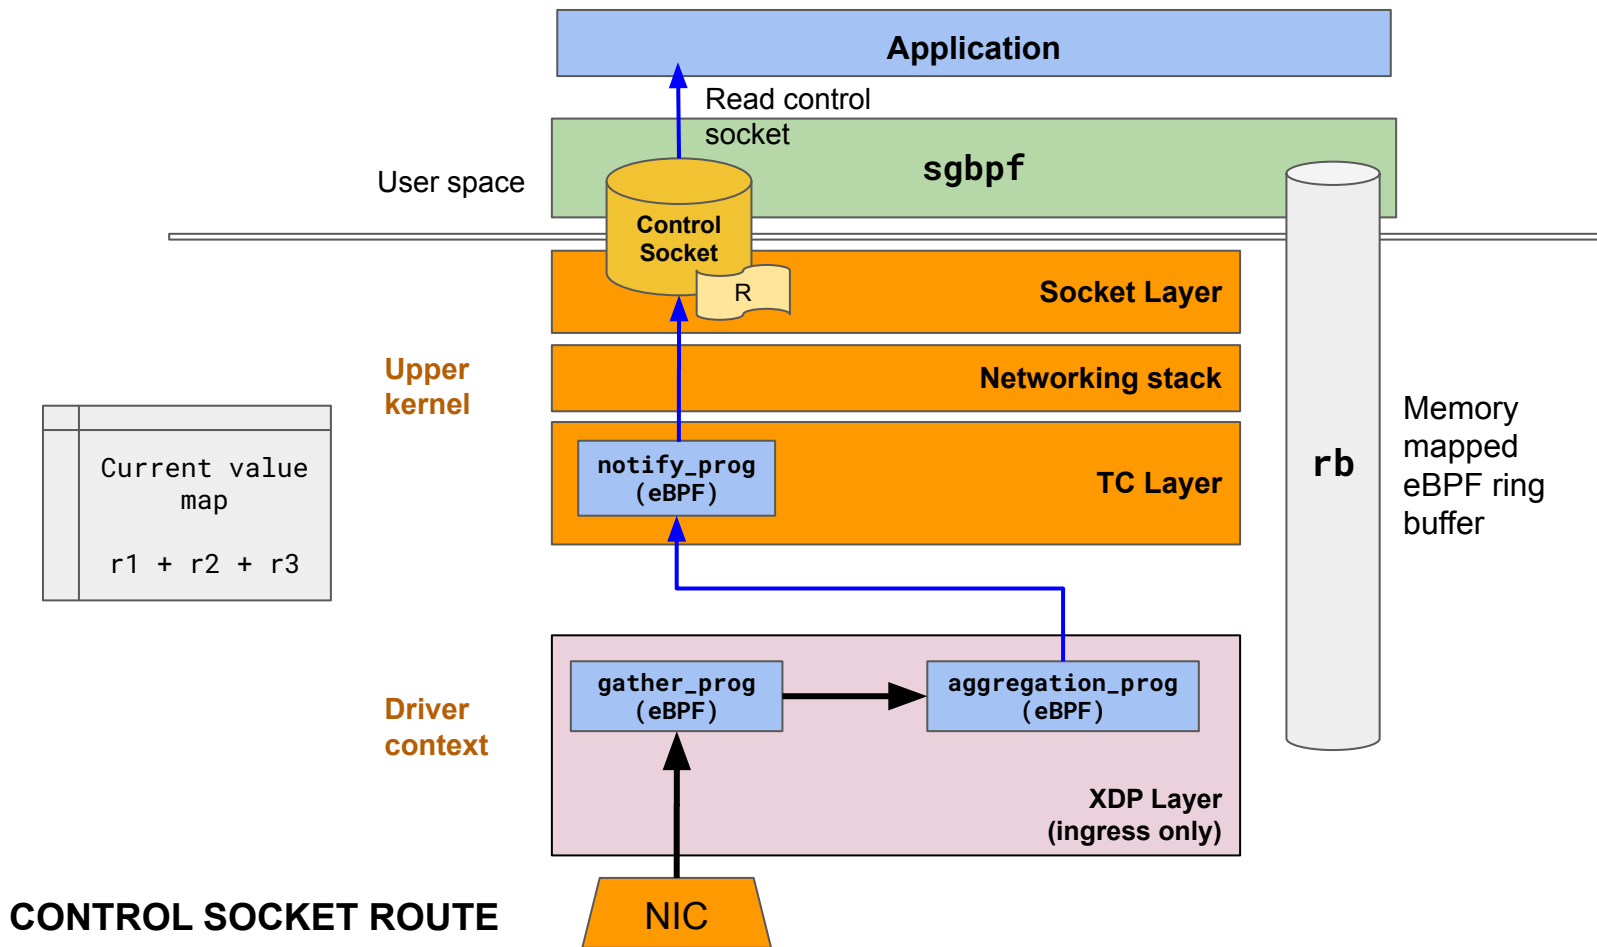


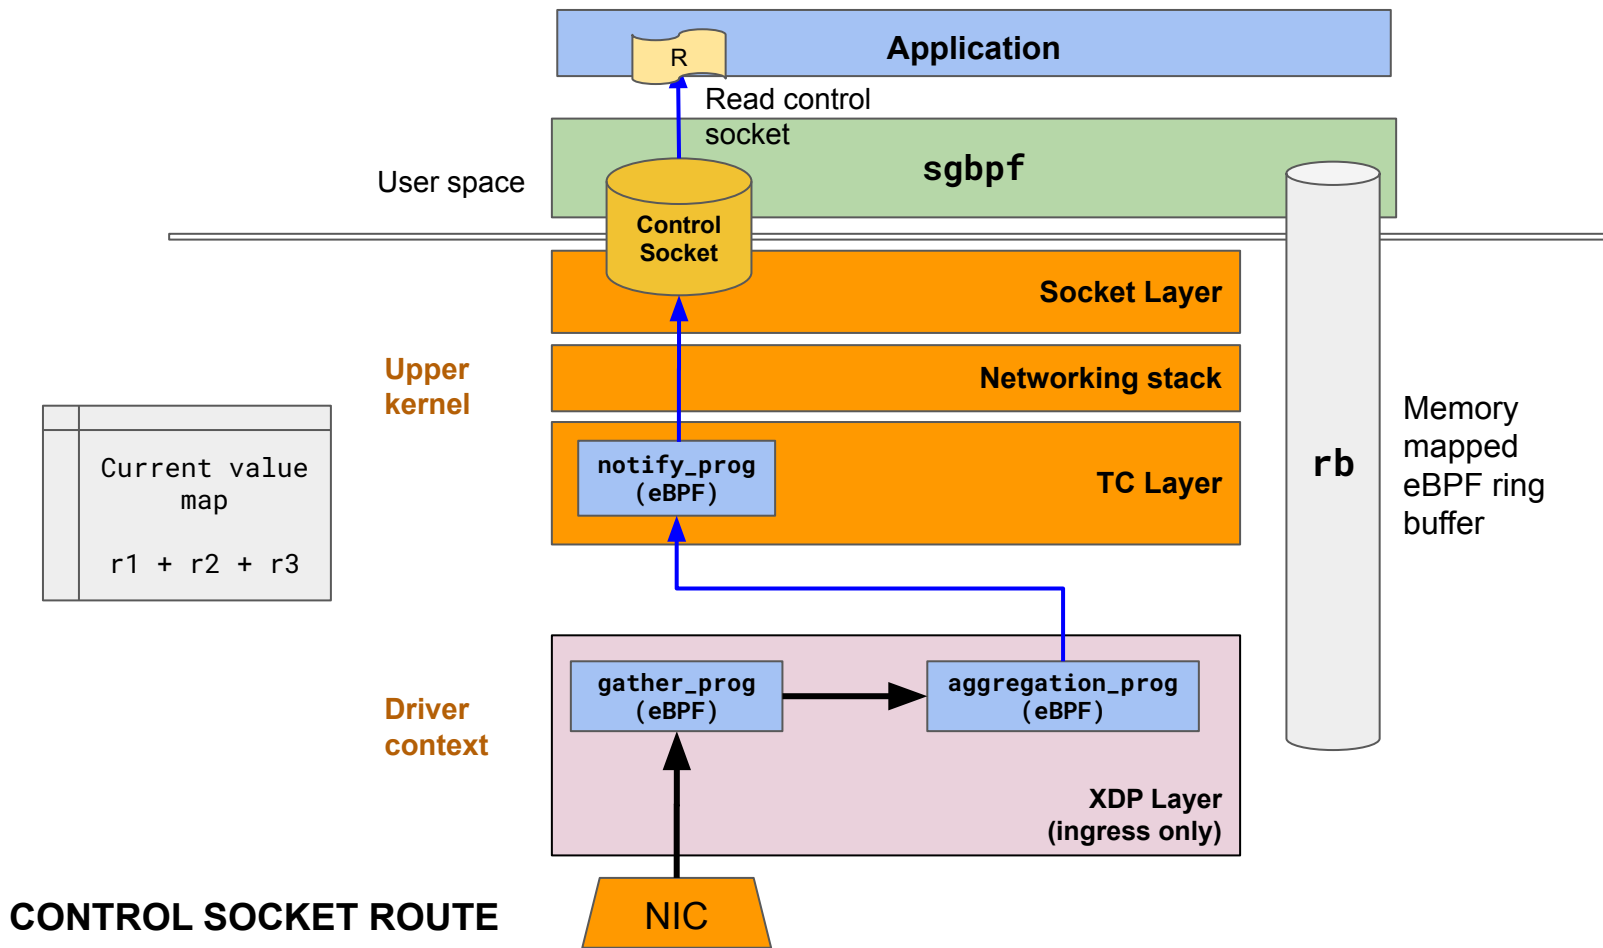


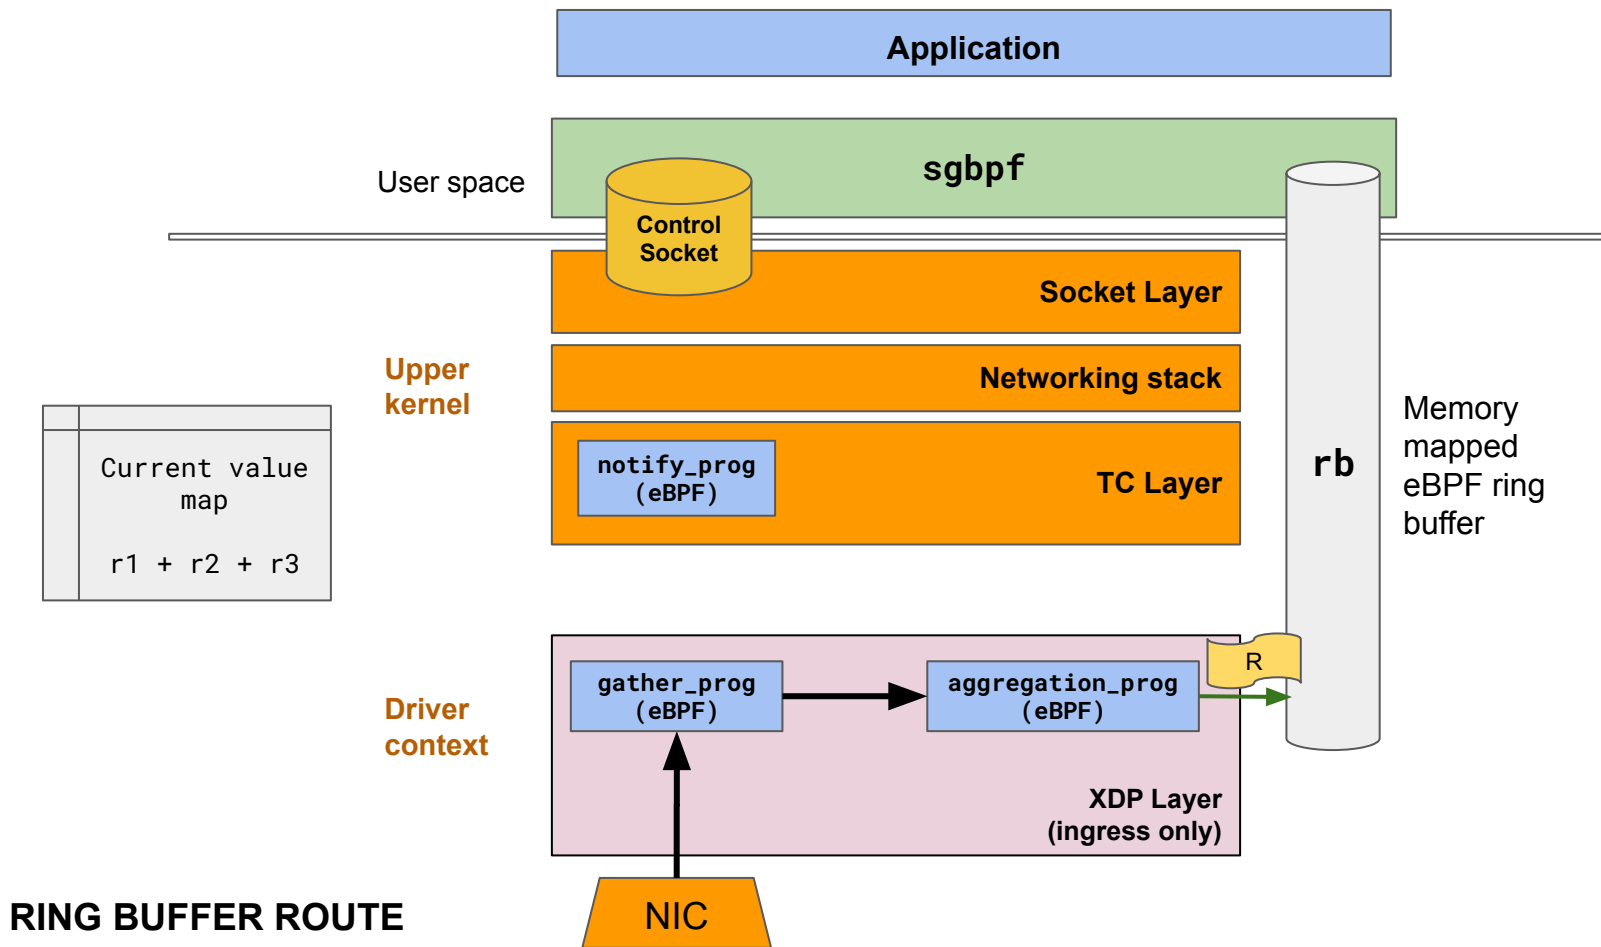


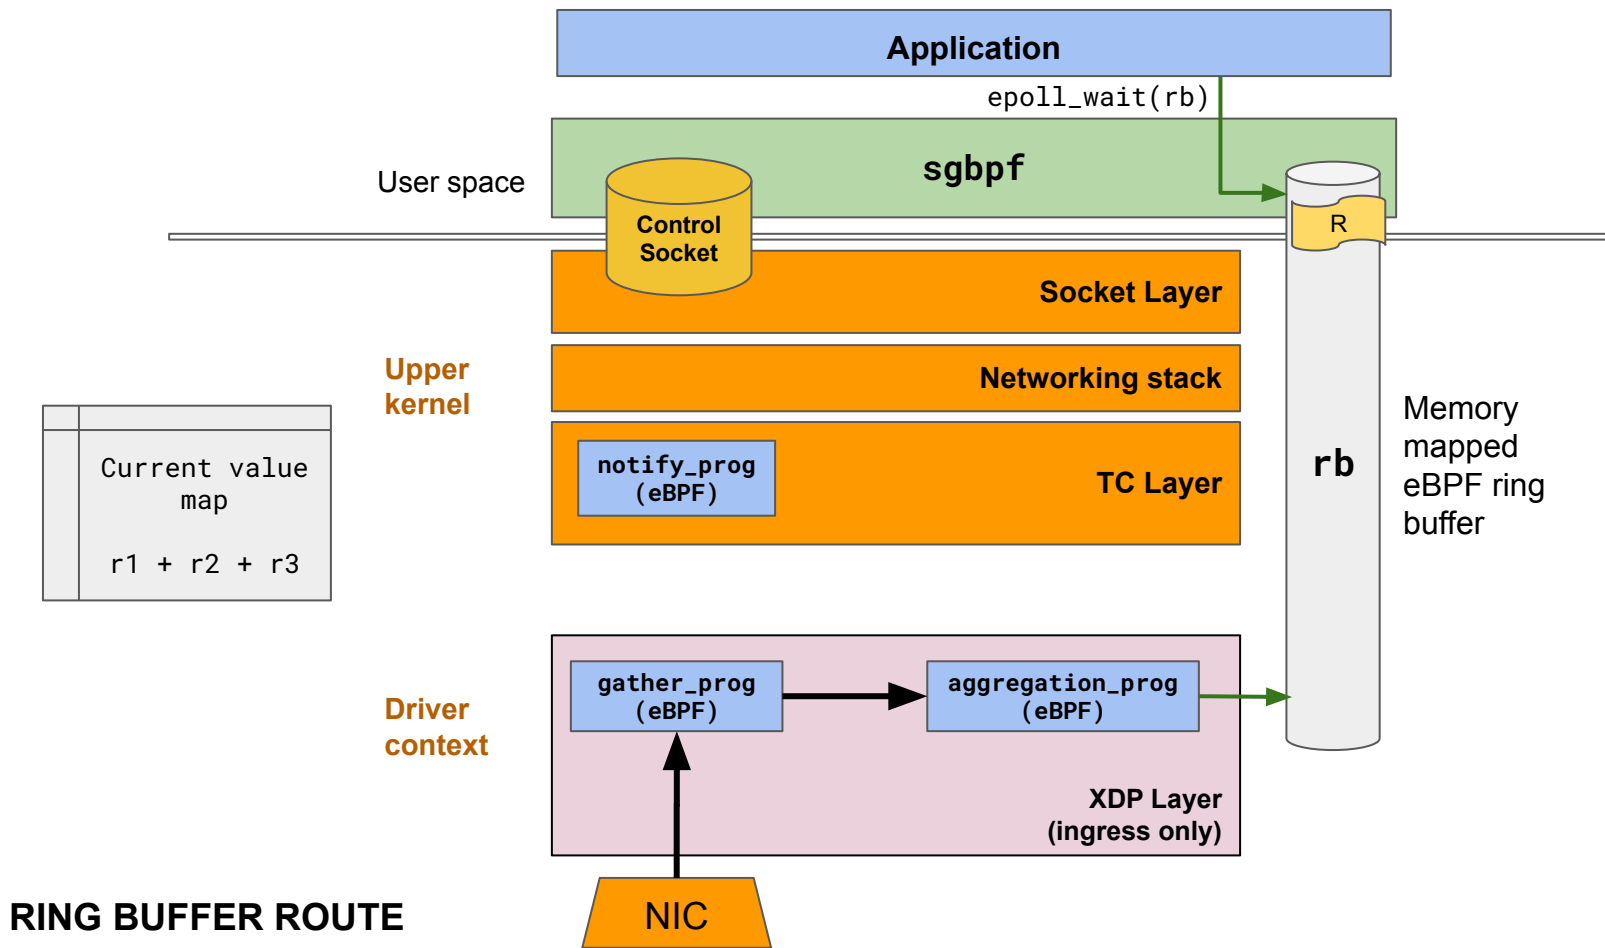


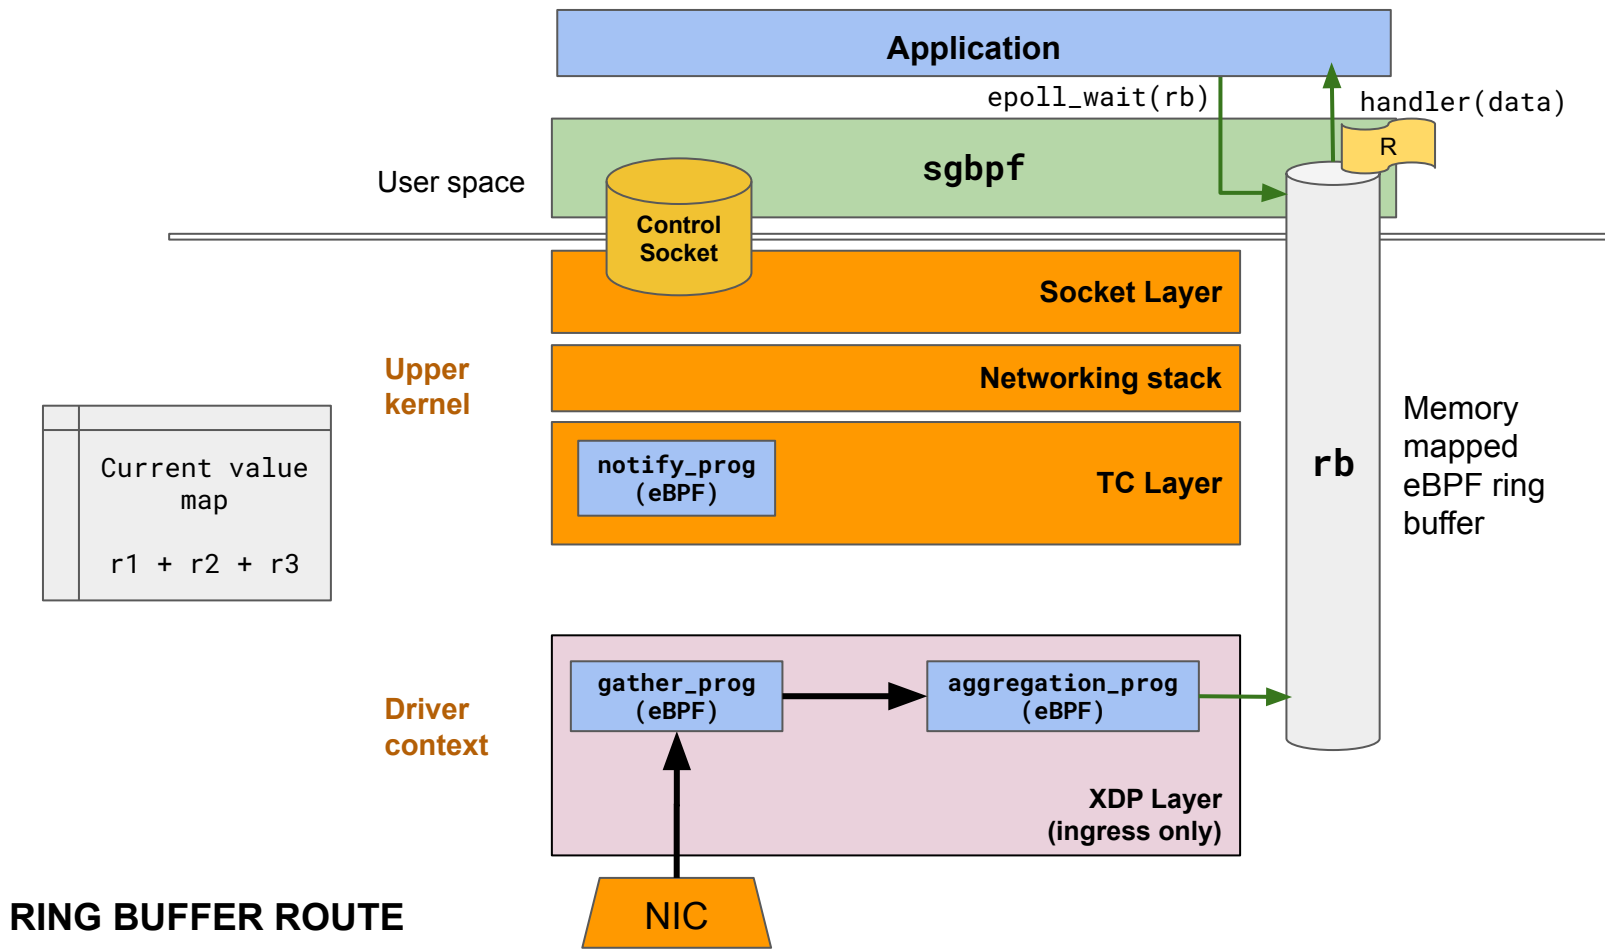


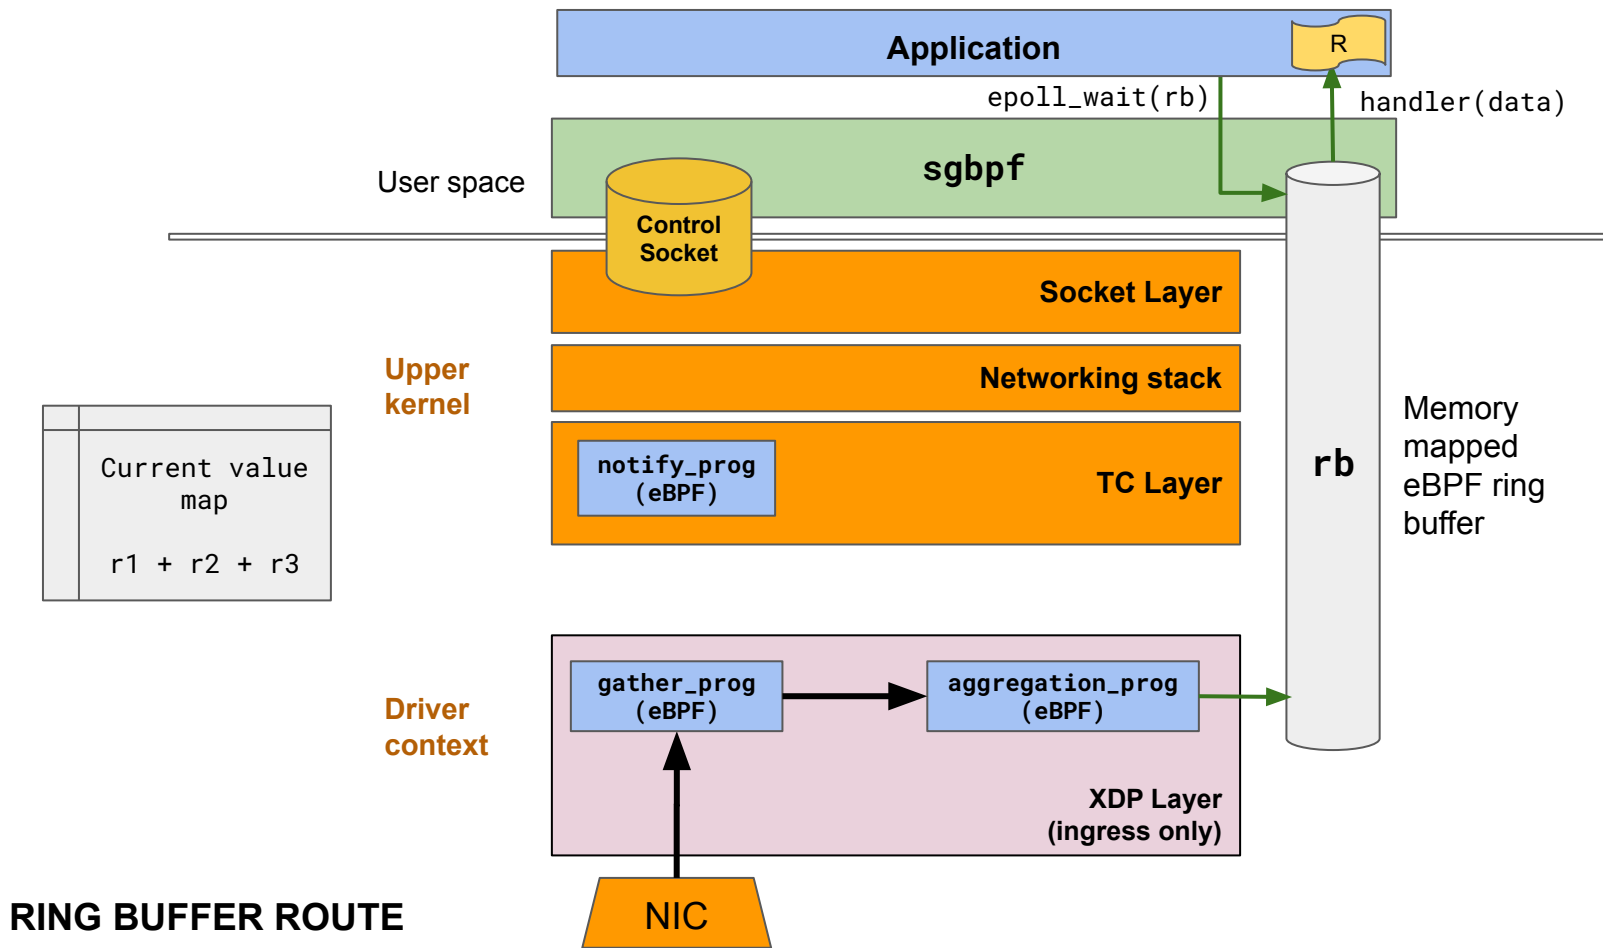












**RING BUFFER ROUTE**



## Example - custom aggregation program

```
#include "bpf_h/helpers.bpf.h"

SEC("xdp")
int aggregation_prog(struct xdp_md* xdp_ctx) {
    struct aggregation_prog_ctx ctx;
    AGGREGATION_PROG_INTRO(ctx, xdp_ctx);

    AGGREGATION_PROG_ACQUIRE_LOCK(ctx);
    for (__u32 i = 0; i < RESP_MAX_VECTOR_SIZE; ++i) {
        ctx.current_value[i] += ((RESP_VECTOR_TYPE*) ctx.pk_msg->body)[i];
    }
    AGGREGATION_PROG_RELEASE_LOCK(ctx);

    AGGREGATION_PROG_OUTRO(ctx, DISCARD_PK);
}
```

## Example - custom aggregation program

```
#include "bpf_h/helpers.bpf.h"

SEC("xdp")
int aggregation_prog(struct xdp_md* xdp_ctx) {
    struct aggregation_prog_ctx ctx;
    AGGREGATION_PROG_INTRO(ctx, xdp_ctx);

    AGGREGATION_PROG_ACQUIRE_LOCK(ctx);
    for (__u32 i = 0; i < RESP_MAX_VECTOR_SIZE; ++i) {
        ctx.current_value[i] += ((RESP_VECTOR_TYPE*) ctx.pk_msg->body)[i];
    }
    AGGREGATION_PROG_RELEASE_LOCK(ctx);

    AGGREGATION_PROG_OUTRO(ctx, DISCARD_PK);
}
```

Helper macros to reduce boilerplate



## Example - custom aggregation program

```
#include "bpf_h/helpers.bpf.h"

SEC("xdp")
int aggregation_prog(struct xdp_md* xdp_ctx) {
    struct aggregation_prog_ctx ctx;
    AGGREGATION_PROG_INTRO(ctx, xdp_ctx);

    AGGREGATION_PROG_ACQUIRE_LOCK(ctx);
    for (__u32 i = 0; i < RESP_MAX_VECTOR_SIZE; ++i) {
        ctx.current_value[i] += ((RESP_VECTOR_TYPE*) ctx.pk_msg->body)[i];
    }
    AGGREGATION_PROG_RELEASE_LOCK(ctx);

    AGGREGATION_PROG_OUTRO(ctx, DISCARD_PK);
}
```

Aggregation takes place holding a per-request spinlock

## Example - custom aggregation program

```
#include "bpf_h/helpers.bpf.h"

SEC("xdp")
int aggregation_prog(struct xdp_md* xdp_ctx) {
    struct aggregation_prog_ctx ctx;
    AGGREGATION_PROG_INTRO(ctx, xdp_ctx);

    AGGREGATION_PROG_ACQUIRE_LOCK(ctx);
    for (__u32 i = 0; i < RESP_MAX_VECTOR_SIZE; ++i) {
        ctx.current_value[i] += ((RESP_VECTOR_TYPE*) ctx.pk_msg->body)[i];
    }
    AGGREGATION_PROG_RELEASE_LOCK(ctx);

    AGGREGATION_PROG_OUTRO(ctx, DISCARD_PK);
}
```

Early dropping specified

## Example - custom aggregation program

```
#include "bpf_h/helpers.bpf.h"

SEC("xdp")
int aggregation_prog(struct xdp_md* xdp_ctx) {
    struct aggregation_prog_ctx ctx;
    AGGREGATION_PROG_INTRO(ctx, xdp_ctx);

    AGGREGATION_PROG_ACQUIRE_LOCK(ctx);
    for (__u32 i = 0; i < RESP_MAX_VECTOR_SIZE; ++i) {
        ctx.current_value[i] += ((RESP_VECTOR_TYPE*) ctx.pk_msg->body)[i];
    }
    AGGREGATION_PROG_RELEASE_LOCK(ctx);

    AGGREGATION_PROG_OUTRO(ctx, DISCARD_PK);
}
```

Provided by the developer as a C file

Compiled into a BPF object file which is loaded into the kernel at runtime

## Example - user-space API to read the control socket

```
sgbpf::Service service{ ← Configure sgbpf
    "path/to/bpjobs",
    "lo",
    sgbpf::Worker::fromFile("workers.cfg"),
    sgbpf::PacketAction::Discard,           // packet fate post aggregation
    sgbpf::CtrlSockMode::DefaultUnix       // data delivery API
};

sgbpf::ReqParams params {
    .completionPolicy = sgbpf::GatherCompletionPolicy::WaitN,
    .numWorkersToWait = 5,
    .timeout           = std::chrono::microseconds{500}
}

sgbpf::Request* req = service.scatter(msg, msgLen, params);
```



## Example - user-space API to read the control socket

```
sgbpf::Service service{
    "path/to/bpjobs",
    "lo",
    sgbpf::Worker::fromFile("workers.cfg"),
    sgbpf::PacketAction::Discard,           // packet fate post aggregation
    sgbpf::CtrlSockMode::DefaultUnix       // data delivery API
};

sgbpf::ReqParams params {
    .completionPolicy = sgbpf::GatherCompletionPolicy::WaitN,
    .numWorkersToWait = 5,
    .timeout          = std::chrono::microseconds{500}
}

sgbpf::Request* req = service.scatter(msg, msgLen, params);
```

Invoke a scatter-gather request

## Example - user-space API to read the control socket

```
// Using sgbpf::CtrlSockMode::DefaultUnix
sg_msg_t buf;
read(service.ctrlSkFd(), &buf, sizeof(sg_msg_t));
handle(buf.body);
```

The application can use the raw Unix file descriptor of the control socket as it wishes  
(including configuring it as a non-blocking file)



## Example - user-space API to read the control socket

```
// Using sgbpf::CtrlSockMode::Ringbuf
service.setRingbufCallback([&](char* data, int reqID) -> void {
    handle(data);
});

while (1) {
    int completions = service.epollRingbuf(E POLL_TIMEOUT_MS);
    if (completions > 0)
        break;
}
```

The application can register a callback function which executes whenever aggregated data is written to the ring buffer (monitored by epoll)

## Example - user-space API to read the control socket

```
// Using sgbpf::CtrlSockMode::BusyWait
while (!req->isReady()) {
    service.processEvents(req->id());
}

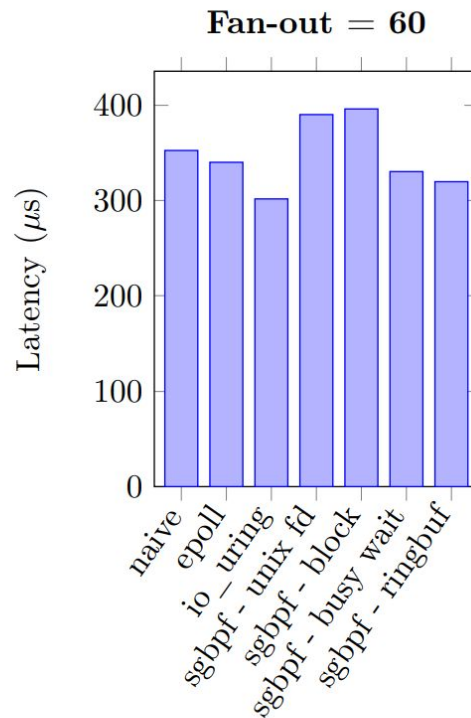
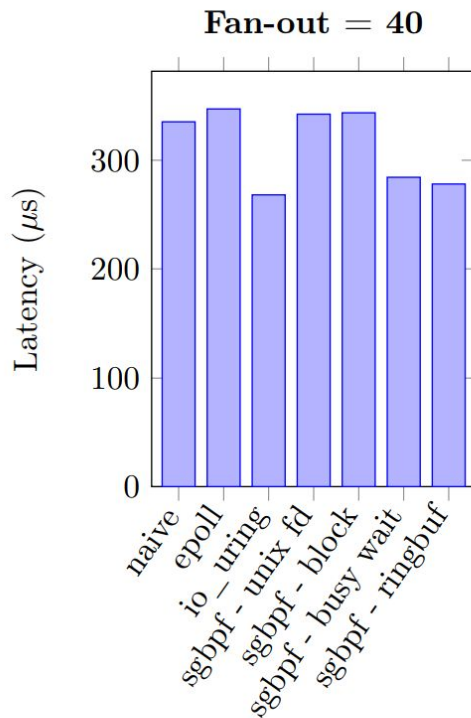
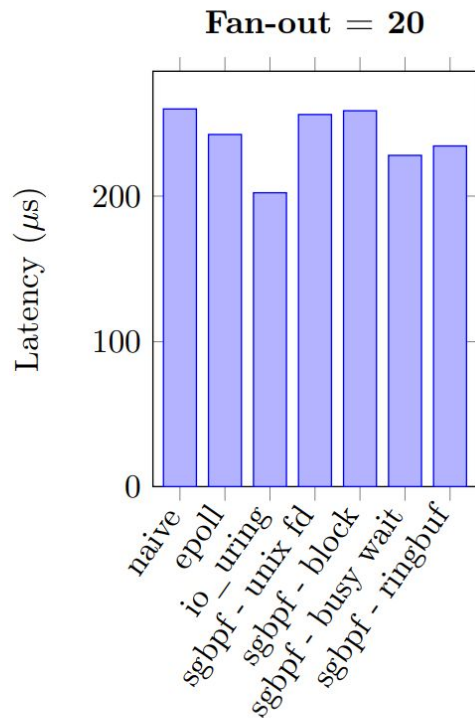
handle(req->ctrlSockData());
```

The application can busy-wait on the control socket using the methods provided in *sgbpf* without incurring any extra system calls

# Performance evaluation

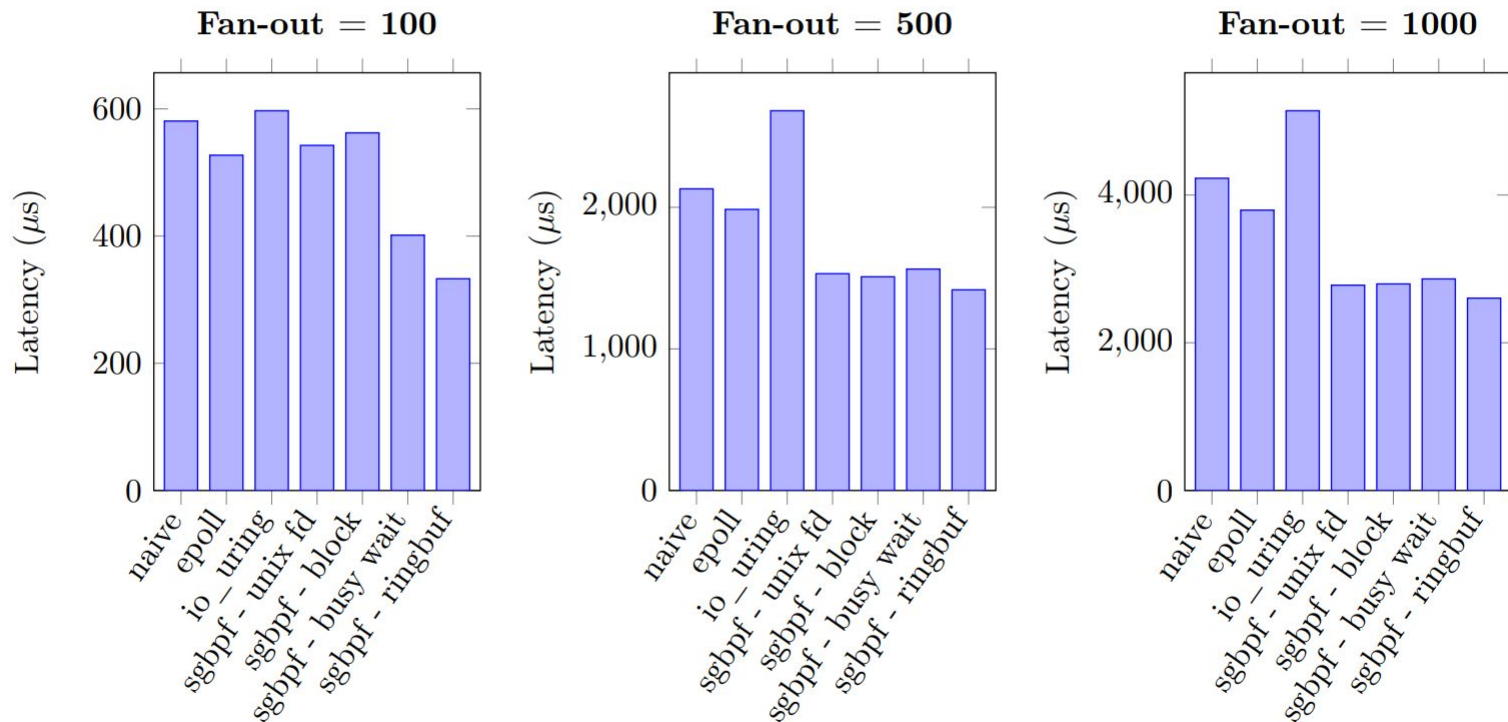
- Metrics of interest:
  - Unloaded latency
  - Throughput under load
- Compare performance with baseline implementations using standard Linux I/O APIs
  - Naive implementation
  - `epoll`-based implementation
  - `io_uring`-based implementation (with provided buffers)
- Results shown are obtained from local benchmarks
  - Workers are executed as local processes on a single machine

# Performance evaluation - unloaded latency



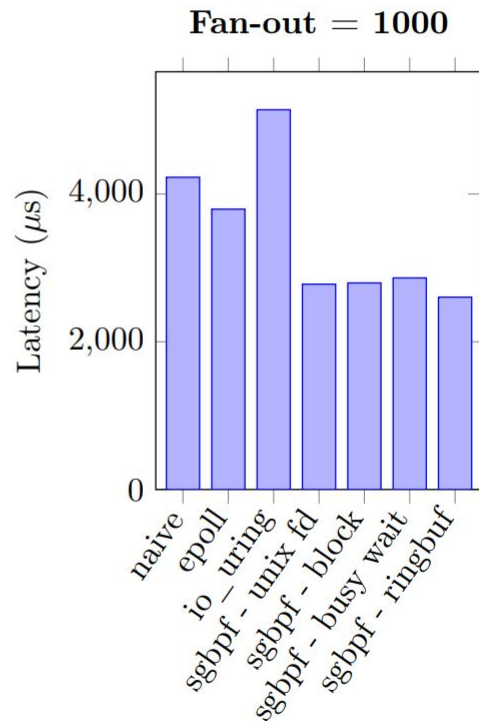
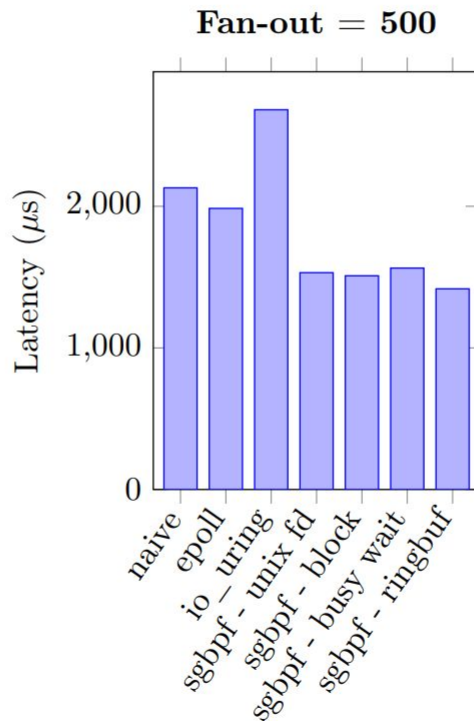
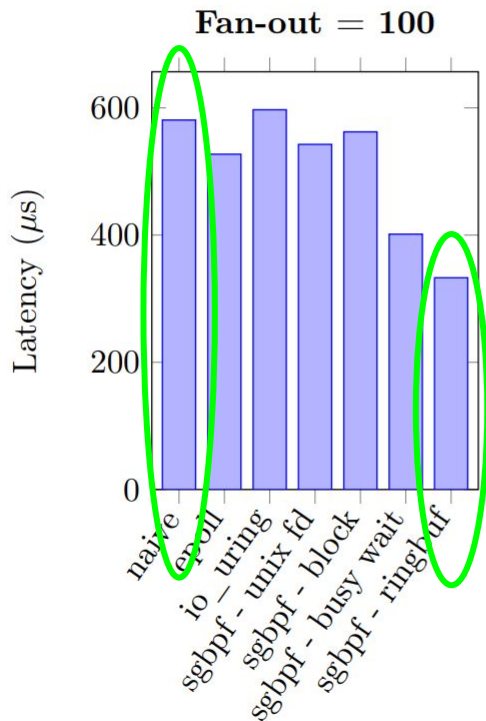
*sgbpf* achieves comparable latency results for small fan-outs

# Performance evaluation - unloaded latency



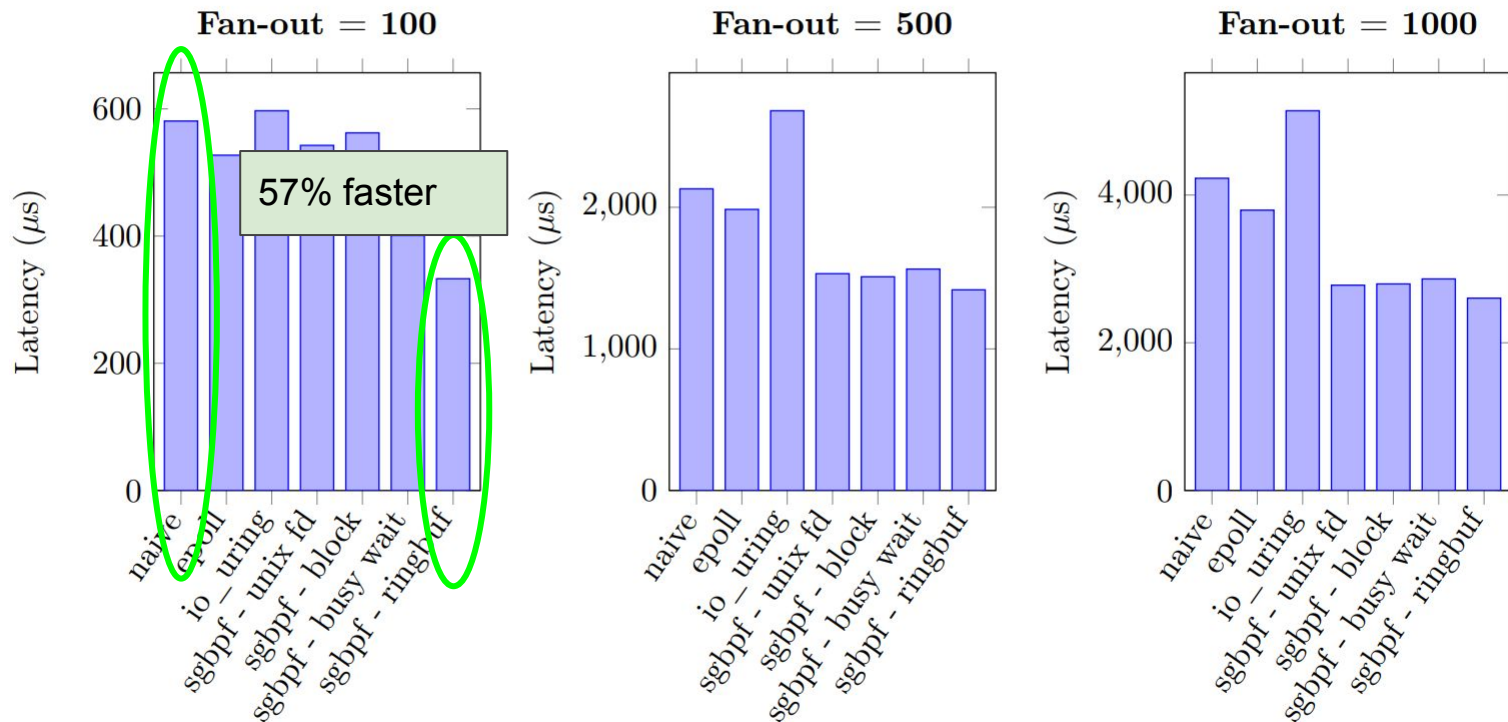
For larger fan-outs, *sgbpf* achieves substantial latency improvements

# Performance evaluation - unloaded latency



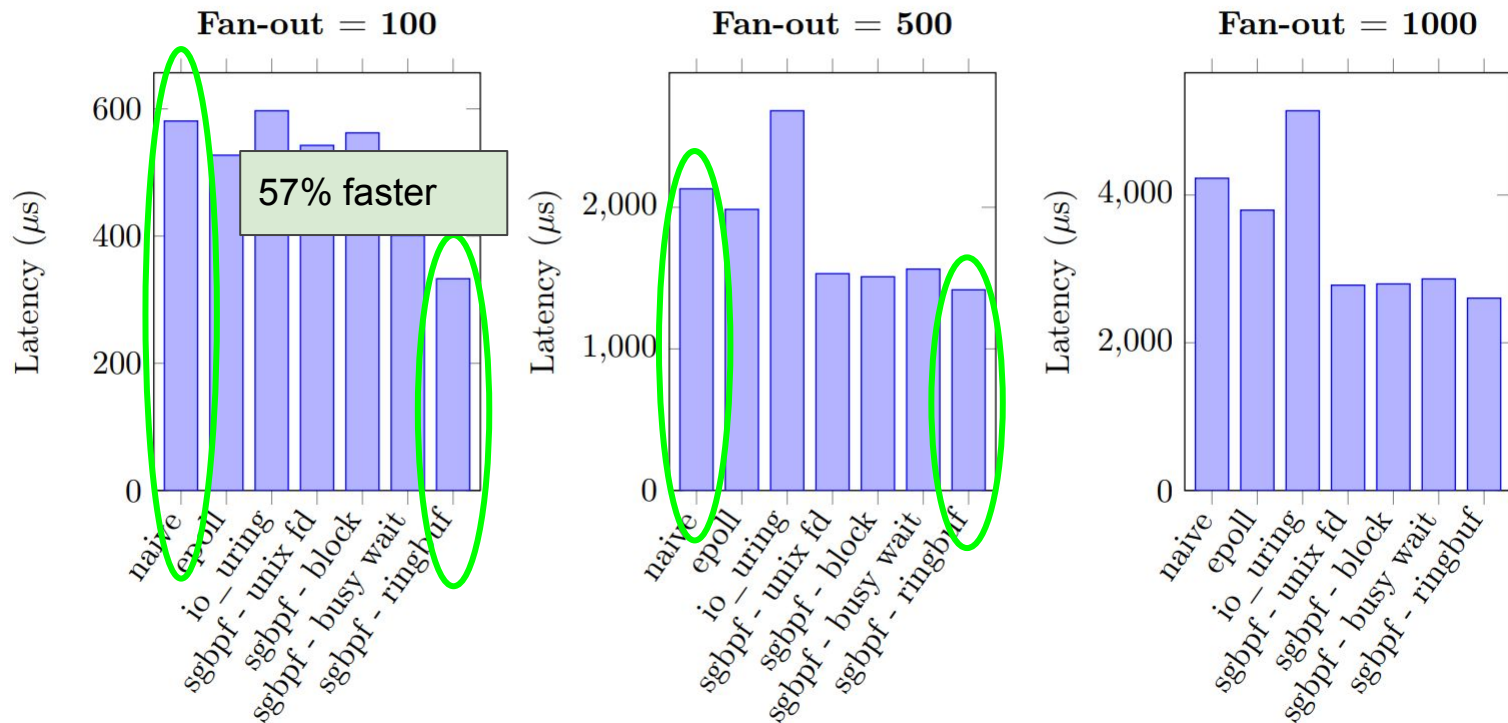
For larger fan-outs, *sgbpf* achieves substantial latency improvements

# Performance evaluation - unloaded latency



For larger fan-outs, *sgbpf* achieves substantial latency improvements

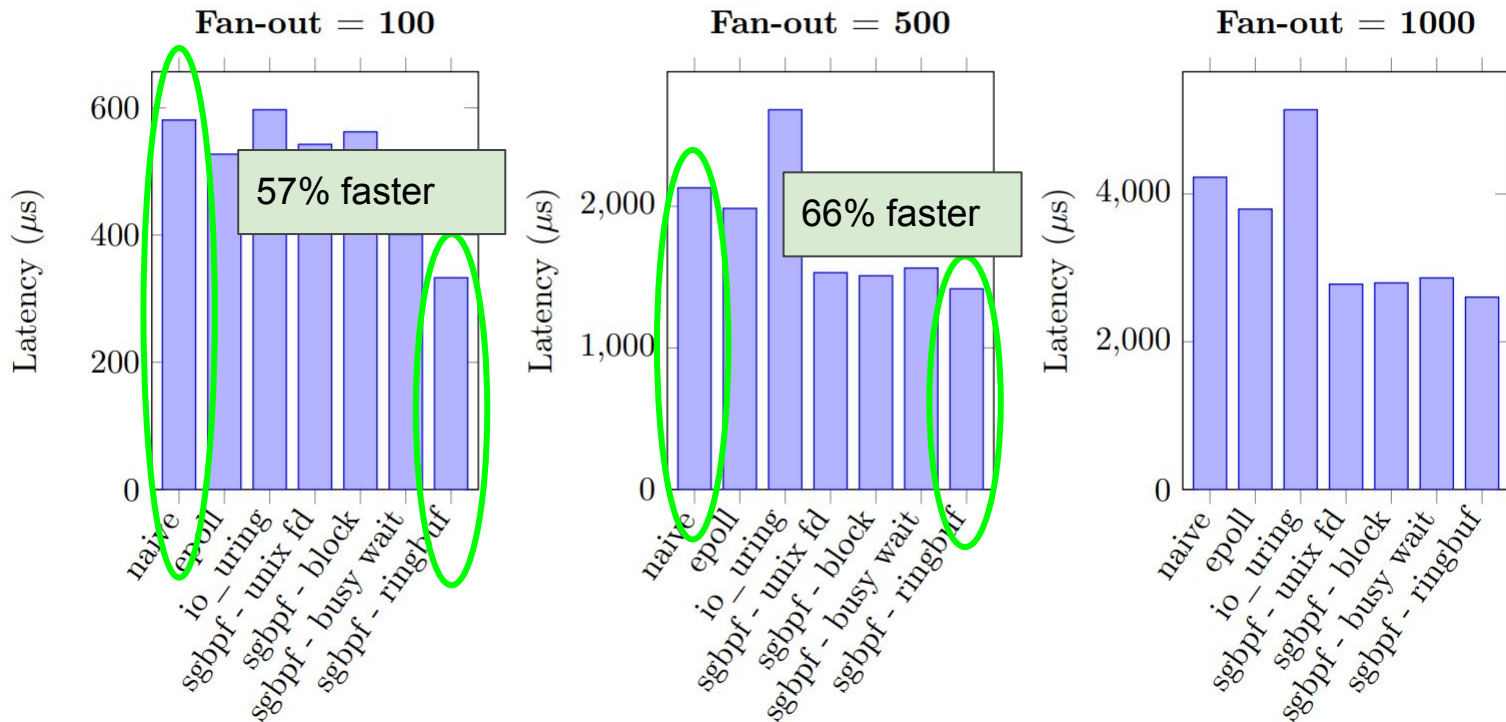
# Performance evaluation - unloaded latency



For larger fan-outs, *sgbpf* achieves substantial latency improvements

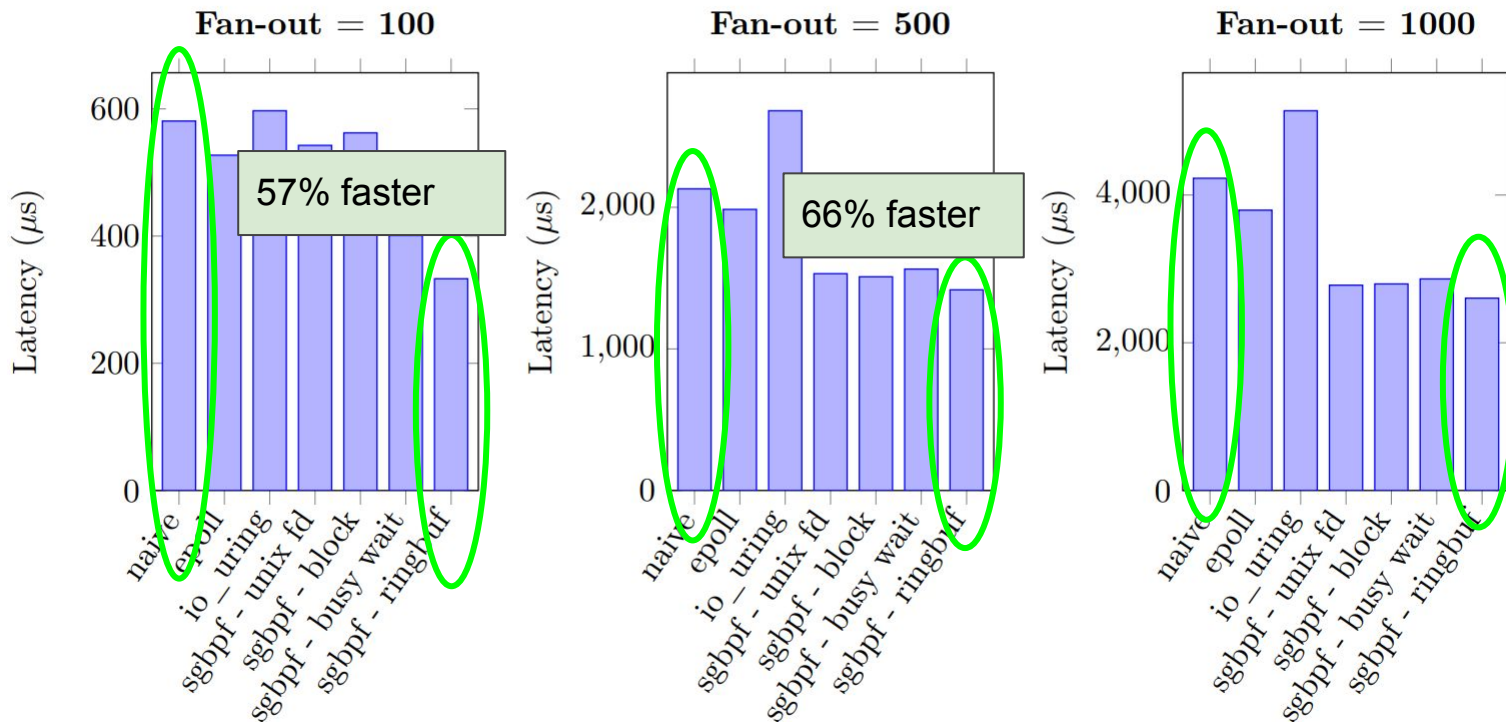


# Performance evaluation - unloaded latency

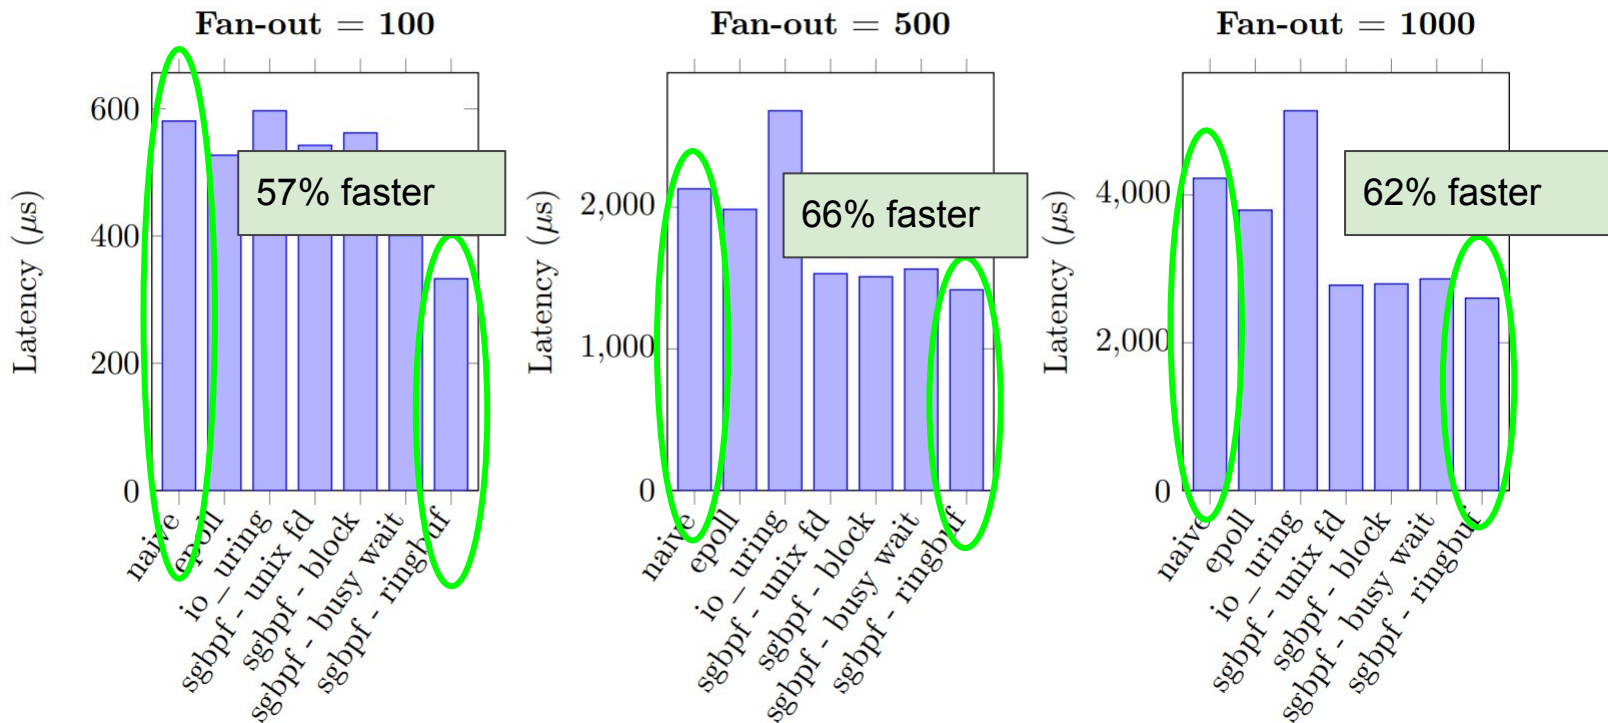


For larger fan-outs, *sgbpf* achieves substantial latency improvements

## Performance evaluation - unloaded latency

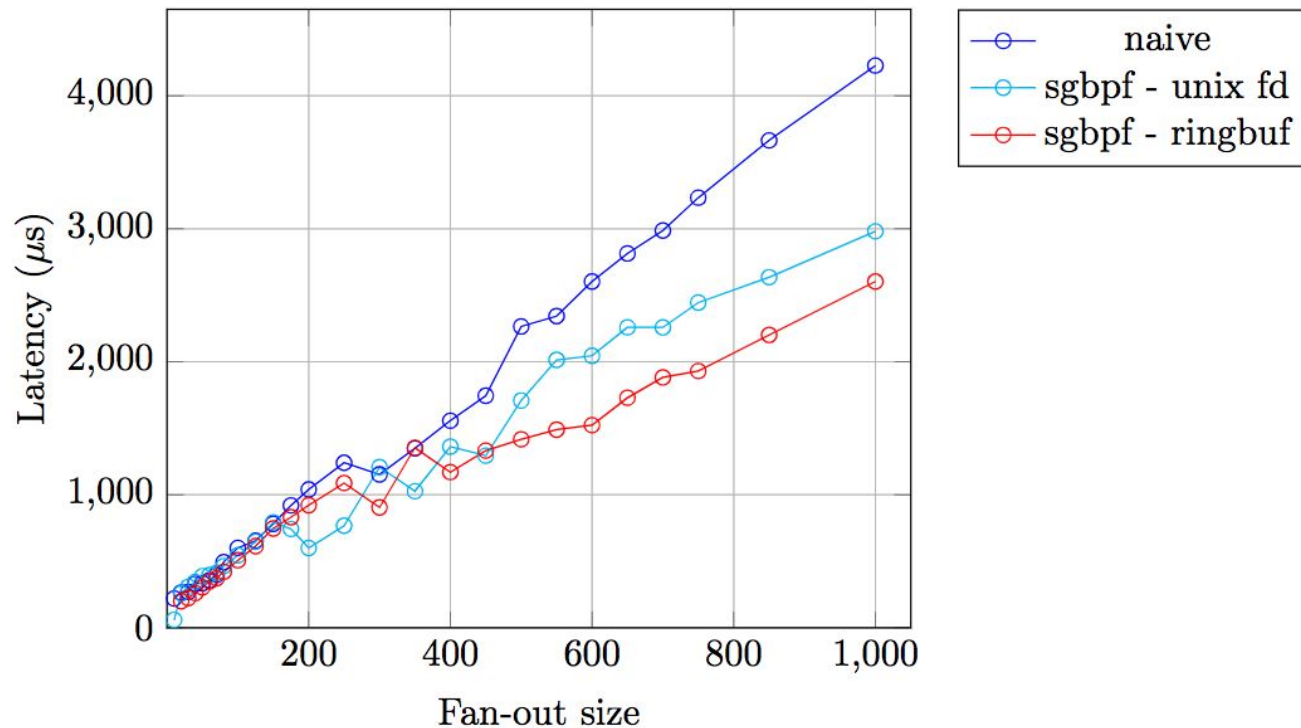


## Performance evaluation - unloaded latency



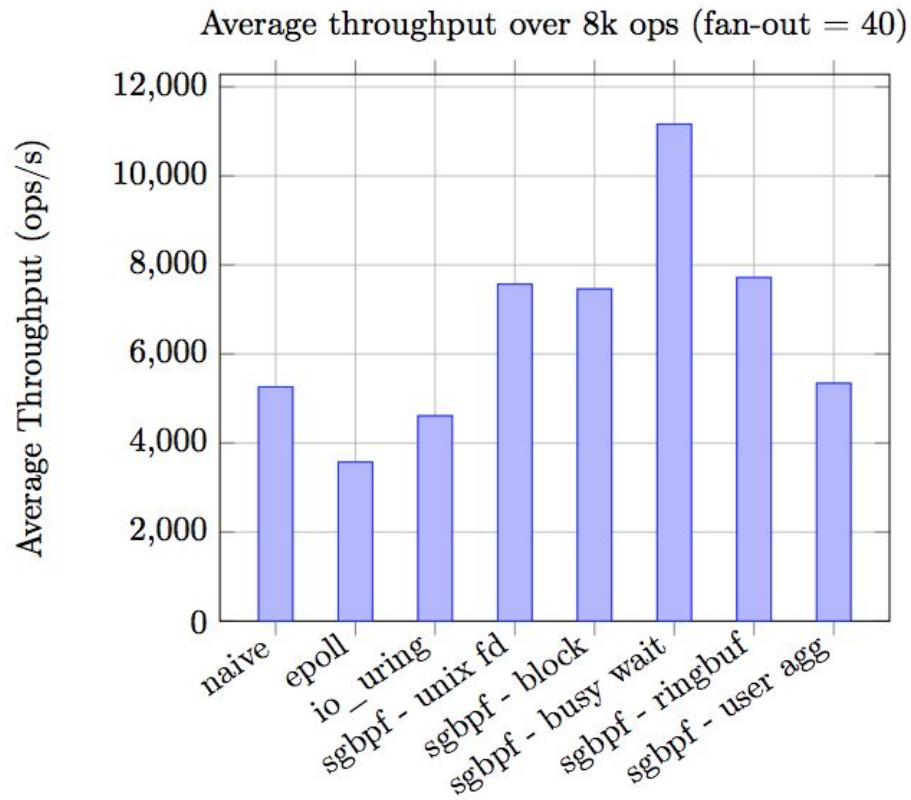
For larger fan-outs, *sgbpf* achieves substantial latency improvements

## Performance evaluation - unloaded latency




*sgbpf* scales better than standard I/O APIs due to total kernel overhead

# Performance evaluation - throughput under load




# Performance evaluation - throughput under load

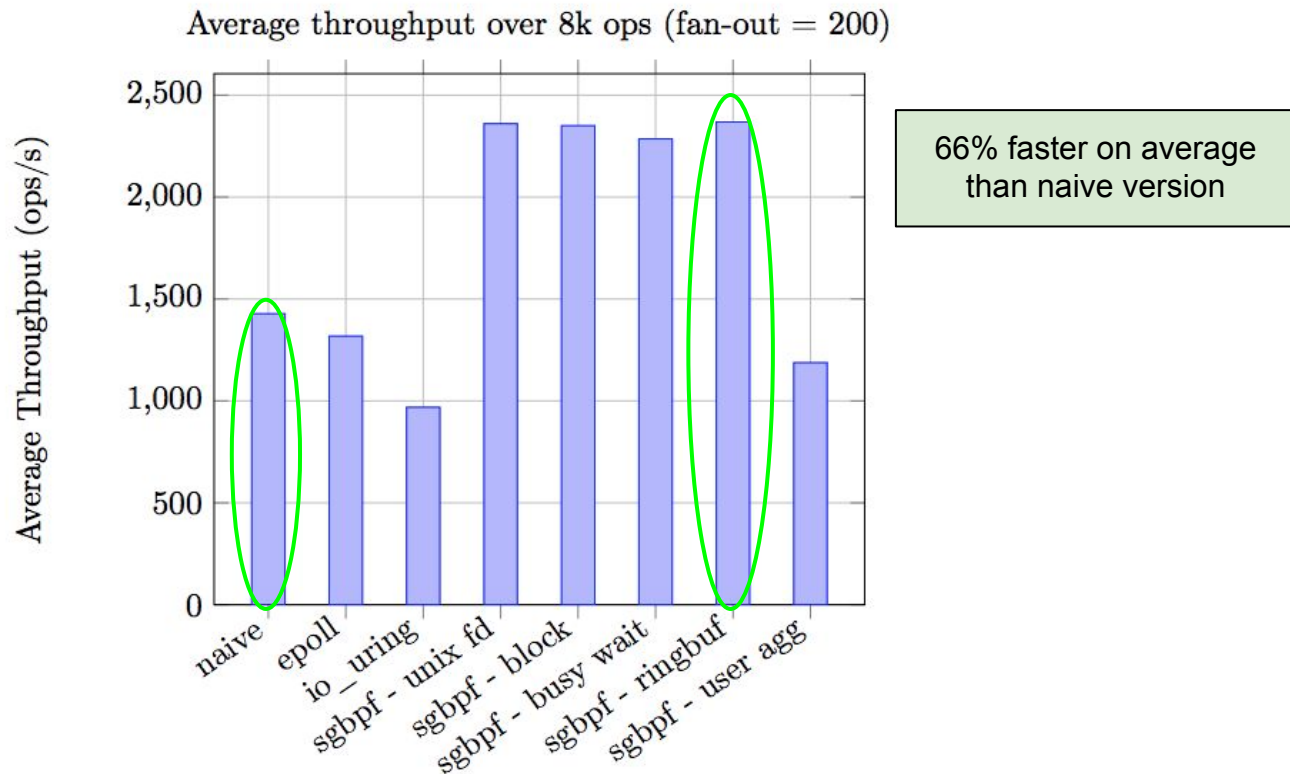


Even at small fan-outs, sgbpf achieves at least 46% better throughput

## Performance evaluation - throughput under load




## Performance evaluation - throughput under load



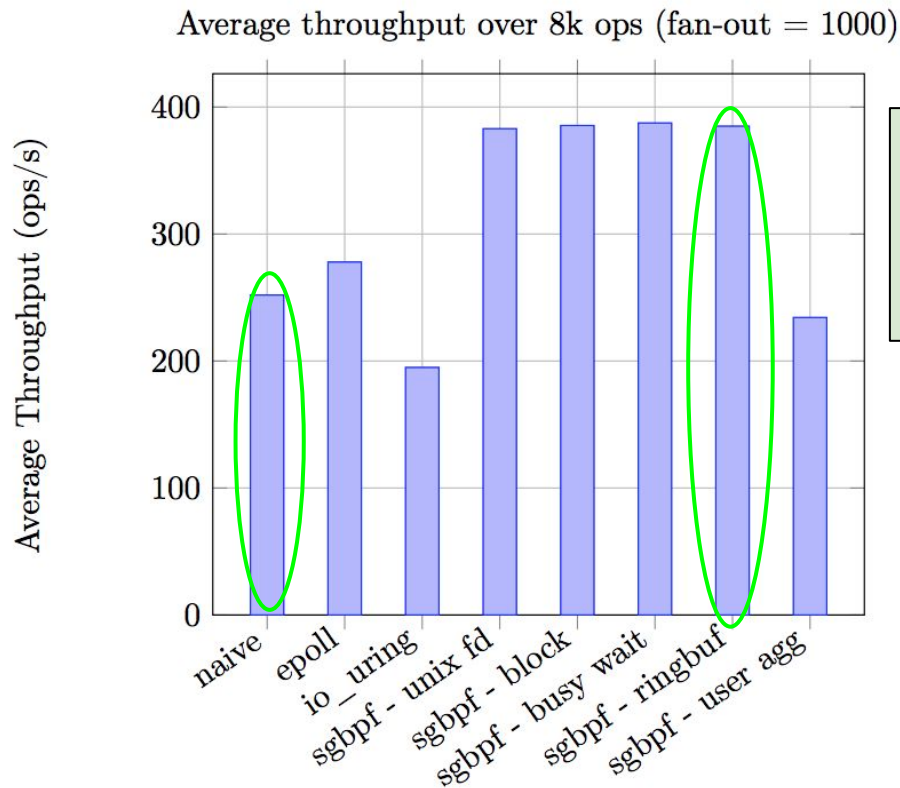
As the fan-out increases, throughput difference starts to increase



# Performance evaluation - throughput under load



# Performance evaluation - throughput under load



53% faster on average than naive version

42% faster than epoll version (best standard baseline)

At extremely large fan-outs, the relative difference decreases slightly but still comfortably beats the standard APIs

# Conclusions

- *sgbpf* reduces the kernel bottleneck in scatter-gather communications using eBPF
  - Minimises user-kernel crossings and kernel network stack traversals
  - Requires at most 2 syscalls (depends on API used)
  - Mainly achieved by performing aggregation as soon as possible inside the kernel
- While previous work shows eBPF can be used to accelerate specific applications, this project shows that generic high-level communication patterns can also incur performance benefits without resorting to extreme solutions.
- The code is open-sourced on [Github](#)

---

## Future work

- Supporting in-kernel aggregation for multi-packet responses
- TCP implementation for responses
- Other optimisations and techniques for higher performance (see report)

# Thank you for listening

*sgbpf* source code is available [here](#) on Github

---

# References

- [1] Dean J, Barroso LA. The tail at scale. Communications of the ACM. 2013 Feb 1;56(2):74-80.
  
  - [2] Yoann Ghigoff, Julien Sopena, Kahina Lazri, Antoine Blin, and Gilles Muller. BMC: Accelerating Memcached using Safe In-kernel Caching and Pre-stack Processing. In 18th USENIX Symposium on Networked Systems Design and Implementation (NSDI 21), volume 3, page 4, 2021.
  
  - [3] Yang Zhou, Zezhou Wang, Sowmya Dharanipragada, and Minlan Yu. Electrode: Accelerating Distributed Protocols with eBPF. In 20th USENIX Symposium on Networked Systems Design and Implementation (NSDI 23), pages 1391–1407, 2023.
-